

# Distributed Systems Direct Communication PART I

Wouter Joosen

DistriNet, KU Leuven

October 03, 2023



# Overview of chapters

---

- Introduction
- Coordination models and languages: direct communication
  - Ch 4: Inter-process communication (only a small part)
  - Ch 5: Remote invocation
  - Ch 8: Distributed objects and components

(Assumes knowledge of computer networks: Chapter 3-4)



note:

## CHAPTER 3 – HELICOPTER VIEW

---

1. Intro (concepts, terms)
2. Types of Network (LAN, WAN etc).
3. Network Principles: packet transmission, data streaming, switching, protocols, routing, congestion control, internetworking
4. Internet Protocols: IP addressing IP protocol, IP routing, IPv6, Mobile IP, TCP and UDP
5. Case studies: Ethernet, WiFi and Bluetooth
6. Summary

All assumed to be studied before.

note:


## CHAPTER 4 – HELICOPTER VIEW

1. Intro
2. The API for the Internet protocols
  - Assumed to be studied before
3. External data representation and marshalling
  - Covered in this part (with chapter 5)
4. Multicast communication (skipped)
5. Network virtualization: Overlay networks (skipped)
6. Case study: MPI (skipped)
7. Summary



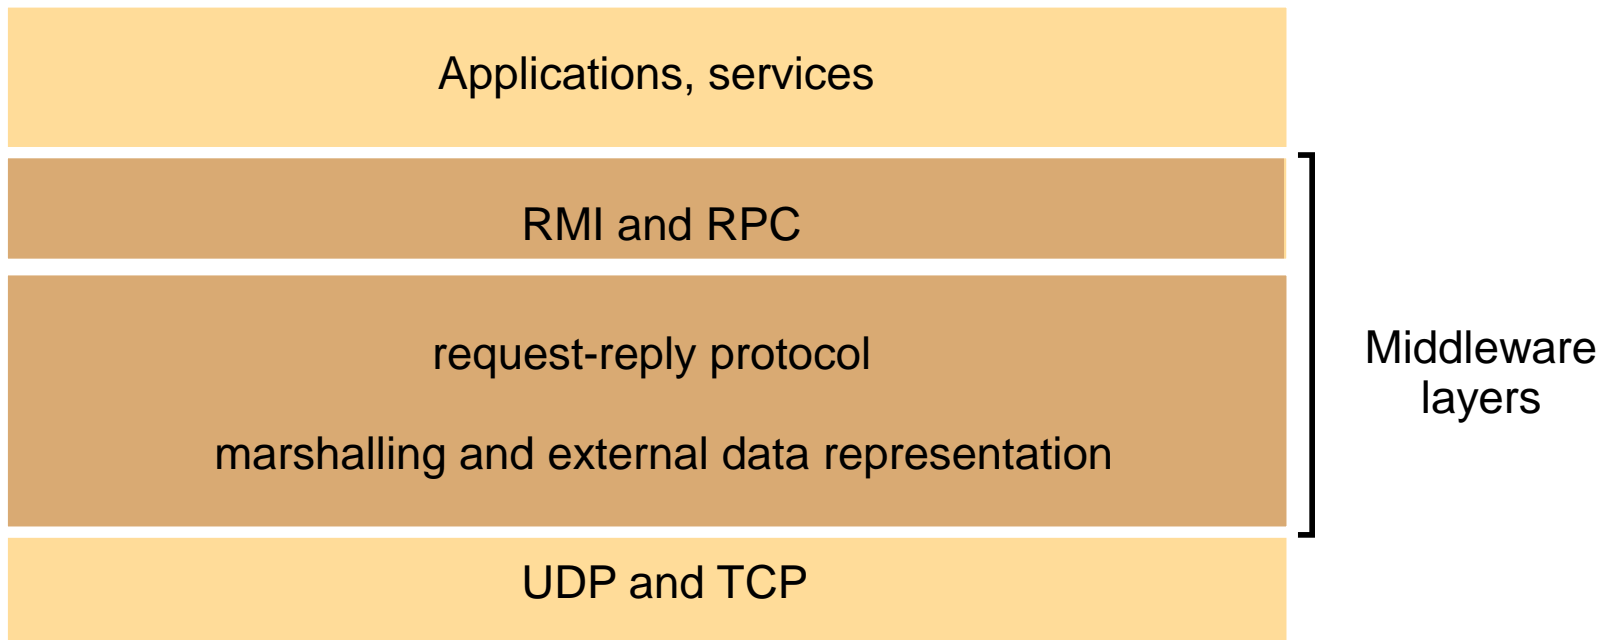
# Introduction

- Communication
  - data representation
- Synchronization
  - how express cooperation?



# Introduction

Distribution service in middleware (*subset of the larger middleware setting*)



# Introduction

---

- Distribution service in middleware
  - shields developers of a distributed application from the complex distributed environment, e.g.
    - Low-level socket API
    - No transfer of structured data
    - ...
  - by offering programming abstraction layers on top of OS
    - Typical programming paradigms incorporated in distributed model (syntactically)
    - Abstraction of heterogeneity of systems



# This lecture: overview

---

- Data representation
- Message passing
- Request-reply protocols
- Remote procedure calls





# Data representation

- Problem



int = 2-complement, 32 bits

int = 1-complement, 40 bits

on the wire: int = ??

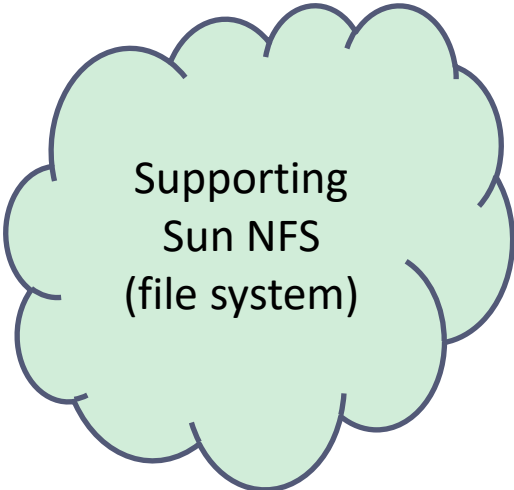
# Data representation *(cont.)*

- Problem:
  - program data: typed, structured, object
  - message data: bit/byte stream
  - different representations of data in heterogeneous systems
- mapping to data items in messages:
  - flattened before transmission
  - rebuilt on arrival



# Data representation (cont.)

- Conversion: different approaches
  - agreed form for transmission (implicit information)
    - e.g. int = 2-complement, 32 bits
    - both partners have full knowledge of transmitted data
    - e.g.: CORBA CDR, Sun XDR
  - full data description transmitted
    - *type, length, value* coding on the wire
    - interpretation at receiving site possible
    - e.g. ASN + BER, Java serialized form
  - Conversion to ASCII text
    - XML



Supporting  
Sun NFS  
(file system)



# Data representation *(cont.)*

- Java serialized form
  - Handles: references to objects (within serialized form)
  - Primitive types: portable binary format

```
Person p = new Person("Smith", "London", 1934);
```

*Serialized values*

Person	8-byte version number		h0
3	int year	java.lang.String name:	java.lang.String place:
1934	5 Smith	6 London	h1

*Explanation*

*class name, version number*

*number, type and name of  
instance variables*

*values of instance variables*

(The true serialized form contains additional type markers; h0 and h1 are handles)

# Data representation *(cont.)*

- XML (e.g. used in Web Services for SOAP)
  - Markup language defined by World Wide Web consortium
  - Data items tagged with 'markup' strings
  - Users can define their own tags

```
<person id='123456789'>  
    <name>Smith</name>  
    <place>London</place>  
    <year>1934</year>  
    <...>  
</person>
```

# Data representation *(cont.)*

- Definitions:
  - **marshalling** = assembling a collection of data items into a form suitable for transmission
  - **unmarshalling** = disassembling a message on arrival to produce the equivalent collection of data items
- operations can be generated from specification
- In Java: serialization & deserialization



# This chapter: overview

---

- Data representation
- Message passing
- Request reply protocols
- Remote procedure calls
- ... Object request brokers (later)



# Message passing

---

- Basic functionality
  - Procedure Send  
(p: PortId; m: Message);
  - Procedure Receive  
(p: PortId; VAR m: Message);





# Message passing *(cont.)*

- **Semantics:**

synchronous  $\Leftrightarrow$  asynchronous communication

- synchronous = blocking
  - **send**: wait for corresponding receive
  - **receive**: wait for message arrival
- asynchronous = no waiting for completion
  - **send**: no wait for message arrival
  - **receive**: announce willingness to accept  
or check for message arrival



# Message passing *(cont.)*

- Semantics: synchronous  $\Leftrightarrow$  asynchronous

type	Blocking Send	Blocking Receive	Language System
Syn	Yes	Yes	occam
Syn	Yes	No	-
Asyn	No	Yes	Mach Chorus
Asyn	No	No	Charlotte



# Message passing *(cont.)*

- Semantics: **synchronous**  $\Leftrightarrow$  asynchronous

Example: Occam style

Sender:

.....

*Send(p, m);*

*{message is accepted!!!}*

.....

Receiver:

....

*Receive(p, b);*

*{sender after send }*

....


Communication  $\rightarrow$  Synchronisation point

**b := m**



# Message passing *(cont.)*

- Semantics: **synchronous**  $\Leftrightarrow$  asynchronous



# Message passing *(cont.)*

- Semantics: synchronous  $\Leftrightarrow$  asynchronous

Example: Mach style

Sender:

.....  
*send(...);*  
*{message is in buffer!!!*  
*arrival??}*

Receiver:


....  
*receive(...);*  
*{message available }*  
....

Communication  $\rightarrow$  NO synchronisation point




# Message passing *(cont.)*

- Semantics: synchronous  $\Leftrightarrow$  asynchronous



# Message passing *(cont.)*

- Semantics: synchronous  $\Leftrightarrow$  asynchronous



# Message passing *(cont.)*

---


- **Semantics:** message destinations
  - message destination = communication identifier
    - preference for location independent identifiers
  - types of message destination:
    - process
    - port
    - mailbox






# Message passing *(cont.)*

- Semantics: message destinations  
**process**



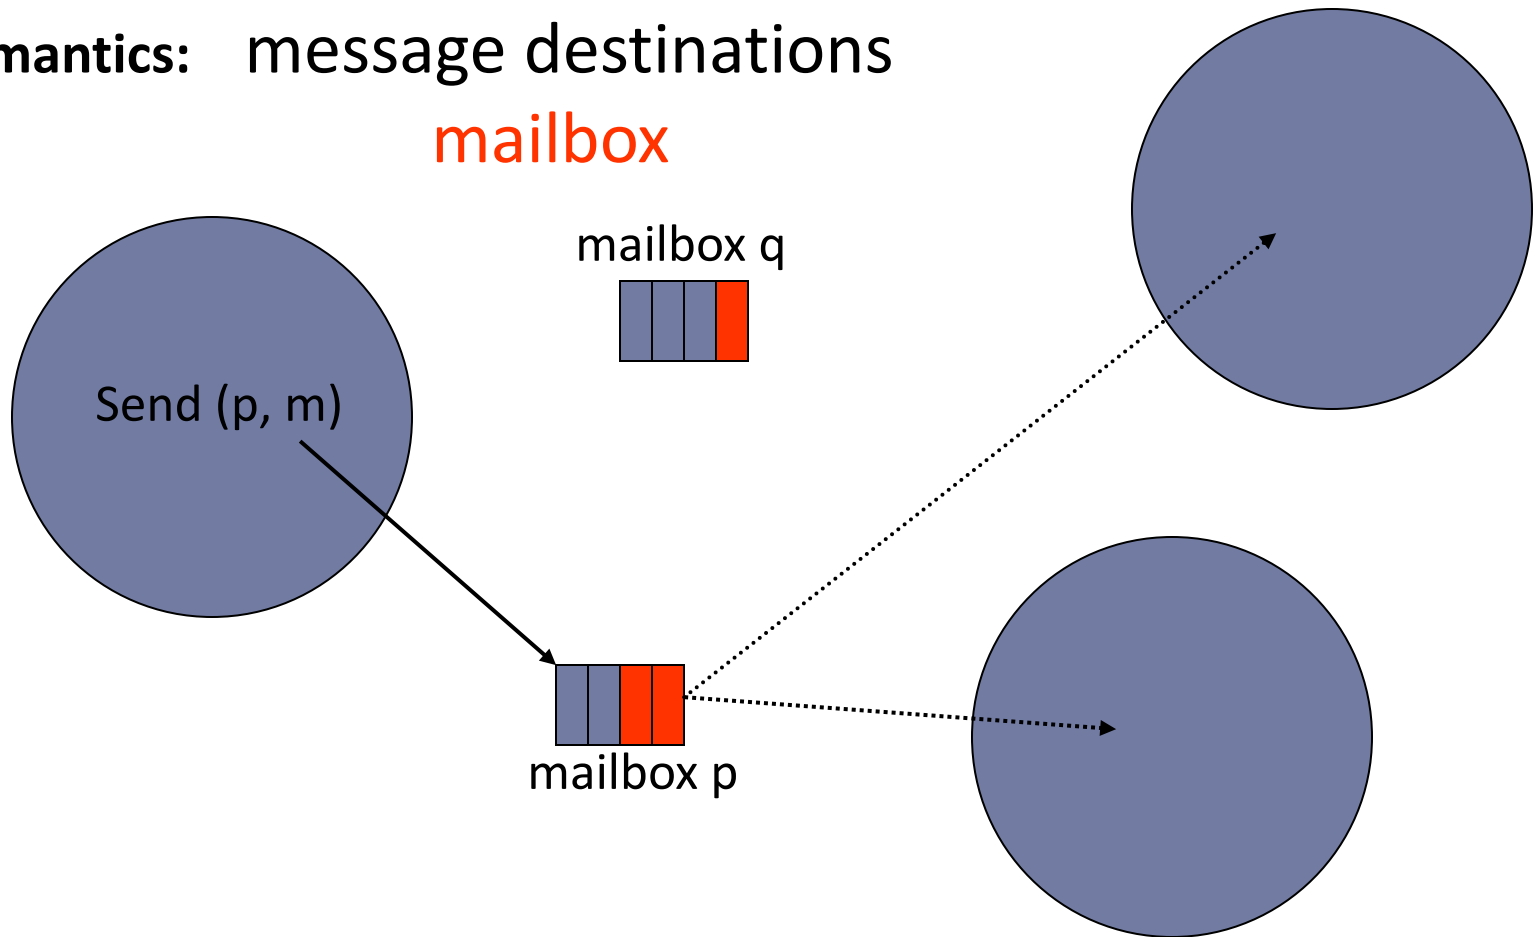
# Message passing *(cont.)*

- Semantics: message destinations  
**port**



# Message passing *(cont.)*

- Semantics: message destinations  
**mailbox**



# Message passing *(cont.)*

- **Semantics:** message destinations
  - types of message destination:
    - process:
      - single entry point per process for all messages
    - port
      - one receiver, many senders
      - may have a message queue
      - many ports per process
    - mailbox
      - may have many receivers
      - message queue



# Message passing (cont.)

- **Semantics:** reliability

- possible failures

- Corrupted messages
    - Duplicate messages
    - Omission: loss of messages
    - Messages out of order
    - Receiver process failure

Communication failure

- Reliable communication

- Delivered uncorrupted, in order, without duplicates
    - Despite a *reasonable* number of packets dropped or lost
    - Perfectly reliable communication can not often be guaranteed



# Message passing *(cont.)*

- **How to implement reliable communication:**
  - Avoiding corruption
    - Include checksum in message
  - Avoids order mistakes and duplicates
    - Include a message number which identifies the message
  - Avoiding omission
    - Sender stores message in buffer, sends it and sets a time-out
    - Receiver replies with acknowledgement
    - Sender retransmits messages after timeout



# Message passing *(cont.)*

---

- **Semantics: reliability**
  - sender of message gets no reply?


no distinction between

- process failure
- communication failure



# Message passing *(cont.)*

- **Case study: UDP/TCP**
  - Sockets  $\Leftrightarrow$  ports
    - Socket bound to (TCP) port + Internet address





# Message passing *(cont.)*

- **Case: UDP**
  - Messages:
    - Restricted packet size:  $< 2^{16}$  ,  $< 2^{13}$  (8Kbytes), truncation
    - No conversion: bit transmission
  - Synchronization semantics:
    - Non-blocking send
    - Blocking receive
  - Timeouts: user can set timeout on receive operation
  - Receive from any: receive returns port + Internet address
    - But sockets can be bound to remote (IP address+port)
  - Unreliable message service
    - lost, out of order, duplicates
    - no message corruption: checksum



# Message passing *(cont.)*

- **Case: TCP**

- Stream communication: ( $\Leftrightarrow$  message passing?)
  - Connect: create a communication channel through communicating sockets
  - Communication: read and write through channel
  - Close
- Implementation:
  - TCP handles all communication
  - Uses buffers at sender and receiver side
  - No conversion: bit transmission
- Synchronization semantics:
  - Non-blocking send, except for flow control  
(when buffers of sender or receiver are full)
  - Blocking receive



# Message passing *(cont.)*

- **Case: TCP**

- Setting up a client server connection:
  - Client sends request for communication to Server port
  - Server accepts client request
  - Typically, server creates new thread which handles communication with client
- Reliable message service
  - Except broken connections
- Overhead compared to UDP:
  - Buffering
  - Creating a connection: 2(?) extra messages
    - Sending a message, returning an acknowledgement
    - May create unacceptable overhead if goal is to send a single message.



# Message passing *(cont.)*

---

- Conclusion: UDP, TCP
  - general purpose communication protocols
  - primitive, low level operations:
    - Setting up a communication
    - No transfer of structured data
  - Difficult to use
  - efficient implementation
  - building blocks used for more complex interactions



# Message passing *(cont.)*

---

- Conclusion: message passing
  - primitive, low level operations
  - difficult, hard to use
  - efficient implementation
  - building blocks used for more complex interactions
- From message passing
  - to **Client-server** (Request reply protocols)
  - to RPC, RMI



# This lecture: overview

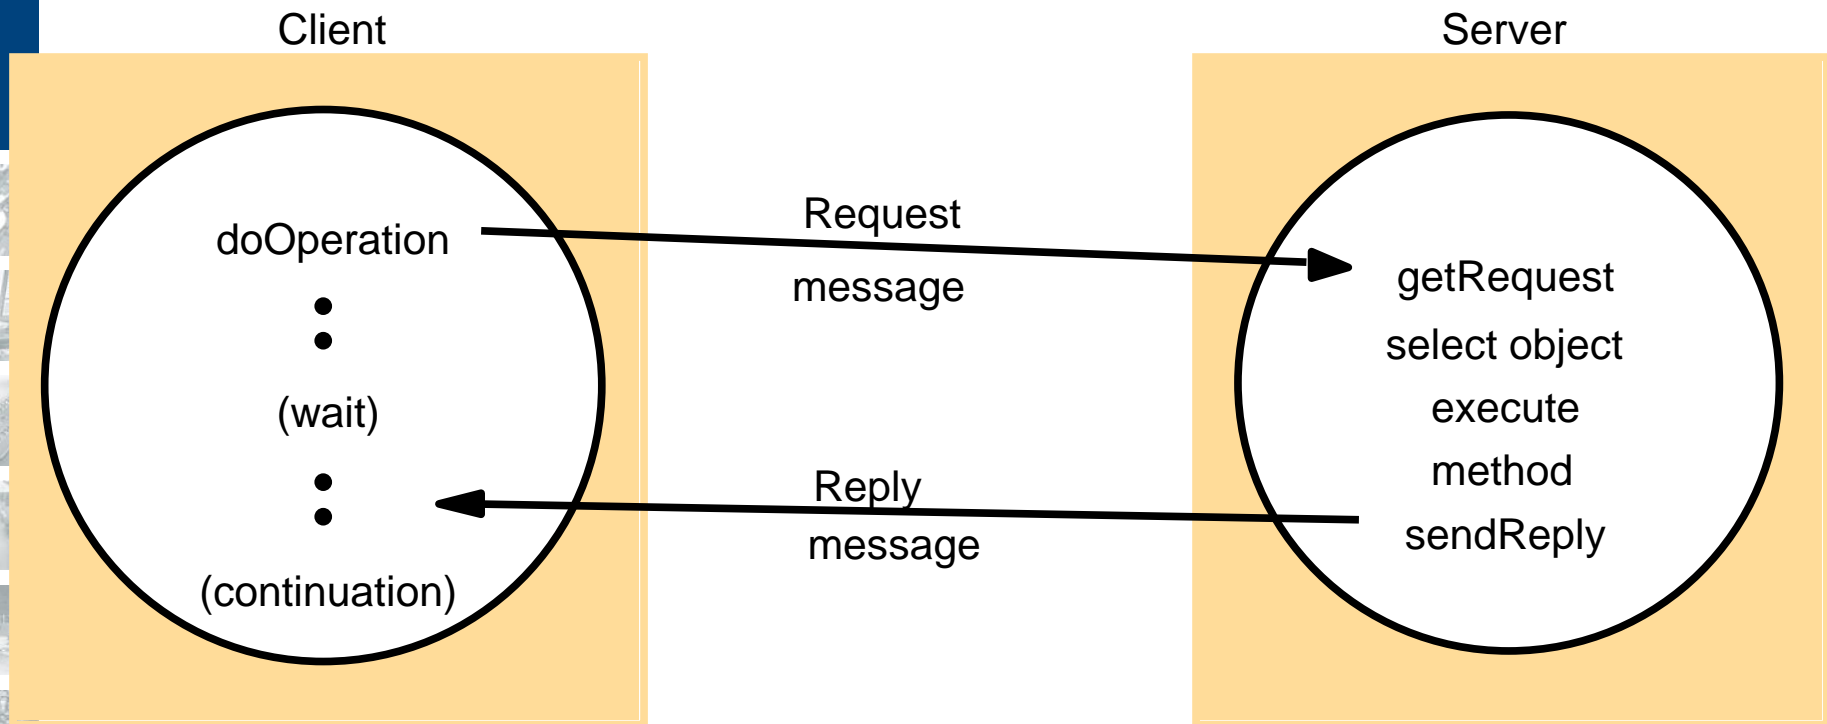
---

- Data representation
- Message passing
- Request-reply protocols
- Remote procedure calls



# Request-Reply protocols

- Client-server communication



# Request-Reply protocols

- Client-server messages/operations
  - Designed to support roles and message exchanges in typical client-server interactions
    - Acks redundant (replies are used)
    - Connections not necessary
    - No flow control
  - Basic operations:
    - Client: **doOperation**: sends request and returns answer to application program
    - Server: **getRequest**, **sendReply**





# Request-Reply protocols (*cont.*)

## *Implementation options: on TCP and UDP*

### – Reliability measures of TCP are an overkill!!

- Acknowledgement of receiver is redundant : the reply message is an acknowledgement!
- Limited size of data packet transfer
- One time communication

=> Making a connection is overhead

=> No stream needed, no flow control

⇒ Use of TCP is (often) an overkill and may cause efficiency problems!

⇒ UDP can be used for building more efficient client server communication.

- What about reliability??



# Request-Reply protocols *(cont.)*

*Comparison: the case of implementing HTTP on top of TCP*

- Case: Client-server communication: HTTP
  - Interactions:
    - Open connection
    - Client sends request
    - Server sends reply
    - Connection closed
    - ✓ New since HTTP1.1: Persistent connections using TCP (for multiple requests at same server)
  - HTTP methods:
    - Get(URL) : request for the resource referred to by URL
    - Post(URL,data): replace or create resource at URL
    - ..



# Message Passing and Request-Reply protocols: Conclusion

---

- So far rather low level implementations:
  - Setting up a connection
  - No transfer of structured data!!!
  - Concerns for synchronisation and failure model
  - No encryption of data
  - ...
- Higher level Message Passing systems exist!!
  - MPI, Mach, ..



# This chapter: overview

---

- Data representation
- Message passing
- Request-Reply protocols
- Remote procedure calls



# This chapter: overview

---

- Data representation
- Message passing
- Request-Reply protocols
- Remote procedure calls



# Remote procedure calls

## Overview

---

- Basic principles
- Design issues
- Implementation aspects
- Asynchronous RPC



# Remote procedure calls

## Basic principles

---


- View on traditional application:
  - no OO yet!
  - application =
    - main program
    - + procedures (functions)
  - ➔ familiar paradigm: **procedure call**



# Remote procedure calls

## Basic principles

- View on traditional application:





# Remote procedure calls

## Basic principles

---

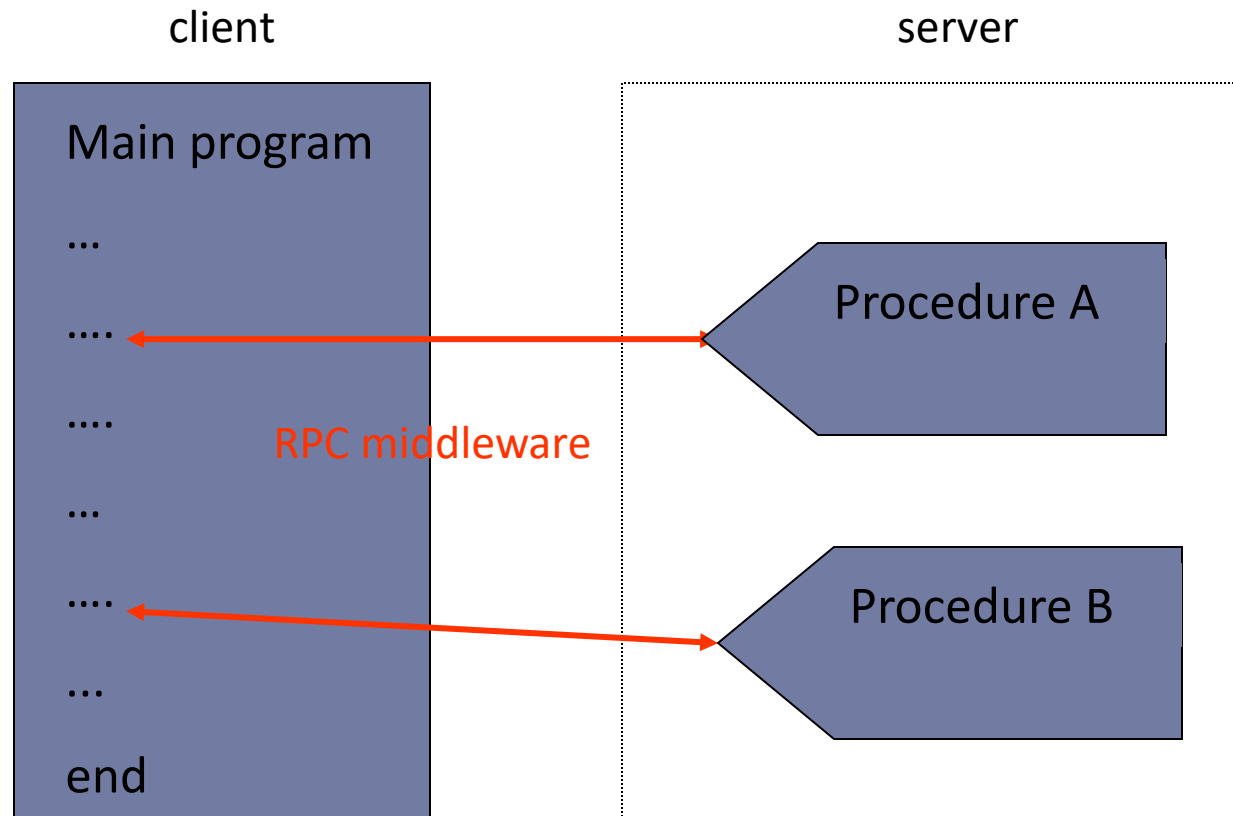
- View on traditional application:
  - main program and procedures (functions)
  - familiar paradigm: procedure call
- approach in distributed systems:
  - group procedures into **servers**
  - main programs become **clients**
  - operations on server look like **conventional procedure calls**



# Remote procedure calls

## Basic principles

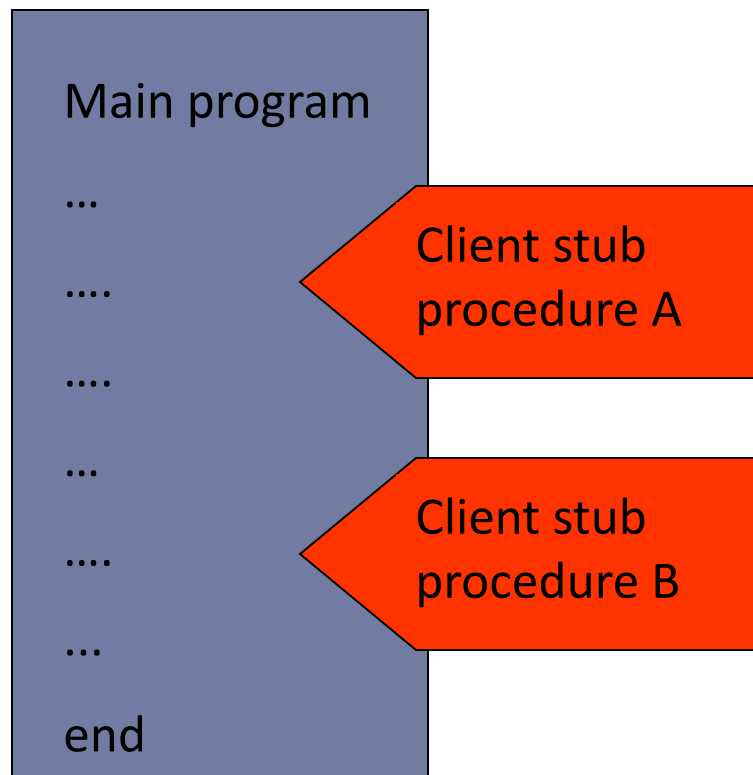
- Basic view on client-server model



# Remote procedure calls

## Basic principles

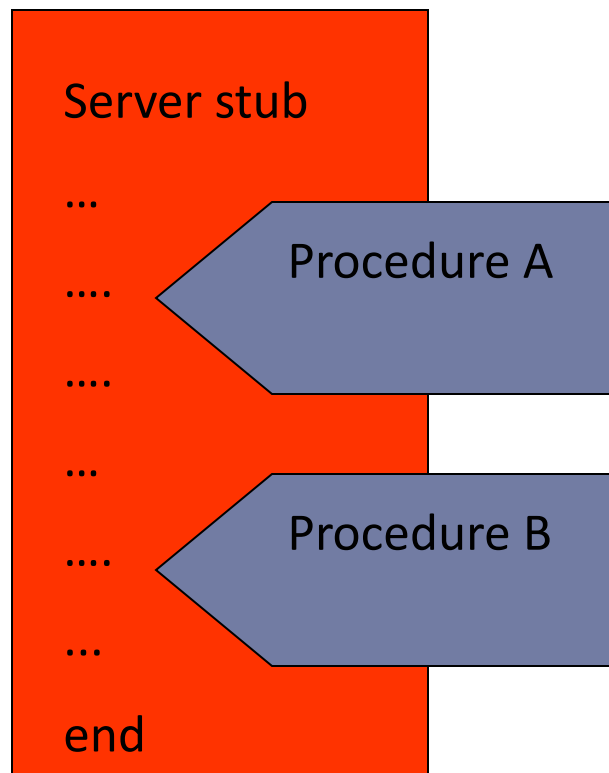
- RPC technology: client side



# Remote procedure calls

## Basic principles

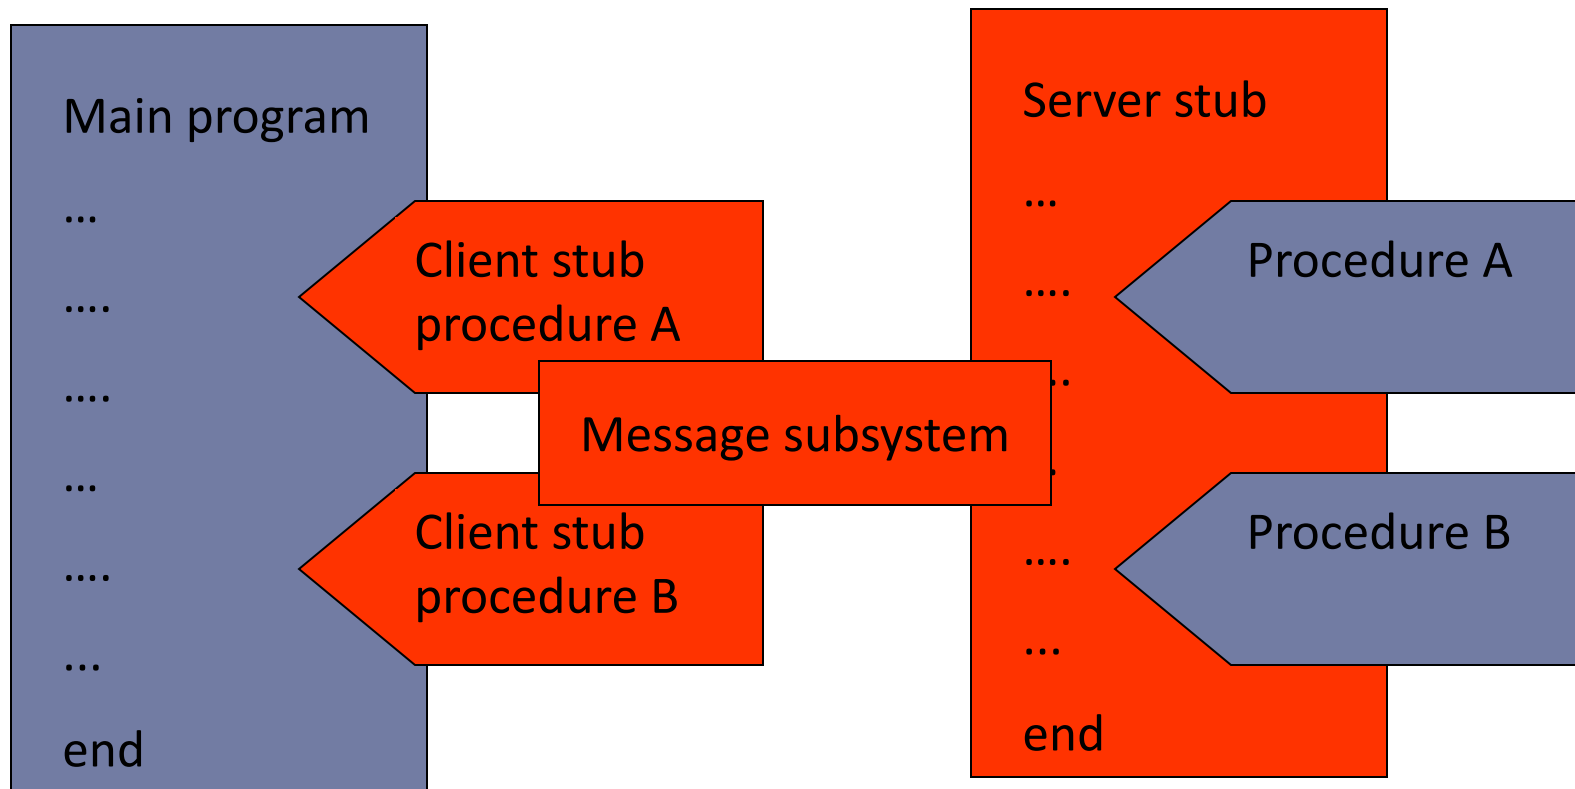
- RPC technology: server side



# Remote procedure calls

## Basic principles

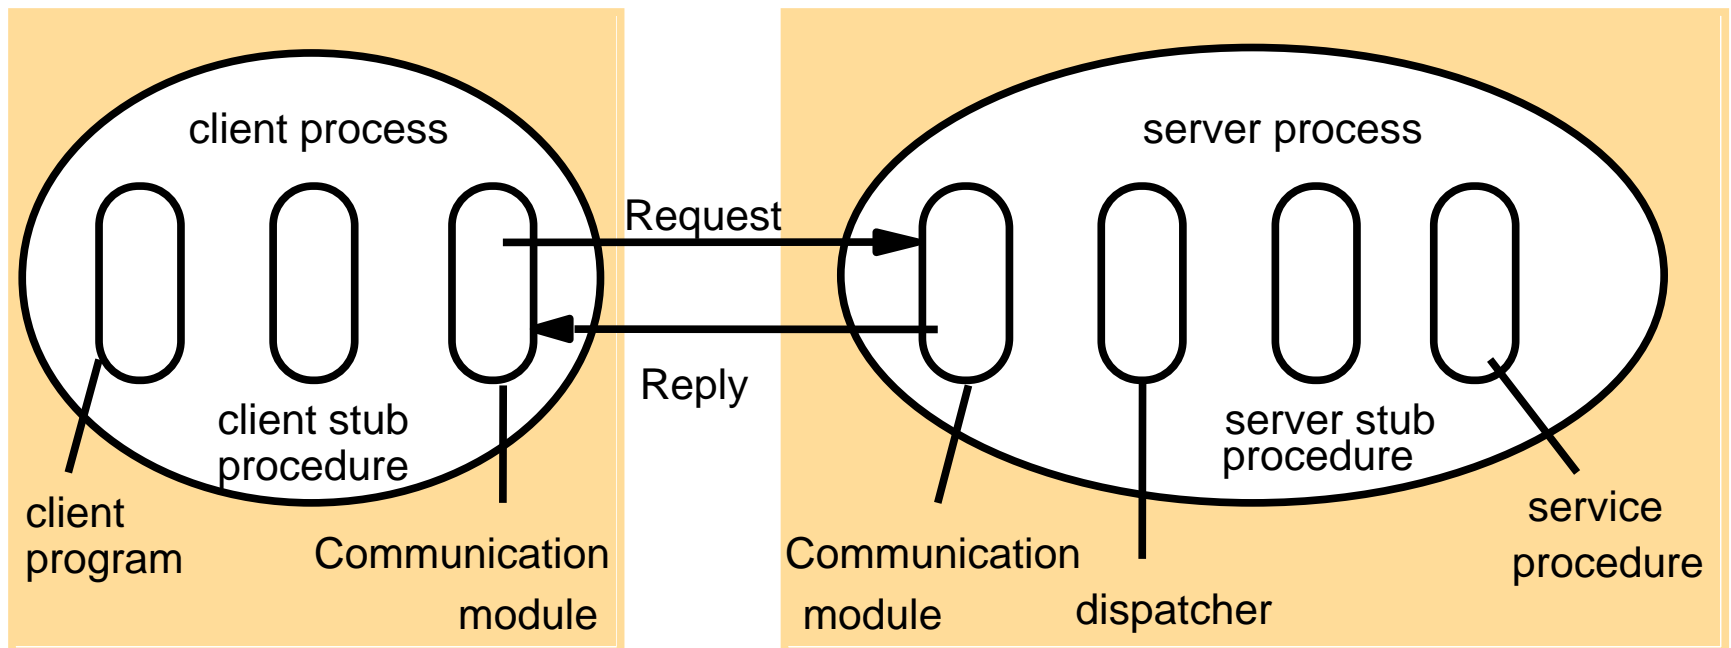
- RPC technology



# Remote procedure calls

## Basic principles

- RPC technology



# Remote procedure calls

- Application program calls client stub procedure
- Client stub procedure marshalls parameters of call and gives it to communication module in client
- Communication module in client transmits a message with the marshalled RPC
- Communication module in server receives message and gives it to dispatcher
- Dispatcher determines which procedure is called and calls correct servers stub procedure with marshalled data
- Server Stub procedure unmarshalls data and calls the server procedure



# Remote procedure calls

- Server procedure returns an answer to Server Stub procedure
- Server Stub procedure marshalls the answer and gives it to the communication module at server side
- Communication module at server side transmits the reply in a message
- Communication module at client side gives data to client stub procedure
- Who unmarshalls data and returns the answer to calling program





# Remote procedure calls

- Primary characteristics:
  - code in client and server independent of communication system
  - familiar paradigm: procedure call
    - no message preparation
    - synchronous interaction
    - semantics?
  - IDL: Interface definition language
    - independent of language used for client or server
    - base for generation of stubs



# Remote procedure calls

## Design issues

---

- Classes of RPC systems
- Interface Definition Language (IDL)
- Exception handling
- Semantics of RPC
- Transparency



# Remote procedure calls

## Design issues

---

- Classes of RPC systems
  - RPC integrated within a particular programming language
    - e.g. Argus (with CLU), Arjuna (C++)
  - RPC based on a special IDL
    - e.g. Sun RPC, ANSA RPC, OSF/DCE



# Remote procedure calls

## Design issues

---

- Interface Definition Language
    - Describes operation signatures
    - Interface compilers for
      - Generating client and server stubs
      - In different languages (e.g. C, Pascal...)
- ➔ Abstraction of heterogeneity



# Remote procedure calls

## Design issues

---

- Exception handling
  - failures cannot be hidden!
    - Network  $\Leftrightarrow$  server failure?
    - Client cannot distinguish
  - approaches to support failures:
    - Language specific
    - using return codes of functions
    - extension provided by IDL



# Remote procedure calls

## Design issues

---

- Semantics of RPC
  - Maybe
  - At-least-once (e.g. Sun RPC)
  - At-most-once (e.g. ANSA RPC)
  - Exactly-once:
    - difficult or impossible given failures



# Remote procedure calls

Delivery guarantees			RPC semantics
retry request	duplicate filtering	re-exec retrans reply	
no	not applic.	not.	Maybe
yes	no	re-exec	At-least-once
yes	yes	retrans reply	At-most-once



# Remote procedure calls

## Design issues

---

- Transparency
  - “*make RPC as much like local procedure calls as possible*”
  - RPC more vulnerable to failure
    - possibility of failure should not be hidden
  - calling instructions different
  - no shared memory between caller and callee

programming convenience

vs.

true transparency





# Remote procedure calls

## Implementation aspects

---

- Tasks for interface compiler: **generate**
  - client stub procedure
  - server stub procedure
  - marshalling and unmarshalling operations for each argument type
  - header for server procedure



# Remote procedure calls

## Implementation aspects

---

- Binding:  
linking client to server at execution time
- Binder interface

*Procedure Register (serviceName: String;  
serverPort: Port; version: integer);*

*Procedure Withdraw (...);*

*Procedure Lookup(serviceName: String;  
version:...): Port;*



# Remote procedure calls

## Implementation aspects

---

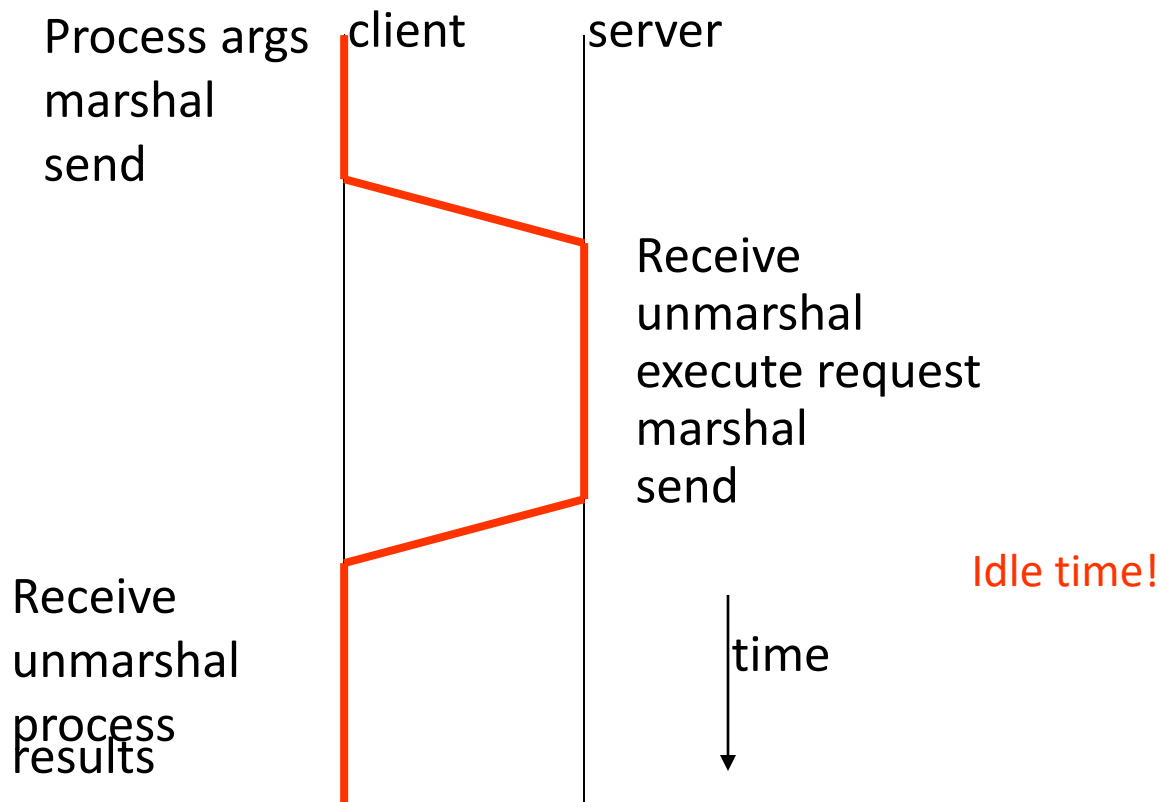
- Binding:
  - server will **register** service at binder
  - client will **lookup** the service
- Locating the binder?
  - well known host address
  - responsibility of OS
  - broadcast message by client



# Remote procedure calls

## Asynchronous RPC

- Problem: throughput RPC?




....

# Remote procedure calls


## Asynchronous RPC

- Timing:

Synchronous RPC



Asynchronous RPC



# Remote procedure calls

## Asynchronous RPC

---

- When to use?
  - Many request, small amount of information, limited processing:  
e.g. windows system
  - parallel requests to several servers
- Additional optimisations:
  - buffer request at client until ...
  - proceed without waiting when no reply is required



# Remote procedure calls

## Asynchronous RPC

---

- Extensions:
  - call streams:
    - mix of synchronous and asynchronous calls
    - message ordering preserved
    - connection oriented
    - connection breaks when semantics cannot be guaranteed



# Remote procedure calls

## Asynchronous RPC

---

- Extensions:
  - Promises:
    - allow clients to continue with (other work) and retrieve result of call later
    - created at the time of call
    - store results of call (object with same type)
    - operations:
      - get, await result
      - result available?
    - Alternative names: futures, tickets, continuations





# Remote procedure calls

---

- Conclusion:
  - familiar paradigm
  - has been basic primitive for distributed programming for many applications and systems...
  - limitations
    - Failures to be handled by clients (hard)
    - No transaction support
    - Only one-to-one communication





fourth edition

## DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris  
Jean Dollimore  
Tim Kindberg



KATHOLIEKE UNIVERSITEIT  
**LEUVEN**

# Distributed Systems: Direct Communication – Part I Questions?

