

Distributed Systems: Introduction to cloud computing

Bert Lagaisse, Stefan Walraven, Wouter Joosen, Tom Van Cutsem

DistriNet, KU Leuven

October, 2023

Overview

- Cloud computing ?
 - = business model + technical architectures
- Case studies:
 - Google AppEngine
 - Microsoft Azure,

Definitions of cloud computing?

Cloud computing: essentials

- Outsourcing, pay per use model
- For on-demand, web-based access
- to shared pool of computing resources
 - Virtual machines: Storage, cpu, network
 - Applications
- 3 types of services in cloud computing model
 - Infrastructure as a service (IaaS):
 - virtual machine , storage, network
 - Platform as a service (PaaS):
 - middleware and web-hosting platform (php, .net, java)
 - Software as a service (SaaS):
 - Zero-install, online applications (CRM, Google Mail, Google Apps)



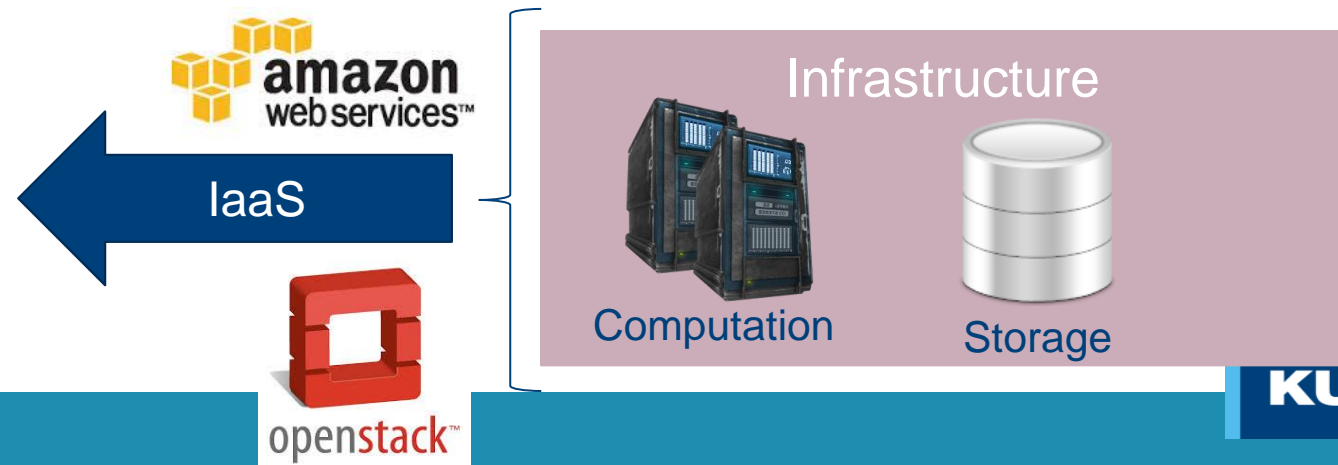
IaaS: private or public ?

- IaaS: infrastructure as a service
 - Virtual machines on shared hardware !
 - With operating system of choice
 - Virtual networks between machines
 - Storage (Blob storage, structured storage)
- Public cloud: Amazon, MS Azure
 - Accessible by any one from anywhere
 - (When you have a credit card. 😊)
 - Many data centers around the world
 - Choose where you want your resources in advance.
- Private cloud: OpenStack, Xen Cloud platform
 - private data centers of companies (in Flanders: integrators such as Cegeka, Telecom operators)
 - Self-service creation of virtual machines etc.

Infrastructure as a Service



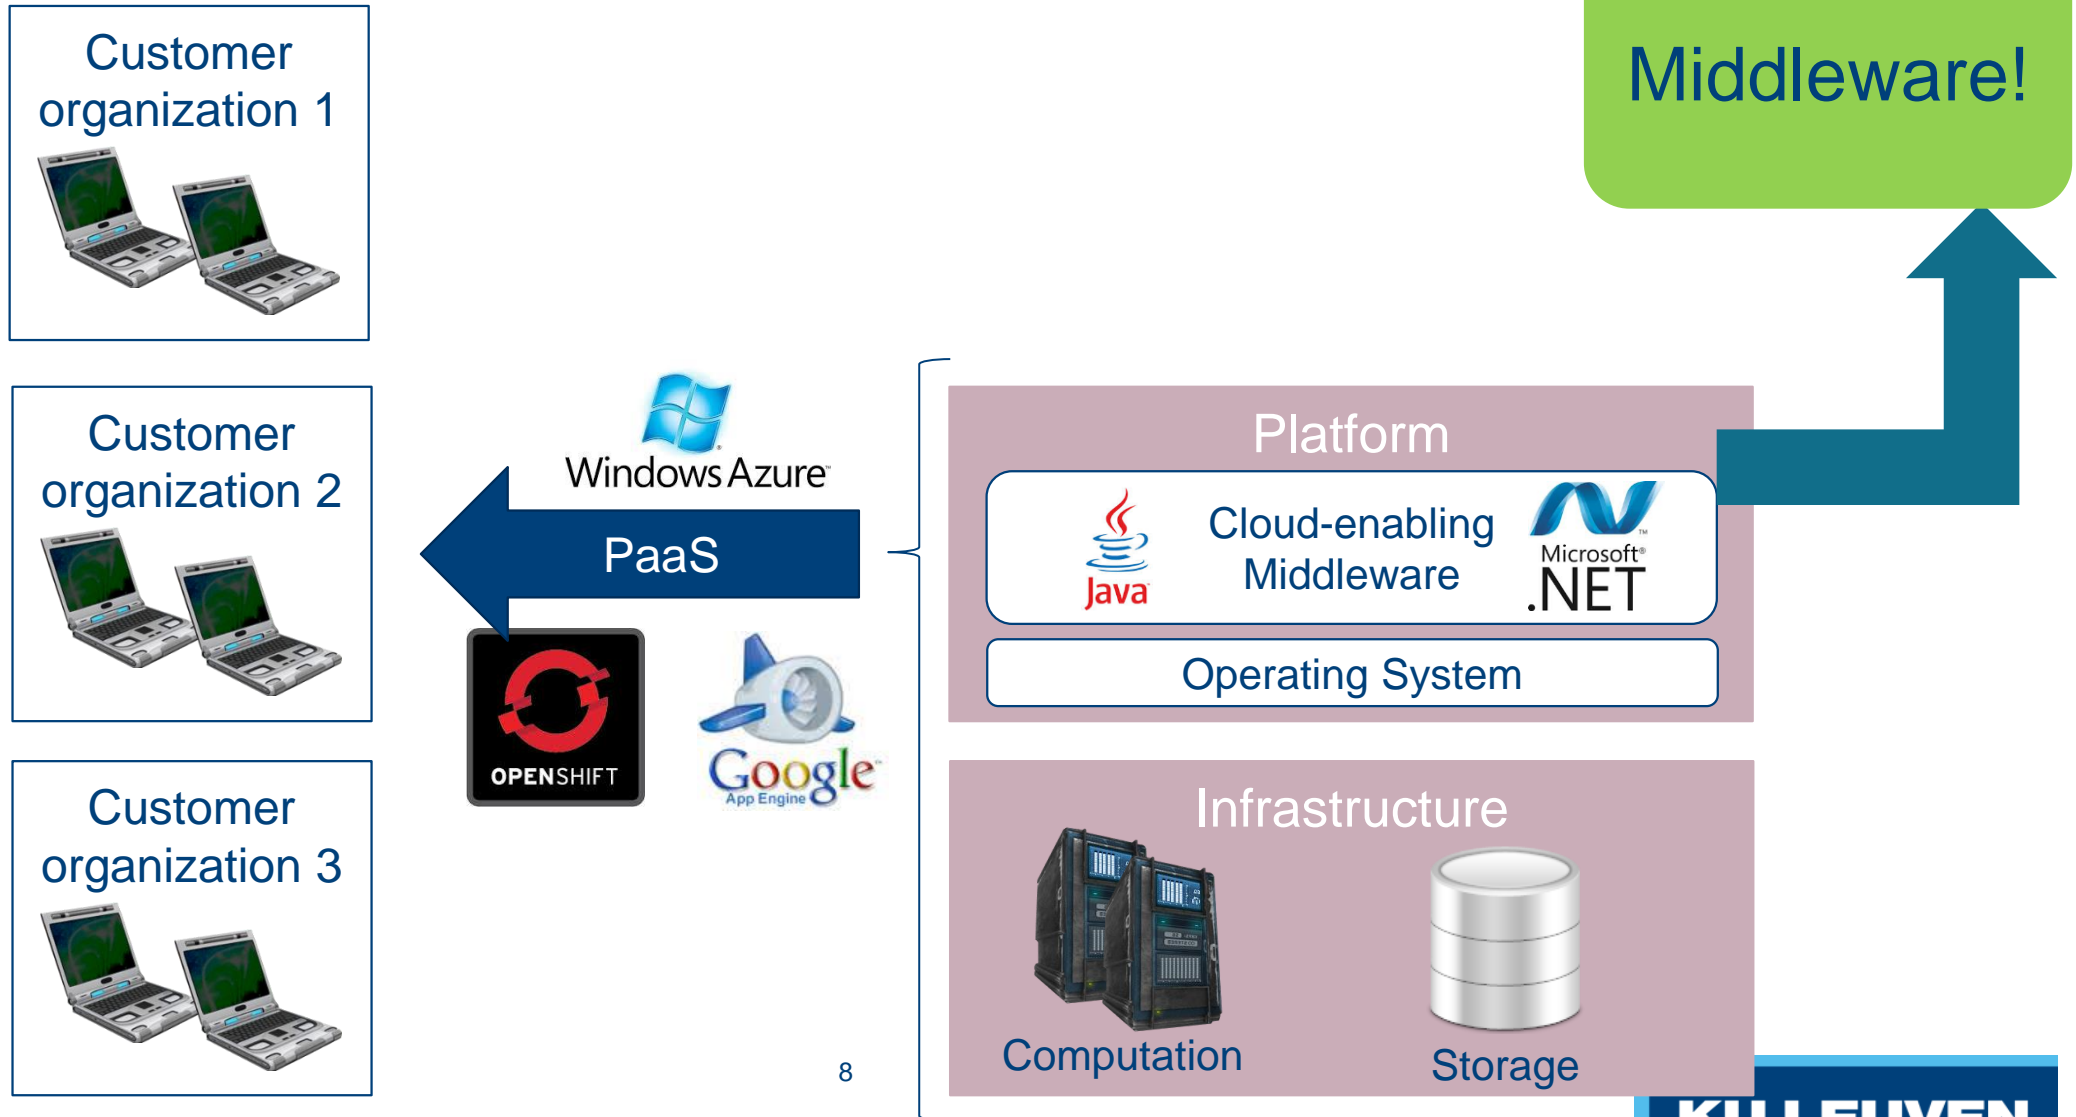
Virtual Systems!



PaaS: focused on scaling-out and resource sharing

- Middleware as a service
 - scalable web server, web container, app server, background workers with triggers,...
 - without the burden of setting up the whole VM underneath
- Google AppEngine, MS Azure, OpenShift.
- Software architecture properties: **scalability** first !
 - Different way of thinking about software entities and their interactions
 - From: object-based, synchronous method invocations
 - To: Task-based, background workers, with asynchronous message passing (not locking resources)
 - Different way of thinking about persistence
 - Due to scalability requirements: must scale-out
 - Need for replicated data storage with simple access (e.g. key based retrieval of simple entities, blobs)
- **Resource sharing** has impact on performance and failure isolation: multiple applications run on the same web server, VM.

Platform as a Service



Next-gen PaaS:

IaaS → VM-based PaaS → serverless model

- Initial cloud hype & popularity
 - IaaS !
 - Easily create VMs anywhere, when you need them
- The systems need to be managed☹ - operational challenges remain, e.g.
 - Installation and maintenance of software: Web server, App server, web frameworks, etc.
 - Configuration of the network between the VMs
 - → Devops = automate the deployment of VMs and software
 - Requires relatively complex scripts for automation, risks to be error prone.
 - Example: robust management of Kubernetes Clusters
- “Serverless” model:
 - Don’t manage the whole (generic) server software stack
 - Focus on application logic, and deploy application software “as smoothly as possible”.
 - Let cloud provider handle automatic scaling of infrastructure
 - Let cloud provider handle the management of the software stack
 - Examples: AWS lambda, (Google AppEngine, Azure Web Apps)

CACM September 2023

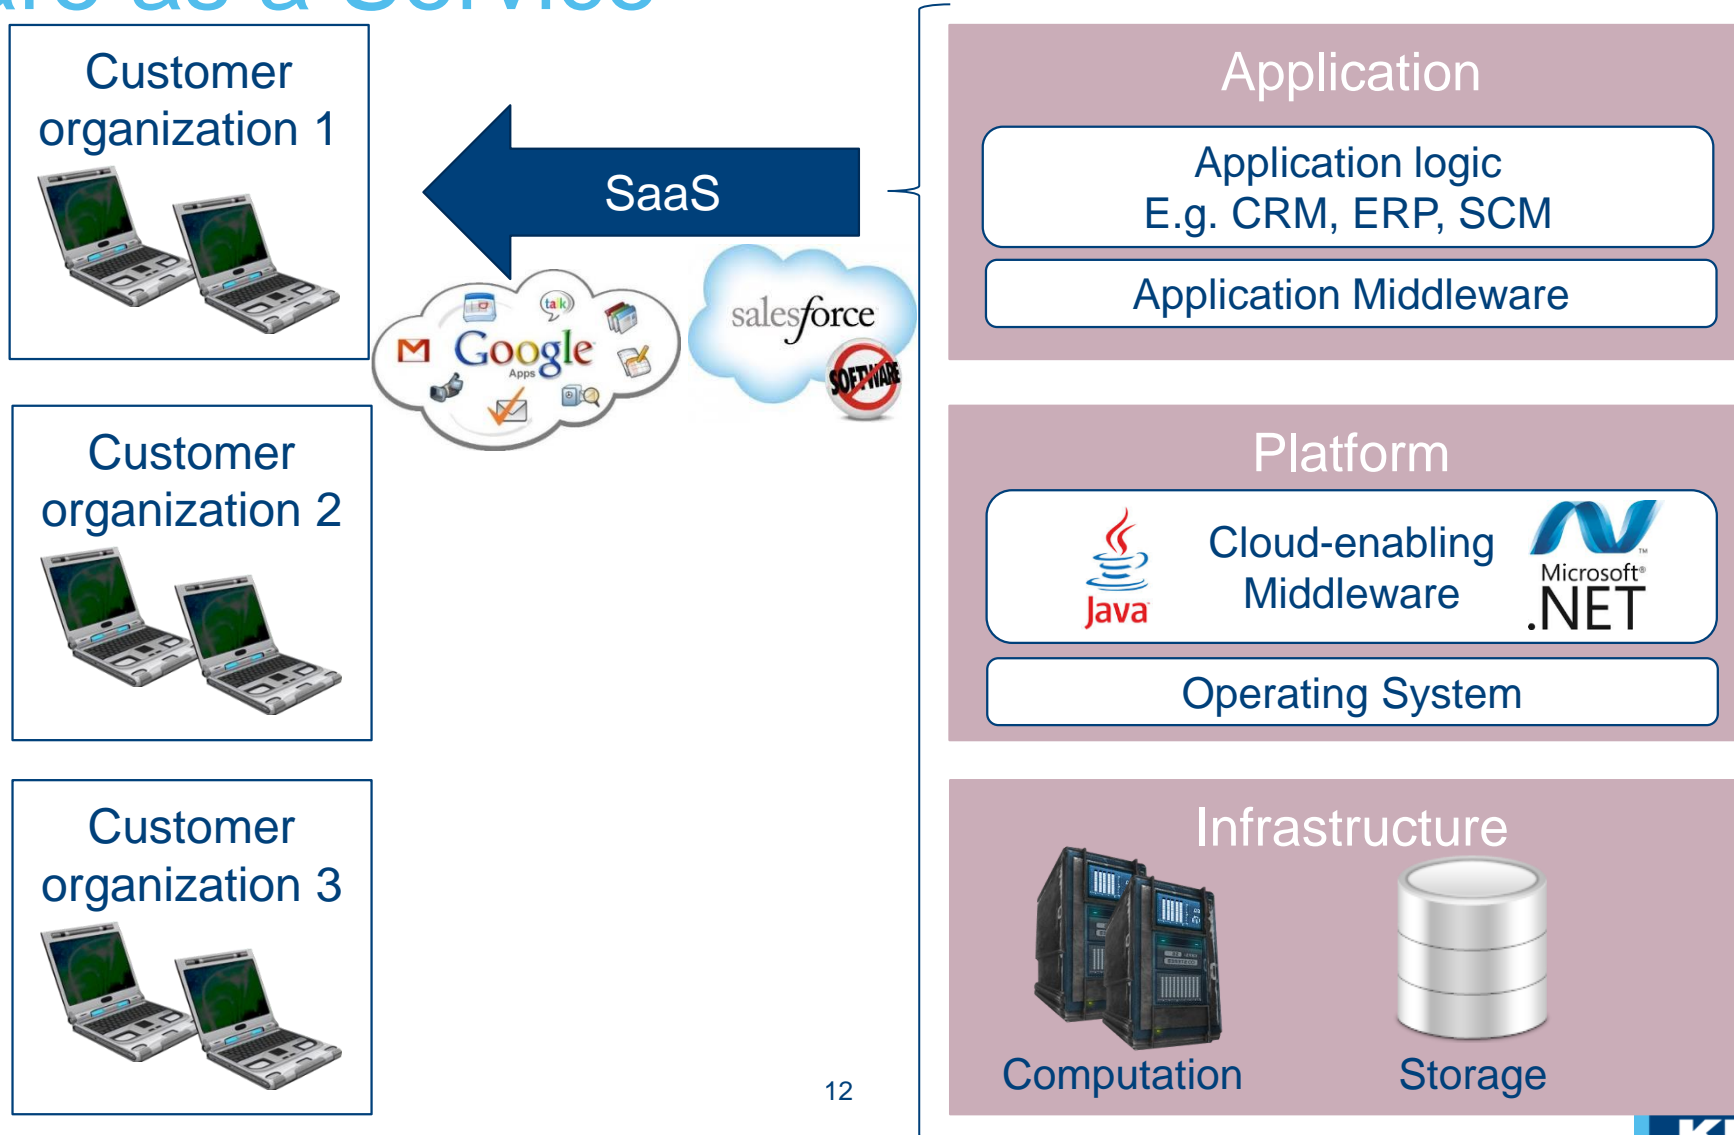
- “*Serverless Computing: What It Is, and What It Is Not?*”, pp 80-92
- → allow to develop deploy and run applications
- Key insights -selected
 - Serverless computing means full automation and fine-grained utilization-based billing
 - Serverless computing supports diverse applications, (e.g.) from enterprise automation to scientific computing
- Note: projected market value of serverless computing \$36.8 billion by the end of 2028 — Ref. 49

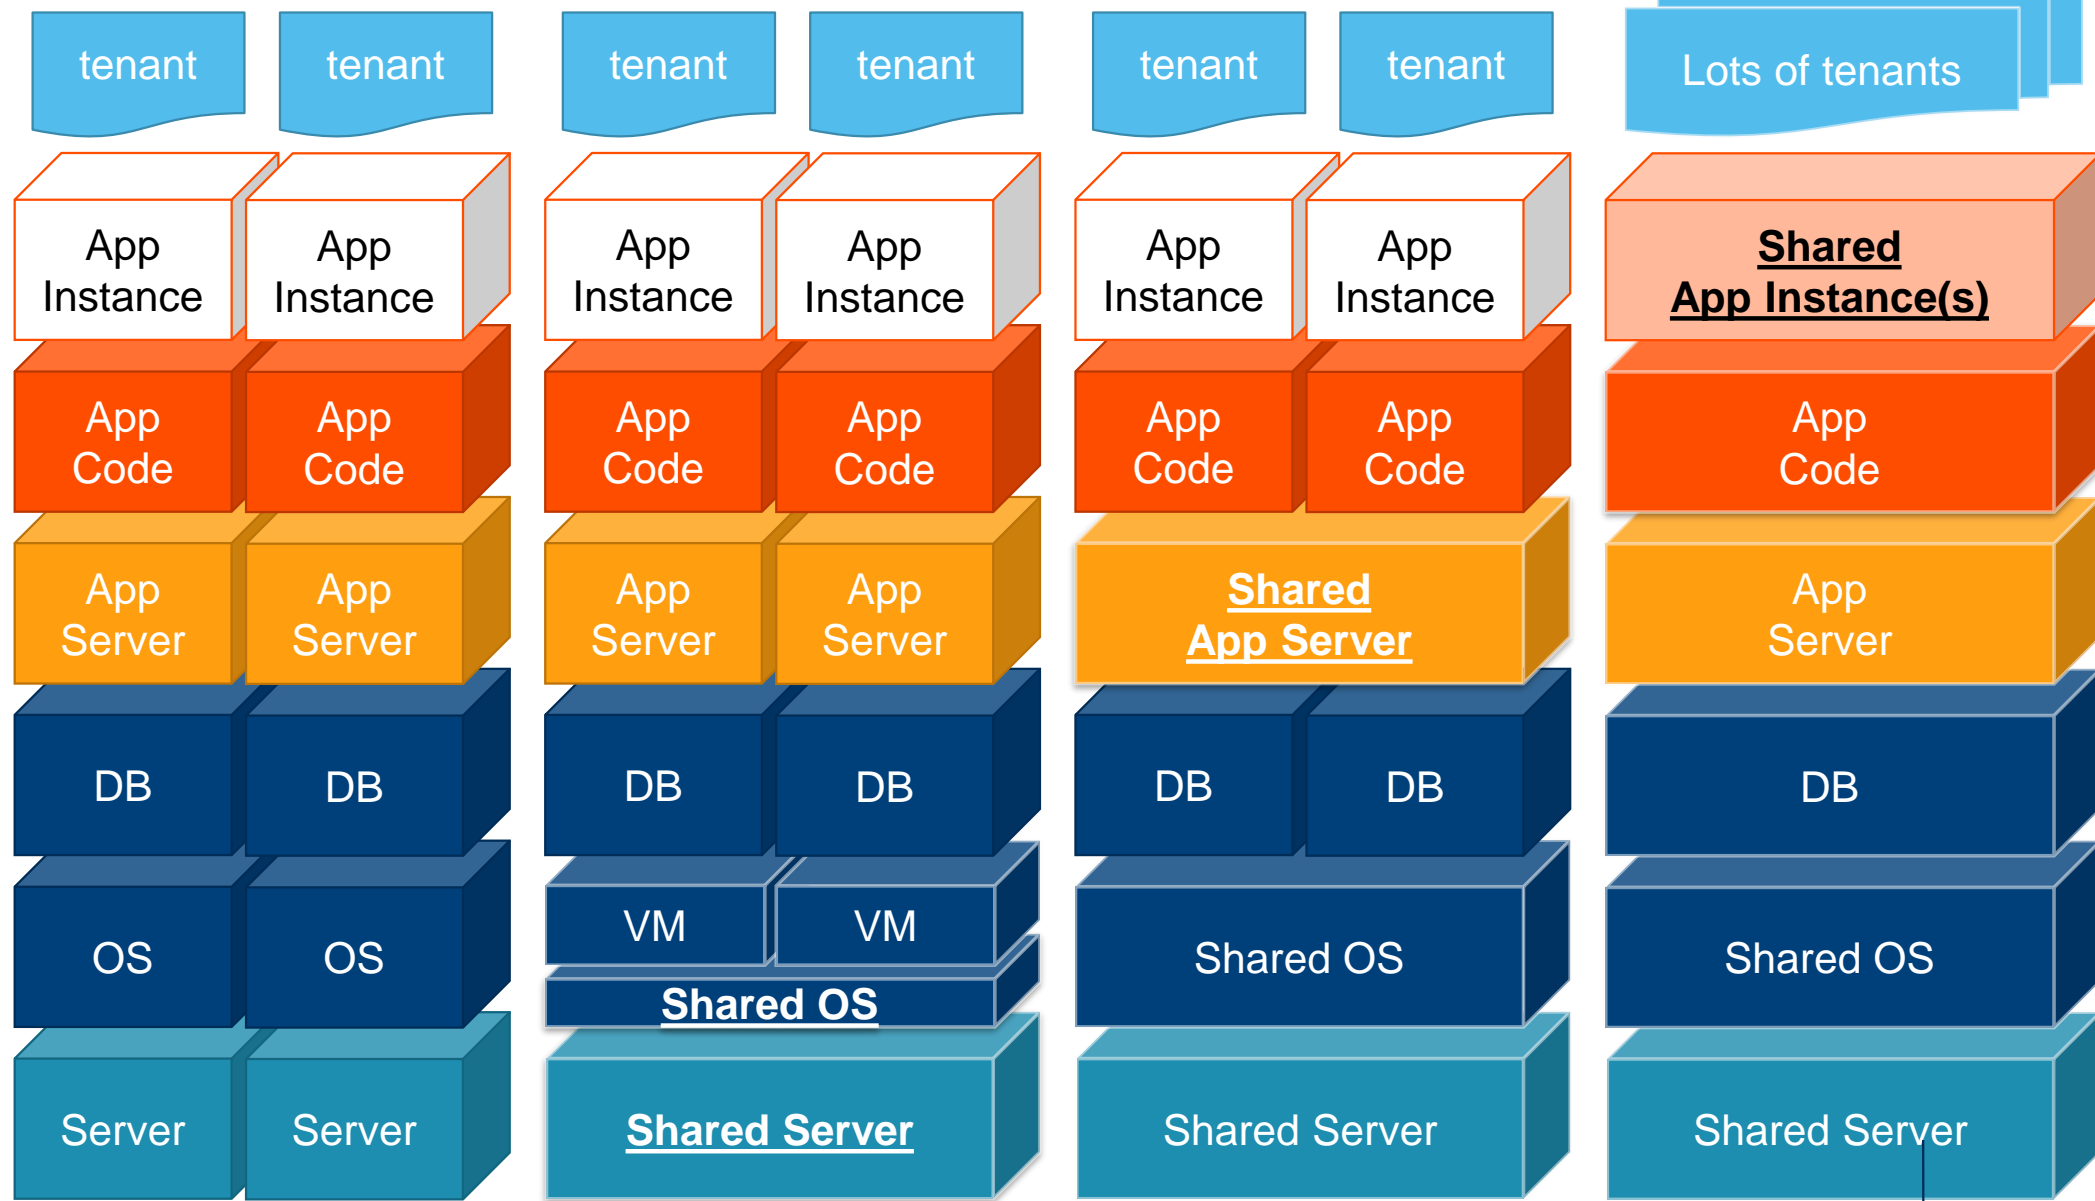
Is Serverless computing in 2020 +

... what cloud computing has been in 2010+?

Is serverless computing indeed the next concept/paradigm?

Software as a Service





Shared Nothing

Shared OS

Shared App Server

Shared Everything

SaaS: Application-level logic to separate tenants:
CHALLENGES: Customization, performance, lots of data, security logic,...



Architectural paradigm shift

- Paradigm shift in business model
 - As discussed before: towards pay per use
- But also at the architectural level: next distributed architecture
 - Scientific and enterprise computing have evolved to a common architecture ... ?

Mainframe	1970s	Centralized bulk data processing, multi-user access on terminals (keyboard-screen)
Client-server	1980s	Fat clients on workstations, centralized database server
Web	1990s	Multi-tier: web tier, application servers, database .com your business
SOA	2000s	Expose automated business process as a web service B2B communication
Cloud	2010+	Centralized provision of dynamically scalable, multi-tenant software as a service over the Internet Builds further upon SOA (→ REST)

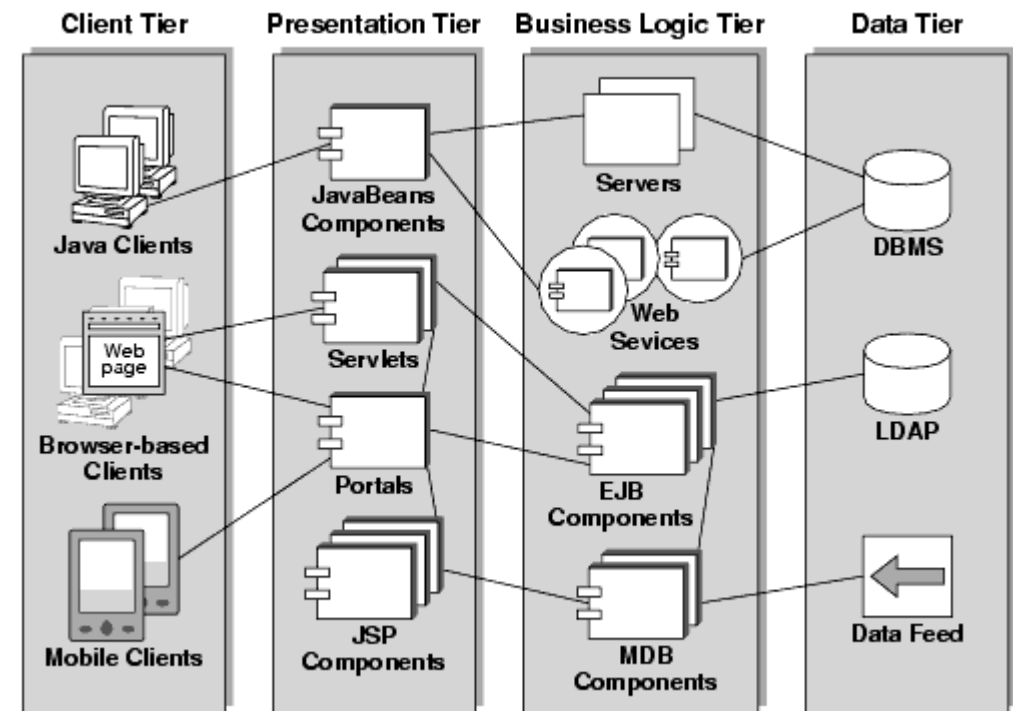
Two Tier architecture

- Two tier architecture
 - Client:
 - Fat client on desktop
 - can contain application logic
 - Server:
 - Data base server with app logic
 - E.g. stored procedures & transactions
- 3 tier architecture
 - user interface (presentation), functional
 - process logic (“business rules”),
 - computer data storage and data access



Multi-tier architecture

- Client tier: client machines
- Presentation tier (server-side)
 - displays information related services
 - communicates with other tiers
 - puts out the results to the browser/client tier and all other tiers in the network.
- Application tier (business logic, logic tier, or middle tier)
 - The logical tier is pulled out from the presentation tier and,
 - controls an application's functionality by performing detailed processing.
- Data tier
 - includes the data persistence mechanisms (database servers, file shares, etc.)
 - Application Programming Interface (API) to the application tier that exposes methods of managing the stored data



Back to cloud

5 core characteristics of cloud

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - Higher degree of distribution
 - Multi-tenancy
 - Elasticity
 - Delivery as a service
 - Self-service

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - **Higher degree of distribution**
 - Execution environment is **distributed at a large scale**
 - Across different data centers (using commodity hardware)
 - At different geographical locations (geo-distribution)
 - to achieve **high availability and high performance** (i.e. high throughput and low latency) via replication
 - accessible from any device and location through standard mechanisms and APIs
 - Multi-tenancy
 - Elasticity
 - Delivery as a service
 - Self-service

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - Higher degree of distribution
 - **Multi-tenancy**
 - Dynamically assigning the available **shared resources** to multiple customer organizations (= tenants)
 - One of the key enablers to realize **economies of scale**:
 - operational costs are significantly reduced by multiplexing (shared) resources among multiple tenants
 - simplifies administration, provisioning and maintenance
 - Elasticity
 - Delivery as a service
 - Self-service

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - Higher degree of distribution
 - Multi-tenancy
 - **Elasticity**
 - Dynamically provisioning and releasing resources on demand
 - by **scaling horizontally**: adding more servers to the system and distributing the load (\Leftrightarrow making servers more powerful)
 - **Flexibility** to scale out rapidly (and automatically) at a fine-grained level based on the current workload
 - without having to constantly provision sufficient resources for peak demand and thus avoiding over- or underprovisioning
 - Delivery as a service
 - Self-service

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - Higher degree of distribution
 - Multi-tenancy
 - Elasticity
 - **Delivery as a service**
 - Management of resources and applications **outsourced** to the cloud provider
 - No upfront commitments and investments in infrastructure
 - Only pay for resources and applications that are used (**pay per use**)
 - Cloud provider has to monitor usage per customer
 - Self-service

Core characteristics

- 5 specific characteristics that differ from what is offered by previous paradigms:
 - Higher degree of distribution
 - Multi-tenancy
 - Elasticity
 - Delivery as a service
 - **Self-service**
 - To preserve **scalability** with an increasing # customers
 - To **rapidly respond** to changing service demand and customer requirements
 - **Automation** (i.e. no human interaction with cloud provider)
 - **Empower customers** to manage cloud services (via APIs, configuration interfaces, tools), e.g. to scale out

Some of the key concerns – still (think *B2B!*)

- Lock-In
 - Hybrid Cloud and Multi-Cloud
- Regulatory Compliance (and Customer needs)
- Security
- SLA/SLO monitoring
 - Service Level Agreements – Service Level Objectives
- → Resilience Management

Going Forward...

- Case studies from large cloud providers
 - Google AppEngine
 - (Considering Azure)