

# Distributed Systems File Systems

Wouter Joosen

DistriNet, KULeuven

October 24 & 31, 2023

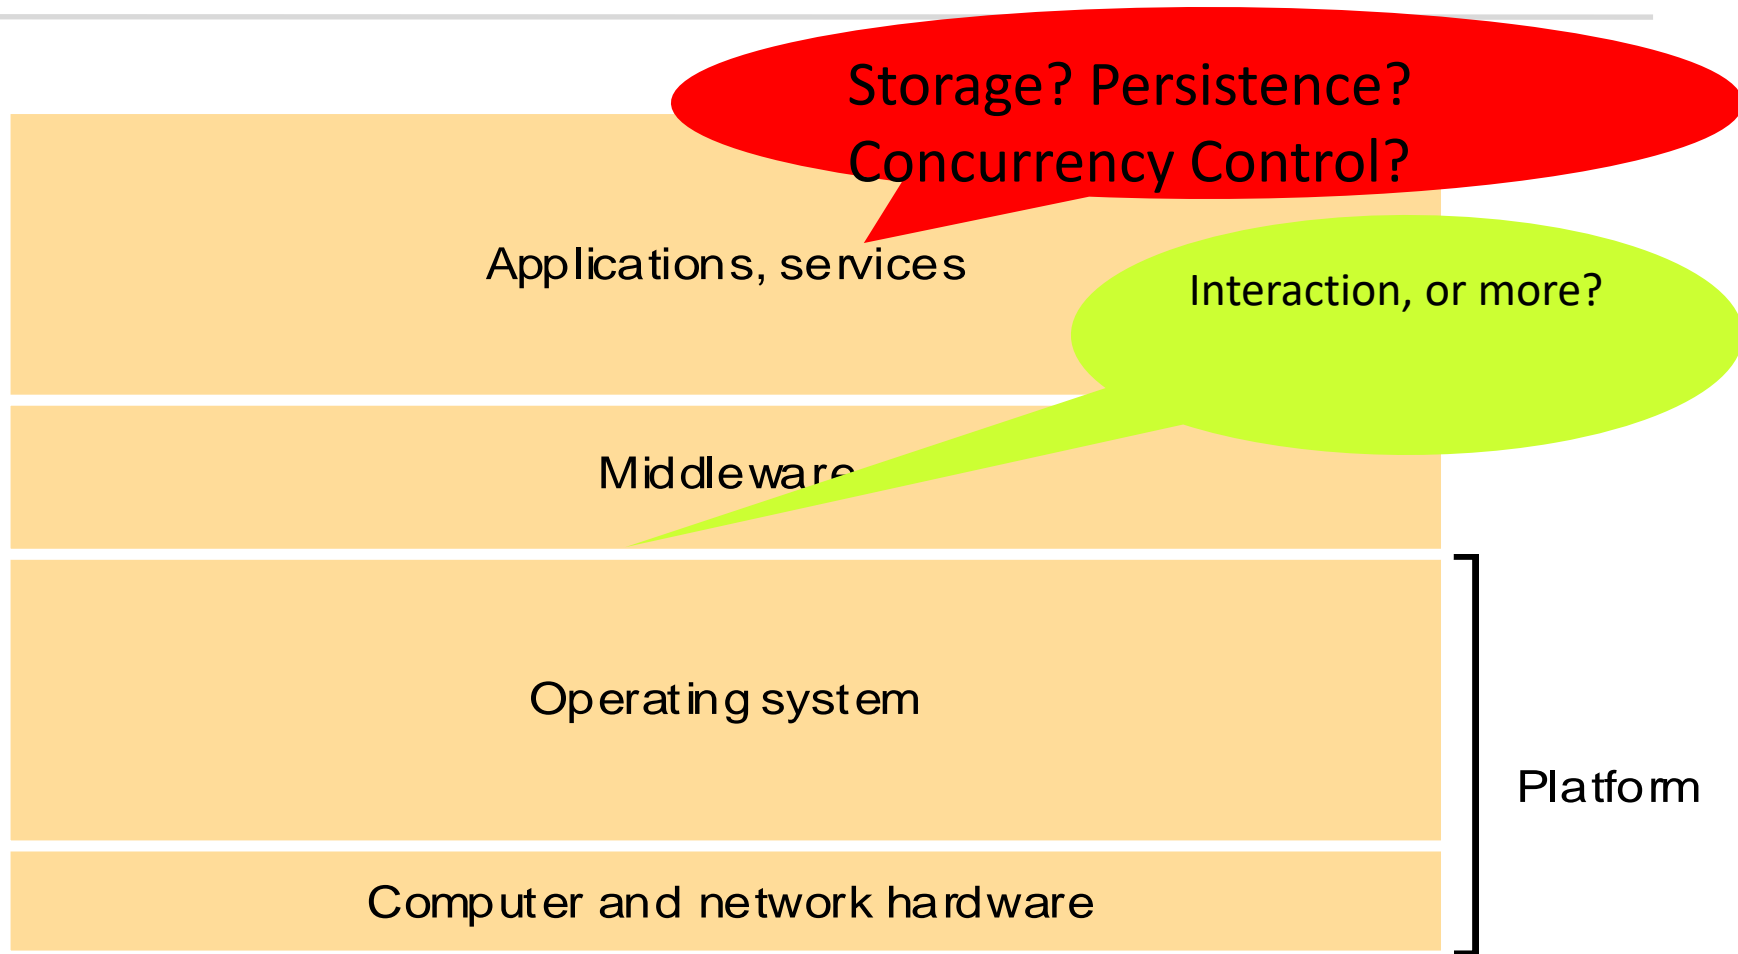


# Introduction

- Objectives of this study
  - *Study existing subsystems: a combination of modern real world, and pedagogical (slightly older)*
    - Services offered
    - Architecture & design trade-offs
    - Limitations
  - Learn about **non-functional requirements**: initial focus is on performance, availability and fault tolerance, data consistency
    - Techniques used
    - Results
    - Limitations
  - Non function requirements: Security & Privacy – ***new & intro in 2023-2024!***



# Where to apply know how?



# Build up experience:

---

- In using middleware (building applications)
  - (system perspective)
- In customizing/adapting middleware
  - (development perspective)
- Eventually in building middleware



# Existing services: *from the textbook*

---

- Distributed file systems
- Name services (see DNS in computer networks!)
- ...





fourth edition

**DISTRIBUTED SYSTEMS**  
CONCEPTS AND DESIGN

George Coulouris  
Jean Dollimore  
Tim Kindberg



# Distributed Systems:

## File Systems

Main focus of this case study—  
caching & performance

# Distributed file systems

---

- Overview
  - File service architecture (quick intro)
  - case study: NFS
  - case study AFS
  - comparison NFS <> AFS



# Distributed file systems

## File service architecture

---

- Definitions:
  - file
  - directory
  - file system
- cf. Operating systems





# Distributed file systems

## File service architecture

---

- Requirements:
  - access transparency
  - location transparency
  - failure transparency
  - performance transparency
  - hardware and operating system heterogeneity
  - Scalability
  - concurrent updates
- Historical case studies (pioneering cases 80'- 90's):
  - Low interest in concurrency,
  - String focus on remote access



# Distributed file systems

## File service architecture

---

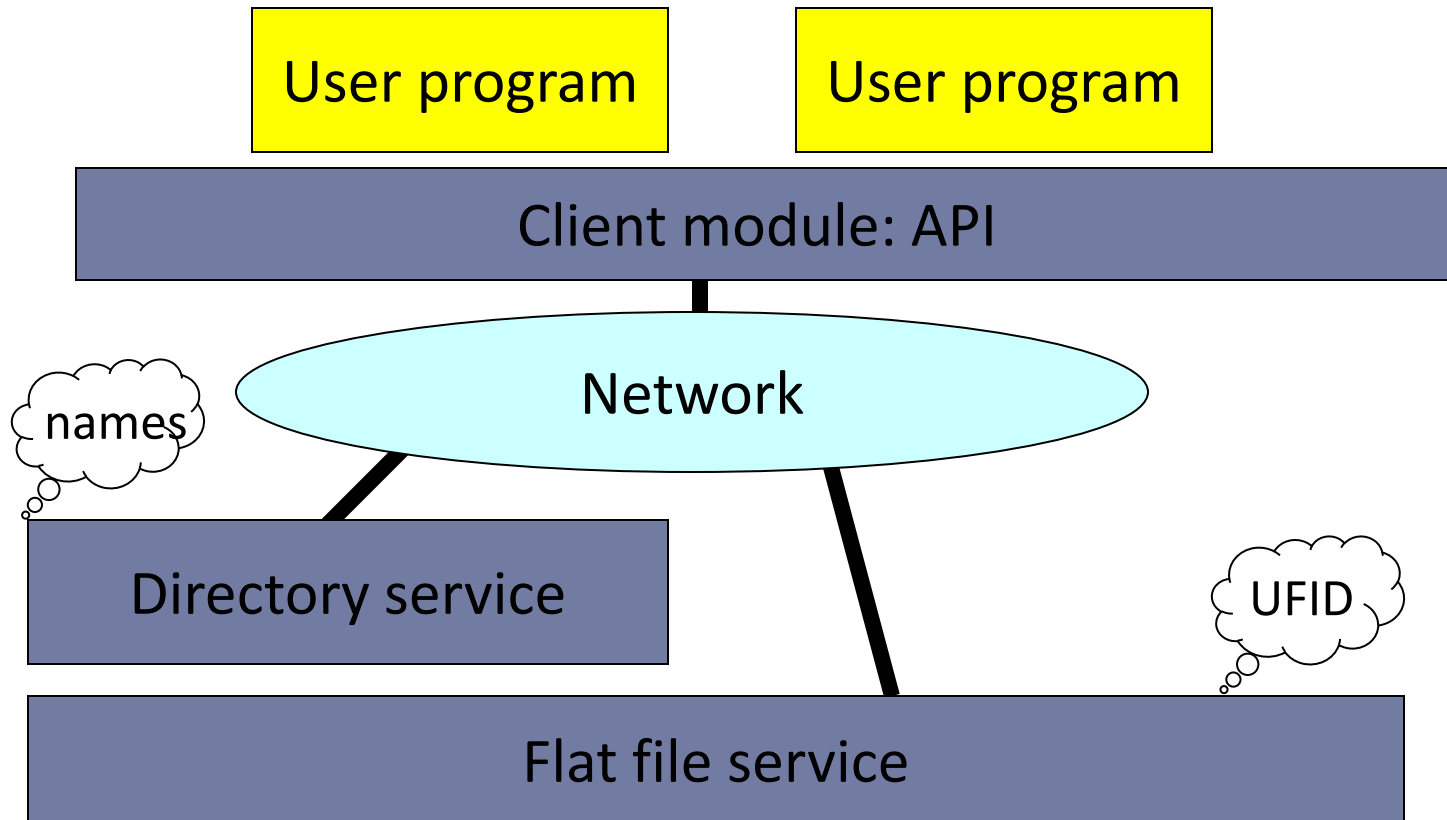
- Requirements (*cont.*) :
  - scalable to a very large number of nodes
    - replication transparency
    - migration transparency
  - future:
    - tolerance to network partitioning



# Distributed file systems

## File service architecture

- File service components



# Distributed file systems

## File service architecture

---

- Flat file service
  - file = data + attributes
  - data = sequence of items
  - operations: simple read & write
  - attributes: in flat file & directory service
  - UFID: unique file identifier



# Distributed file systems

## File service architecture

---

- Flat file service: operations
  - *Read (FileID, Position, n) → Data*
  - *Write (FileID, Position, Data)*
  - *Create () → FileID*
  - *Delete (FileID)*
  - *GetAttributes (FileID) → Attr*
  - *SetAttributes (FileID, Attr)*



# Distributed file systems

## File service architecture

---

- Flat file service: in light of *fault tolerance (easy recovery)*
  - straightforward for simple servers
  - idempotent operations
  - stateless servers



# Distributed file systems

## File service architecture

---

- Directory service
  - translate file name in UFID
  - substitute for open
  - responsible for access control



# Distributed file systems

## File service architecture

---

- Directory service: operations
  - *Lookup (Dir, Name, AccessMode, UserID)*  
→ UFID
  - *AddName (Dir, Name, FileID, UserID)*
  - *UnName (Dir, Name)*
  - *ReName (Dir, OldName, NewName)*
  - *GetNames (Dir, Pattern)* → *list-of-names*





# Distributed file systems

## File service architecture

---

- Implementation techniques:
  - known techniques from OS experience
  - remain important
  - distributed file service: comparable in
    - performance
    - reliability



# Distributed file systems

## File service architecture

---

- Implementation techniques: overview (in line with text book)
  - file groups
  - space leaks
  - capabilities and access control (*no priority at this point*)
  - access modes
  - file representation
  - file location
  - group location
  - file addressing
  - Caching: *main focus of this case study*



# Distributed file systems

## File service architecture

---

- Implementation techniques: **file groups**
  - (similar: file system, partition)
  - = collection of files mounted on a server computer
  - unit of distribution over servers
  - transparent migration of file groups
  - once created file is locked in file group
  - UFID = file group identifier + file identifier
- Note: operational management and administrative aspects often ignored, underestimated in importance....



# Distributed file systems

## File service architecture

---

- Implementation techniques: **space leaks**
  - 2 steps for creating a file
    - create (empty) file and get new UFID
    - enter name + UFID in directory
  - failure after step 1:
    - file exists in file server
    - unreachable: UFID not in any directory
  - ➔ lost space on disk
  - detection requires co-operation between
    - file server
    - directory server



# Distributed file systems

## File service architecture

---

- Implementation techniques: **capabilities**
  - = digital key: access to resource granted on presentation of the capability
    - request to directory server: file name + user id + mode of access
      - ➔ UFID including permitted access modes
    - construction of UFID
    - unique
    - encode access
    - unforgeable



# Distributed file systems

## File service architecture

- Implementation techniques: **capabilities**



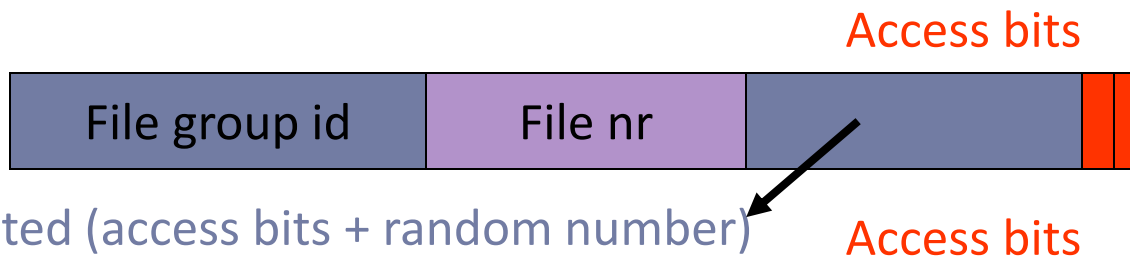
Reuse?



Access check?



Forgeable?



Encrypted (access bits + random number)

Access bits

Access bits

# Distributed file systems

## File service architecture

---

- Implementation techniques: **file location**
  - from UFID → location of file server
  - use of replicated group location database  
**file group id, PortId**
  - why replication?
  - why location not encoded in UFID?



# Distributed file systems

## File service architecture

---

- Implementation techniques: **caching**
  - server cache: reduce delay for disk I/O
    - selecting blocks for release
    - coherence:
      - dirty flags
      - write-through caching
  - client cache: reduce network delay
    - always use write-through
    - synchronisation problems with multiple caches





# Distributed file systems

---

- Overview
  - File service model
  - case study: Network File System -- NFS
  - case study: AFS
  - comparison NFS <> AFS



# Distributed file systems

## NFS

- Background and aims
  - first file service product
  - emulate UNIX file system interface
  - de facto standard
    - key interfaces published in public domain
    - source code available for reference implementation
  - supports diskless workstations
    - not important anymore



# Distributed file systems

## NFS

- Design characteristics
  - client and server modules can be in any node  
Large installations include a few servers
  - Clients:
    - On Unix: emulation of standard UNIX file system
    - for MS/DOS, Windows, Apple, ...
  - Integrated file and directory service
  - Integration of remote file systems in a local one: **mount** →  
**remote mount**



# Distributed file systems

## NFS: configuration

---


- Unix mount system call
  - each disk partition contains hierarchical FS
  - how integrate?
    - Name partitions  
a:/usr/students/john
    - glue partitions together
      - invisible for user
      - partitions remain useful for system managers



# Distributed file systems

## NFS: configuration

- Unix mount system call




Directory staff  $\equiv$  root of c: `/usr/staff/ann/network`

# Distributed file systems

## NFS: configuration

- Remote mount



Directory staff  $\equiv$  users: `/usr/staff/ann/...`

# Distributed file systems

## NFS: configuration

---

- Mount service on server
  - enables clients to integrate (part of) remote file system in the local name space
  - exported file systems in `/etc/exports` + access list (= hosts permitted to mount; secure?)
- On client side
  - file systems to mount enumerated in `/etc/rc`
  - typically mount at start up time



# Distributed file systems

## NFS: configuration

---

- Mounting semantics
  - hard
    - client waits until request for a remote file succeeds
    - eventually forever
  - soft
    - failure returned if request does not succeed after n retries
    - breaks Unix failure semantics





# Distributed file systems

## NFS: configuration

---

- Automounter
  - principle:
    - *empty* mount points on clients
    - mount on first request for remote file
  - acts as a server for a local client
    - gets references to empty mount points
    - maps mount points to remote file systems
    - referenced file system mounted on mount point via a symbolic link, to avoid redundant requests to automounter



# Distributed file systems

## NFS: implementation

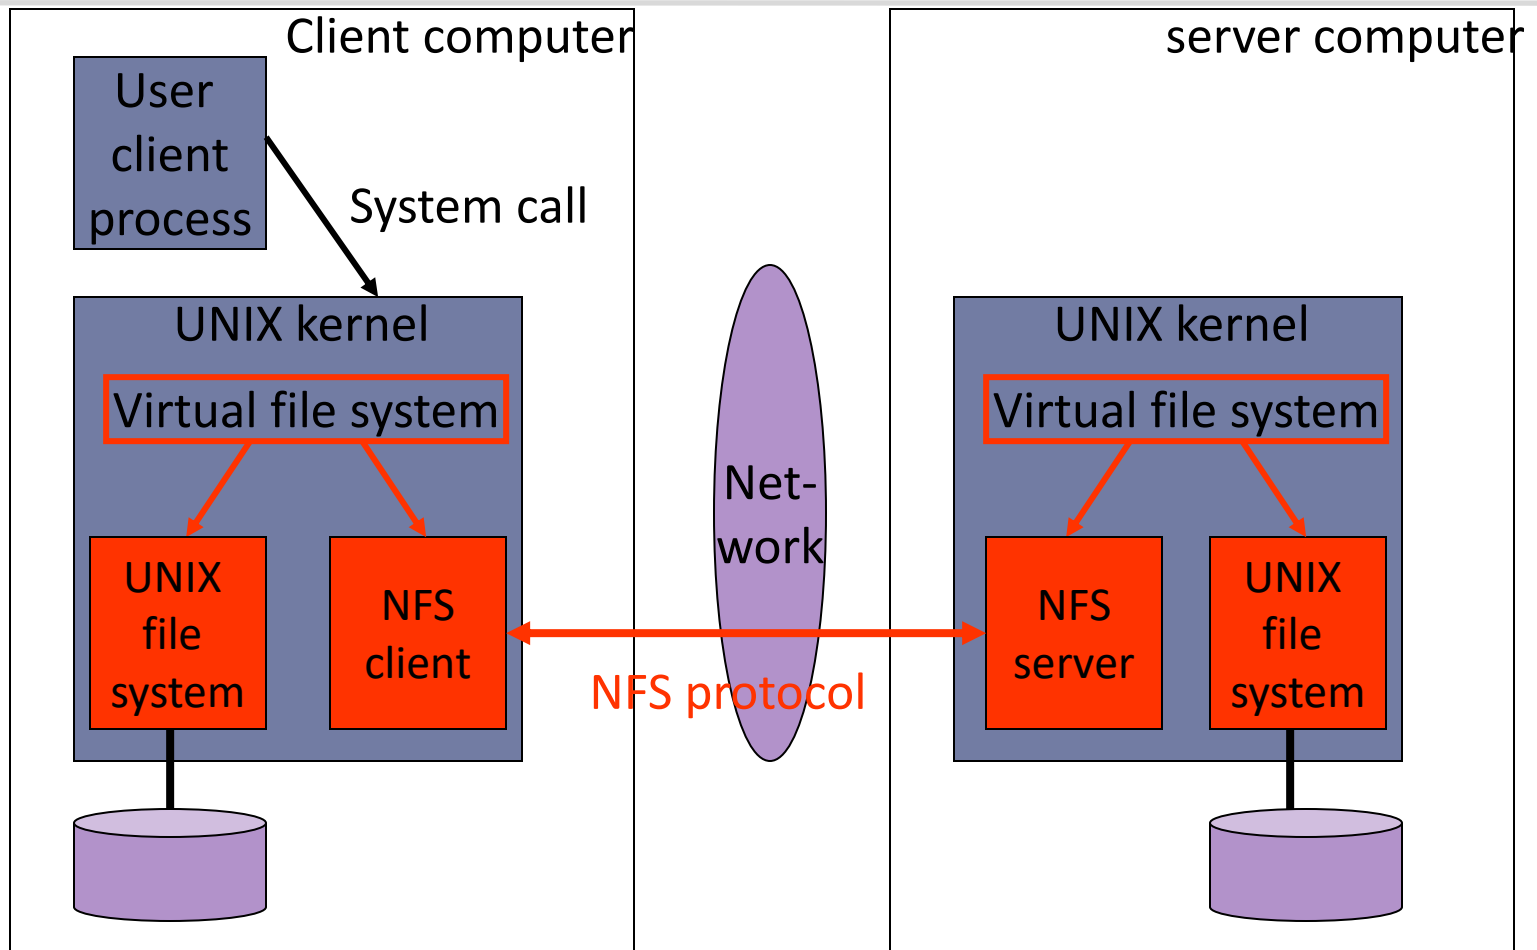
---

- In UNIX: client and server modules implemented in kernel
- virtual file system:
  - internal key interface, based on file handles for remote files



# Distributed file systems

## NFS: implementation



# Distributed file systems

## NFS: implementation

---

- Virtual file system
  - Added to UNIX kernel to make distinction
    - Local files
    - Remote files
  - File handles: file ID in NFS
    - Base: inode number
      - File ID in partition on UNIX system
    - Extended with:
      - File system identifier
      - inode generation number (to enable reuse)



# Distributed file systems

## NFS: implementation

---

- Client integration:
  - NFS client module integrated in kernel
    - offers standard UNIX interface
    - no client recompilation/reloading
    - single client module for all user level processes
    - encryption in kernel
- server integration
  - only for performance reasons
  - user level = 80% of kernel level version



# Distributed file systems

## NFS: implementation

---

- Directory service
  - name resolution co-ordinated in client
  - step-by-step process for multi-part file names
  - mapping tables in server: high overhead reduced by caching
- Access control and authentication
  - based on UNIX user ID and group ID
  - included and checked for every NFS request
  - secure NFS 4.0 thanks to use of DES encryption



# Distributed file systems

## NFS: implementation

---

- Caching
  - Unix caching
    - based on disk blocks
    - delayed write
    - read ahead
    - periodically sync to flush buffers in cache
  - Caching in NFS
    - Server caching
    - Client caching



# Distributed file systems

## NFS: implementation

---

- Server caching in NFS
  - based on standard UNIX caching: 2 modes
  - write-through (instead of delayed write)
    - failure semantics ↗
    - performance ↘
  - delayed write
    - Data stored in cache, till commit operation is received
      - Close on client → commit operation on server
    - Failure semantics?
    - Performance ↗





# Distributed file systems

## NFS: implementation

---

- Client caching
  - cached are results of
    - *read, write, getattr, lookup, readdir*
  - problem: multiple copies of same data at different NFS clients
  - NFS clients use read-ahead and delayed write



# Distributed file systems

## NFS: implementation

---

- Client caching (*cont.*)
  - handling of writes
    - block of file is fetched and updated
    - changed block is marked as dirty
    - dirty pages of files are flushed to server asynchronously
      - on close of file
      - sync operation on client
      - by bio-daemon (when block is filled)
    - dirty pages of directories are flushed to server
      - by bio-daemon without further delay



# Distributed file systems

## NFS: implementation

---

- Client caching (*cont.*)
  - consistency checks
    - based on time-stamps indicating last modification of file on server
    - validation checks
      - when file is opened
      - when a new block is fetched
      - assumed to remain valid for a fixed time (3 sec for file, 30 sec for directory)
      - next operation causes another check
  - costly procedure



# Distributed file systems

## NFS: implementation

---

- Caching
  - Cached entry is valid
$$(T - T_c) < t \text{ or } (T_{m_{client}} = T_{m_{server}})$$
    - T: current time
    - T<sub>c</sub>: time when cache entry was last validated
    - t: freshness interval (3 .. 30 secs in Solaris)
    - T<sub>m</sub>: time when block was last modified at ...
  - consistency level
    - acceptable
    - most UNIX applications do not depend critically on synchronisation of file updates



# Distributed file systems

## NFS: implementation

---

- Performance
  - reasonable performance
    - remote files on fast disk better than local files on slow disk
    - RPC packets are 9 Kb to contain 8Kb disk blocks
  - Lookup operation covers about 50% of server calls
  - Drawbacks:
    - frequent *getattr* calls for time-stamps (cache validation)
    - poor performance of (relative infrequent) writes (because of write-through)



# Distributed file systems

---

- Overview
  - File service model
  - case study: NFS
  - case study: Andrew File System -- AFS
  - comparison NFS <> AFS



# Distributed file systems

## AFS

- Background and aims
  - Base: observation of UNIX file systems (*data driven observations at that time 198X*)
    - files are small ( < 10 Kb)
    - read is more common than write
    - sequential access is common, random access is not
    - most files are not shared
    - shared files are often modified by one user
    - file references come in bursts
  - Aim: combine best of personal computers and time-sharing systems



# Distributed file systems

## AFS

- Background and aims (*cont.*)
  - Assumptions about environment
    - secured file servers
    - public workstations
    - workstations with local disk
    - no private files on local disk
  - Key targets: scalability and security
    - CMU 1991: 800 workstations, 40 servers





# Distributed file systems

## AFS

- Design characteristics
  - whole-file serving
  - whole-file caching
    - ➔ entire files are transmitted, not blocks
  - client cache realised on local disk (relatively large)
    - ➔ lower number of open requests on the network
  - separation between file and directory service



# Distributed file systems

## AFS


- Configuration
  - single (global) name space
  - local files
    - temporary files
    - system files for start-up
  - volume = unit of configuration and management
  - each server maintains a replicated location database (volume-server mappings)



# Distributed file systems

## AFS

- File name space



Directory `afs`  $\equiv$  root of shared file system

# Distributed file systems

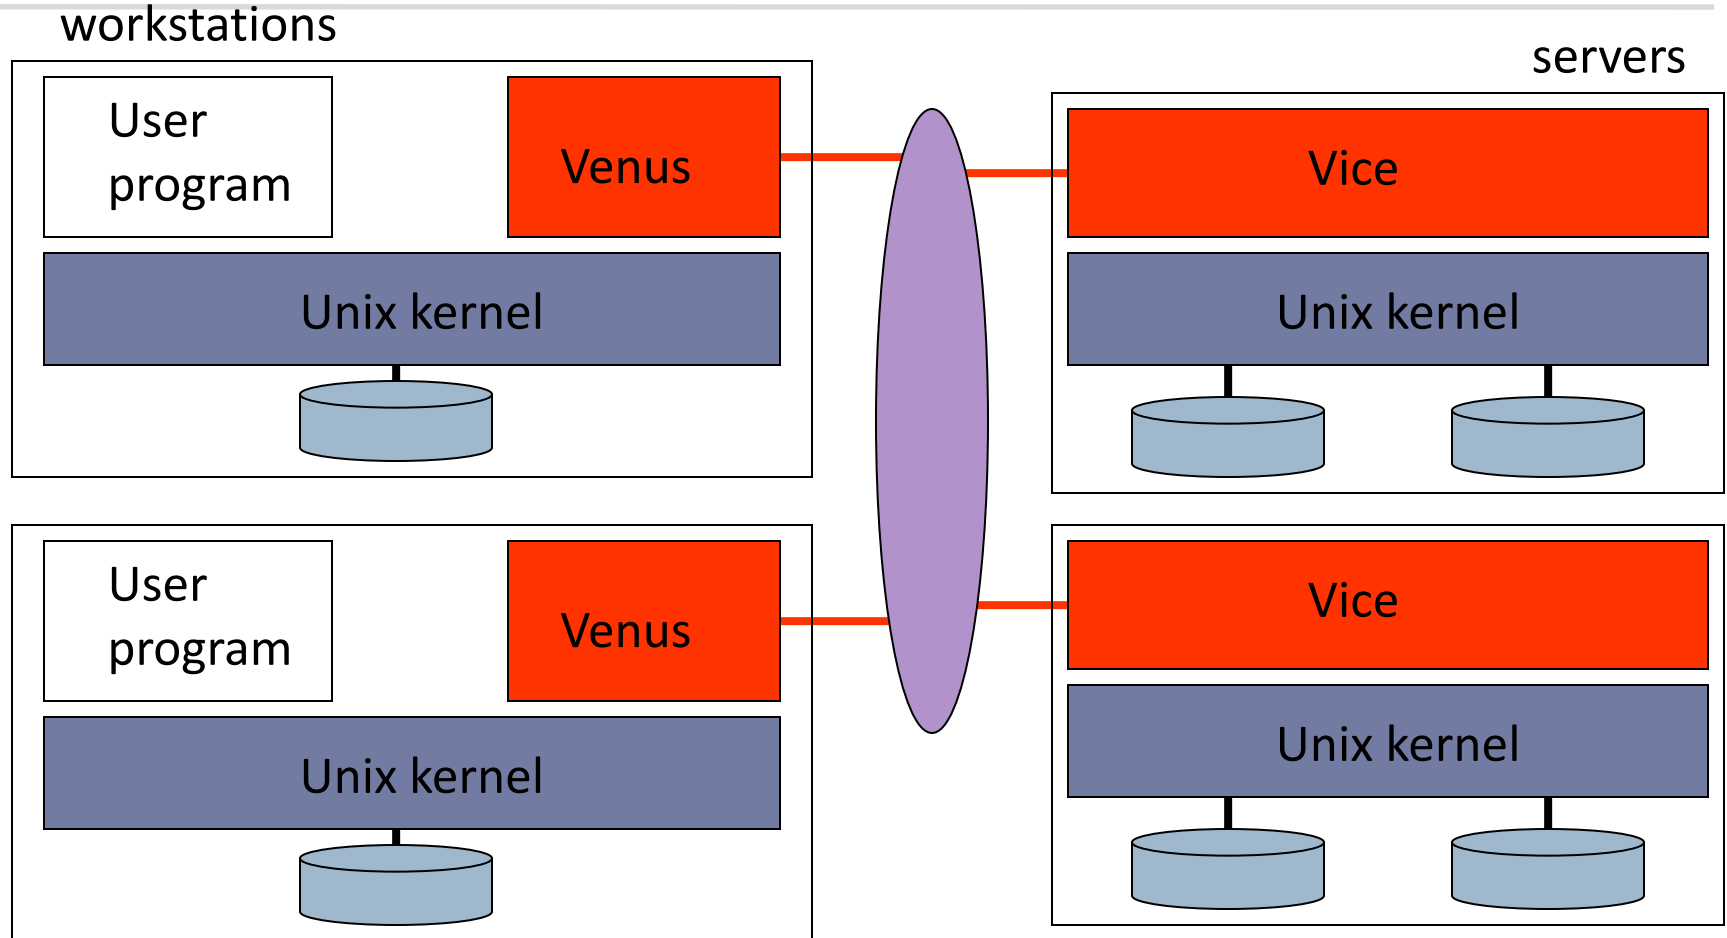
## AFS

- Implementation
  - Vice = file server
    - secured systems, controlled by system management
    - understands only file identifiers
    - runs in user space
  - Venus = user/client software
    - runs on workstations
    - workstations keep autonomy
    - implements directory services
  - kernel modifications for *open* and *close*




# Distributed file systems

## AFS



# Distributed file systems

## AFS



# Distributed file systems

## AFS

- Implementation of file system calls
  - open, close:
    - UNIX kernel of workstation
    - Venus (on workstation)
    - VICE (on server)
  - read, write:
    - UNIX kernel of workstation



# Distributed file systems

## AFS

- Open system call

`open( FileName,...)`

kernel

if FileName in shared space /afs then pass request to Venus

venus

if file not present in local cache OR file present with invalid callback then pass request to Vice server

network

vice

transfer copy of file and valid callback

place copy of file in local file system and store FileName

open local file and return descriptor to user process

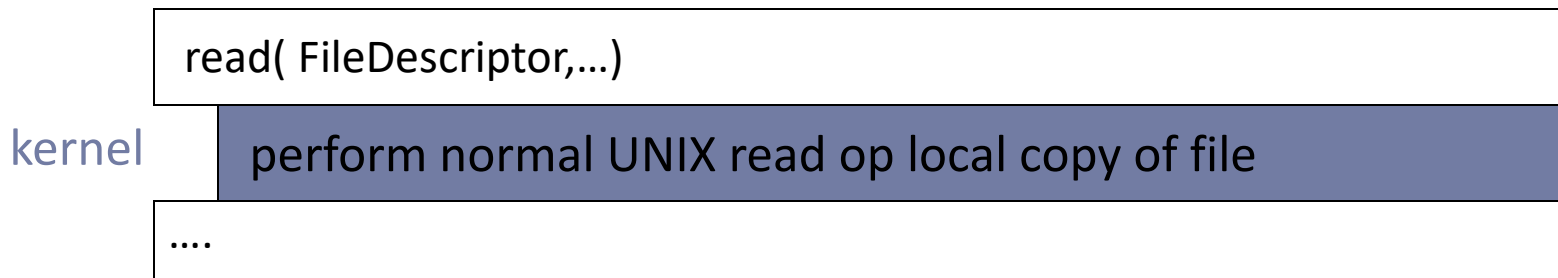




# Distributed file systems

## AFS

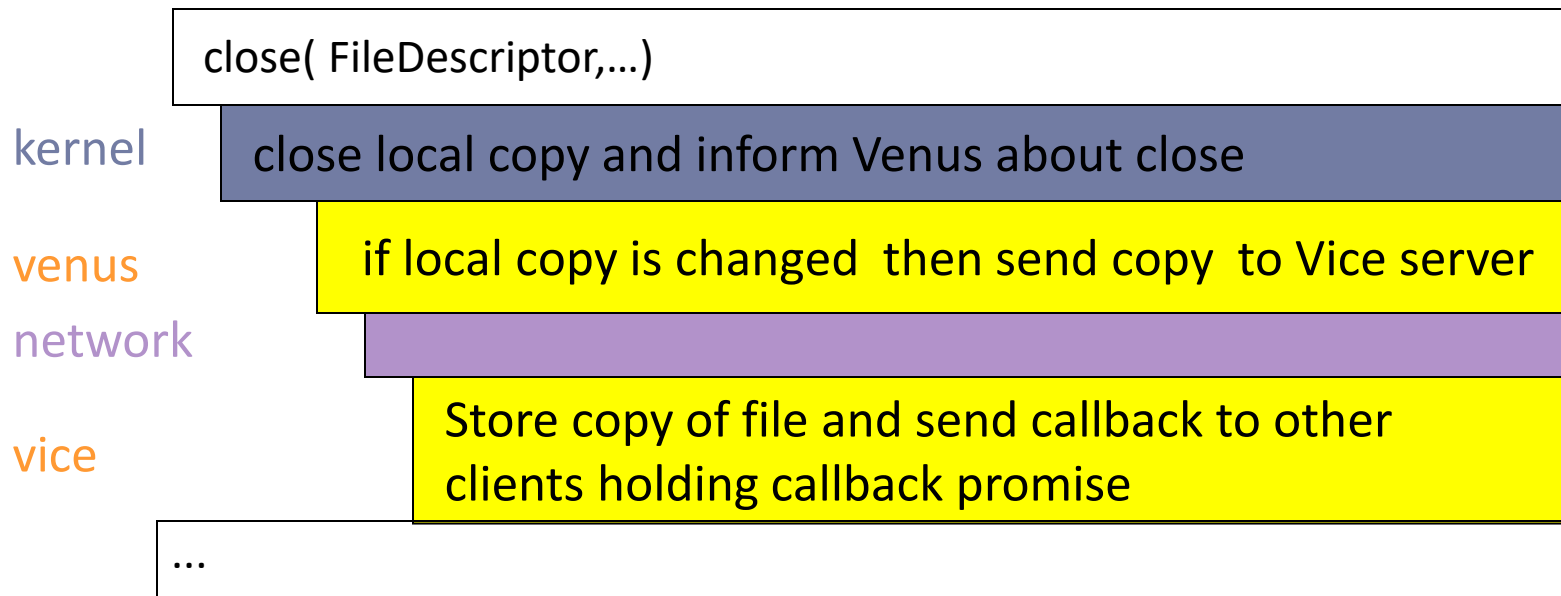
- Read Write system call



# Distributed file systems

## AFS

- Close system call



# Distributed file systems

## AFS

- Caching
  - callback principle
    - service supplies “callback promise” at open
    - “promise” (state) is stored with file in client cache
  - callback promise can be *valid* or *cancelled* (*terminology from original papers; in fact, a boolean value indicating the validity of the file in client-side cache*).
    - initially valid
    - server sends message to cancel callback promise
      - to all clients that cache the file
      - whenever file is updated



# Distributed file systems

## AFS

- Caching maintenance
  - when client workstation reboots
    - cache validation necessary because of missed messages
    - cache validation requests are sent for each valid promise
  - valid callback promises are renewed
    - on *open*
    - when no communication has occurred with the server during a period T



# Distributed file systems

## AFS

- Update semantics
  - guarantee after successful *open* of file **F** on server **S**:  
latest(F, S, 0)  
or  
lostcallback(S, T) **and** incache(F) **and** latest(F, S, T)
  - no other concurrency control
    - 2 copies can be updated at different workstations
    - all updates except from last close are (silently) lost
    - <> normal UNIX operation



# Distributed file systems

## AFS

- Performance: impressive <> NFS
  - benchmark: load on server
    - 40% AFS
    - 100% NFS
  - whole-file caching:
    - reduction of load on servers
    - minimises effect of network latency
  - read-only volumes are replicated (master copy for occasional updates)

Notice: Andrew has been optimised for a specific pattern of use!!



# Distributed file systems

---

- Overview
  - File service model
  - case study: NFS
  - case study AFS
  - comparison NFS <> AFS



# Distributed file systems

## NFS <> AFS

- Access transparency
  - both in NFS and AFS
  - Unix file system interface is offered
- Location transparency
  - uniform view on shared files in AFS
  - in NFS
    - mounting freedom
    - same view possible if same mounting; discipline!





# Distributed file systems

## NFS <> AFS

- Failure transparency
  - NFS
    - no state of clients stored in servers
    - idempotent operations
    - transparency limited by soft mounting
  - AFS
    - state about clients stored in servers
    - cache maintenance protocol handles server crashes
    - limitations?



# Distributed file systems

## NFS <> AFS

- Performance transparency
  - NFS
    - acceptable performance degradation
  - AFS
    - only delay for first open operation on file
    - better than NFS for small files



# Distributed file systems

## NFS <> AFS

- Migration transparency
  - limited: update of **locations** required
  - NFS
    - update configuration files on all clients
  - AFS
    - update of replicated database



# Distributed file systems

## NFS <> AFS

- Replication transparency
  - NFS
    - not supported
  - AFS
    - limited support; for read-only volumes
    - one master for exceptional updates; manual procedure for propagating changes to other volumes



# Distributed file systems

## NFS <> AFS

---

- Concurrency transparency
  - not supported (in UNIX)
- Scalability transparency
  - AFS better than NFS



# Distributed file systems

---

- Overview
  - File service architecture
  - case study: NFS
  - case study AFS
  - comparison NFS <> AFS



# Distributed file systems

- Conclusion beyond the specific of the case study (this often applies)
  - Standards <> quality
  - evolution to standards and common key interface
    - AFS-3 incorporates Kerberos, vnode interface, large messages for file blocks (64Kb)
  - network (mainly access) transparency causes inheritance of weakness: e.g. no concurrency control
  - evolution mainly performance driven



# What if?

---

- Files are read-only, immutable
- Files are append-only?
- Files are relatively/extremely large?
- Devices are extremely poor? (e.g. certain categories of IoT)

*... everything has been thought off before 😊.*







fourth edition

**DISTRIBUTED SYSTEMS**  
CONCEPTS AND DESIGN

George Coulouris  
Jean Dollimore  
Tim Kindberg



# Distributed Systems: File Systems Questions?

