



fourth edition

**DISTRIBUTED SYSTEMS**  
**CONCEPTS AND DESIGN**

George Coulouris  
Jean Dollimore  
Tim Kindberg



KATHOLIEKE UNIVERSITEIT  
**LEUVEN**

# Distributed Systems: Transactions – Part 3

# Overview

- Distributed transactions
  - Flat and nested distributed transactions
  - Atomic commit protocols
  - Concurrency in distributed transactions
  - Distributed deadlocks
  - Transaction recovery

# Distributed transactions

## Locking

- Locks are maintained **locally** (at each server)
  - it decides whether
    - to grant a lock
    - to make the requesting transaction wait
  - it cannot release the lock until it knows whether the transaction has been
    - committed
    - aborted
- at all servers
- deadlocks can occur

# Distributed transactions

## Locking

- Locking rules for nested transactions
  - child transaction inherits locks from parents
  - when a nested transaction commits, its locks are inherited by its parents
  - when a nested transaction aborts, its locks are removed
  - a nested transaction can get a read lock when all the holders of write locks (on that data item) are ancestors
  - a nested transaction can get a write lock when all the holders of read and write locks (on that data item) are ancestors

# Overview

- Transactions
- Distributed transactions
  - Flat and nested distributed transactions
  - Atomic commit protocols
  - Concurrency in distributed transactions
  - Distributed deadlocks
  - Transaction recovery
- Replication

# Distributed deadlocks

- Single server approaches

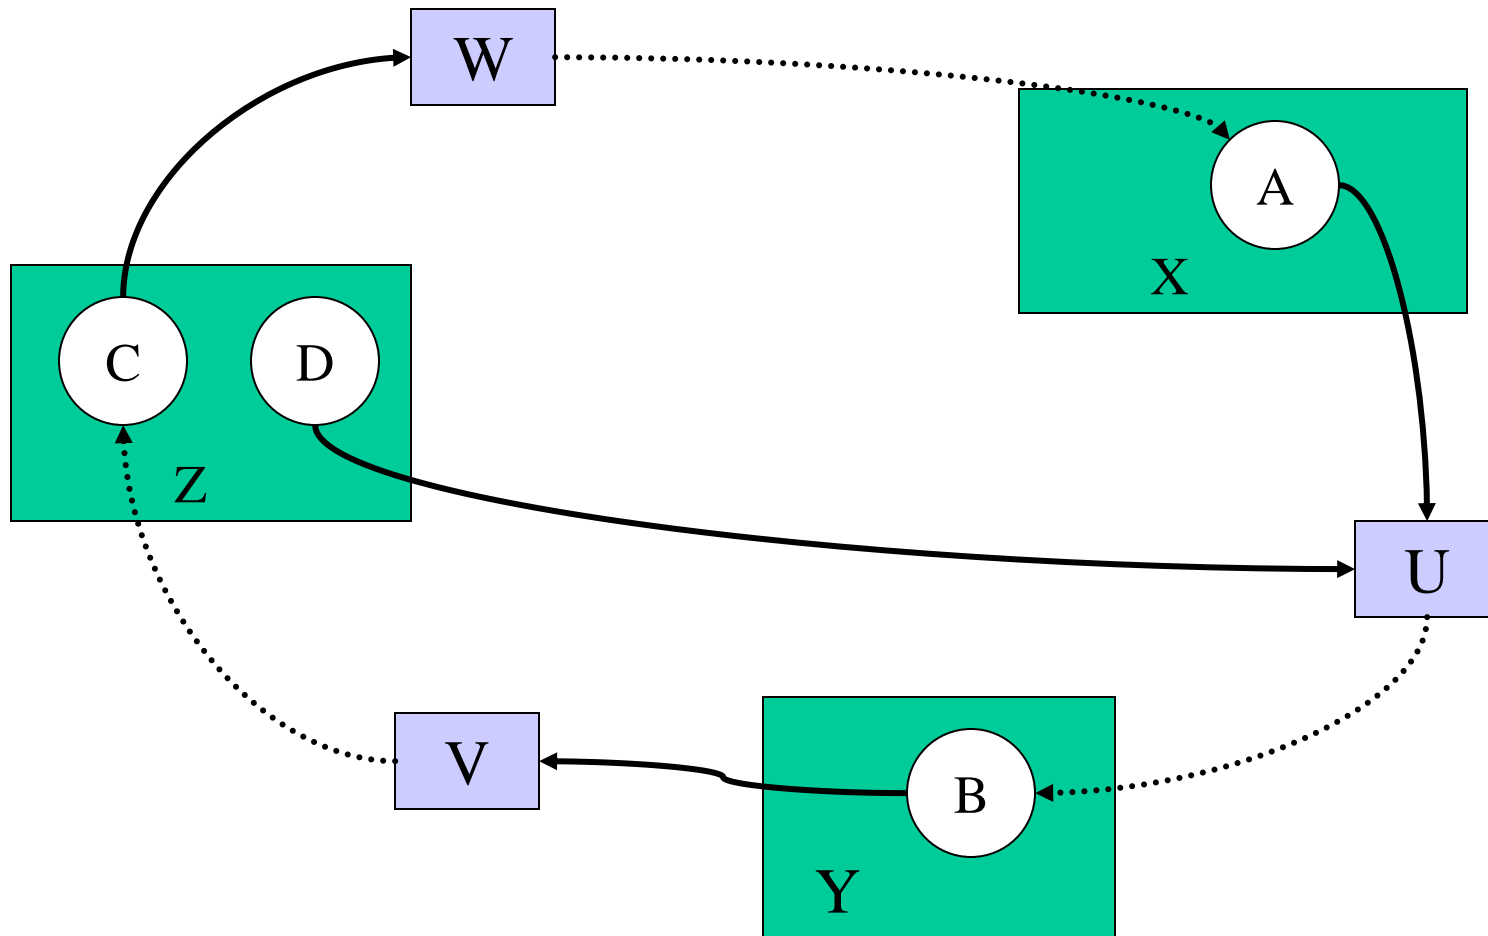
- prevention: difficult to apply
- timeouts: value with variable delays?

## ➔ Detection

- global wait-for-graph can be constructed from local ones
- cycle in global graph possible without cycle in local graph

# Distributed transactions

## Deadlocks



# Distributed transactions

## Deadlocks

- Algorithms:
  - centralised deadlock detection: not a good idea
    - depends on a single server
    - cost of transmission of local wait-for graphs
  - distributed algorithm:
    - Relatively Complex
    - In this course: edge chasing approach



# Distributed transactions

## Deadlocks

- Phantom deadlocks
  - deadlock detected that is not really a deadlock
  - during deadlock detection
    - while constructing global wait-for graph
    - waiting transaction is aborted

# Distributed transactions

## Deadlocks

- Edge Chasing
  - distributed approach to deadlock detection:
    - no global wait-for graph is constructed
  - servers attempt to find cycles
    - by forwarding probes (= messages) that follow edges of the wait-for graph throughout the distributed system

# Distributed transactions

## Deadlocks

- Edge Chasing
  - three steps:
    - initiation: transaction starts waiting
      - new probe constructed
    - detection: probe received
      - extend probe
      - check for loop
      - forward new probe
    - resolution

# Distributed transactions

## Deadlocks

- Edge Chasing: initiation

- send out probe



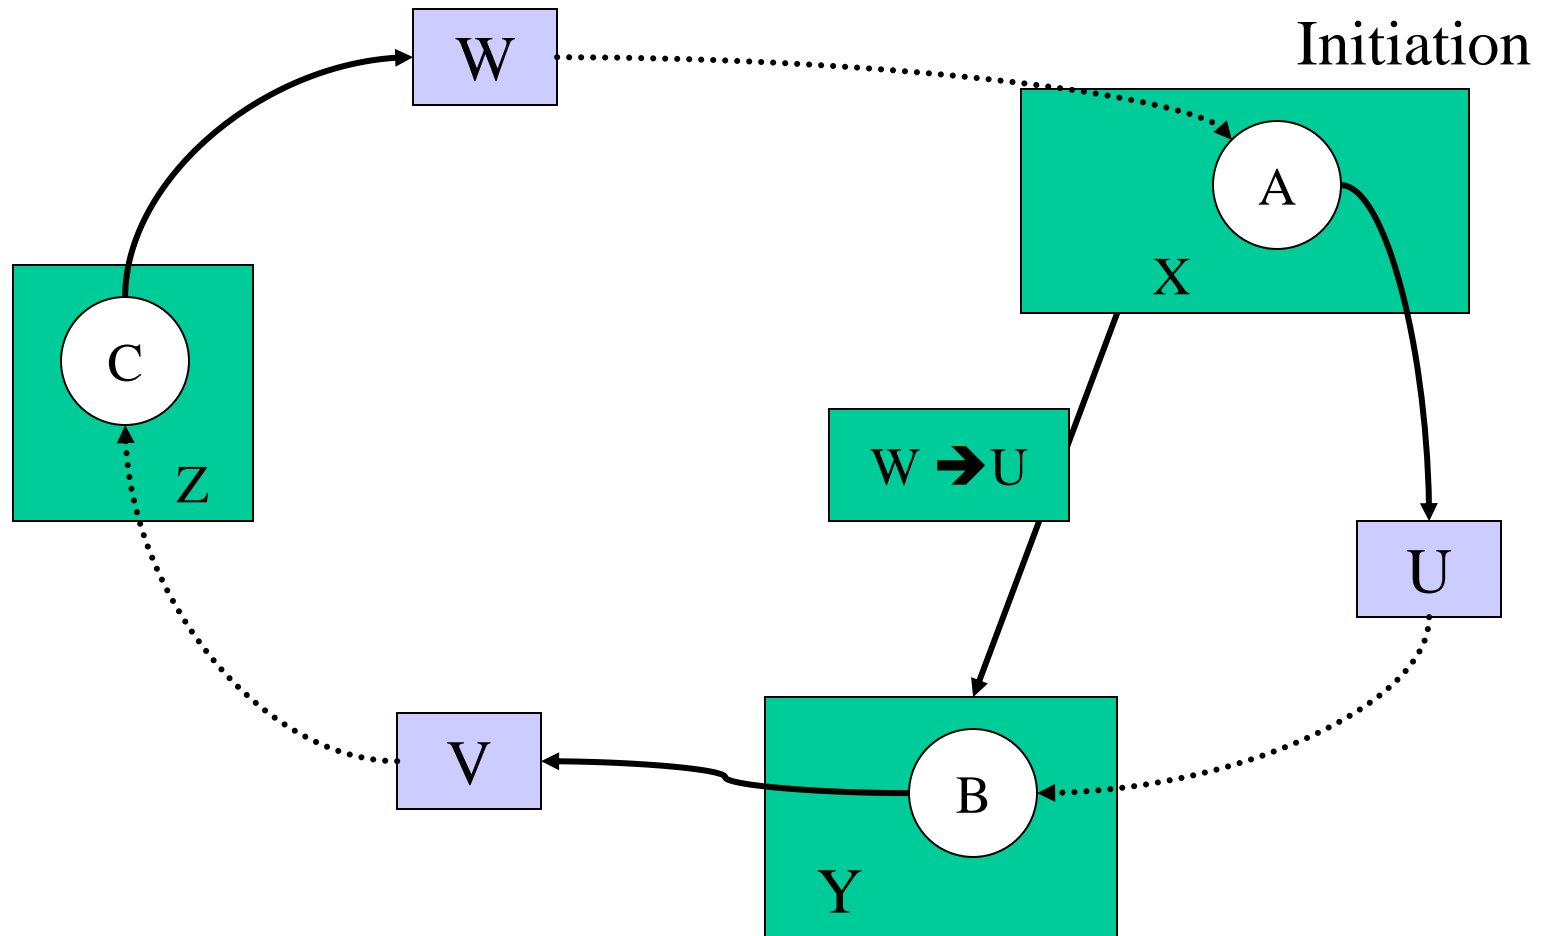
T → U

when transaction T starts waiting for U (and U is already waiting for ...)

- in case of lock sharing, different probes are forwarded

# Distributed transactions

## Deadlocks



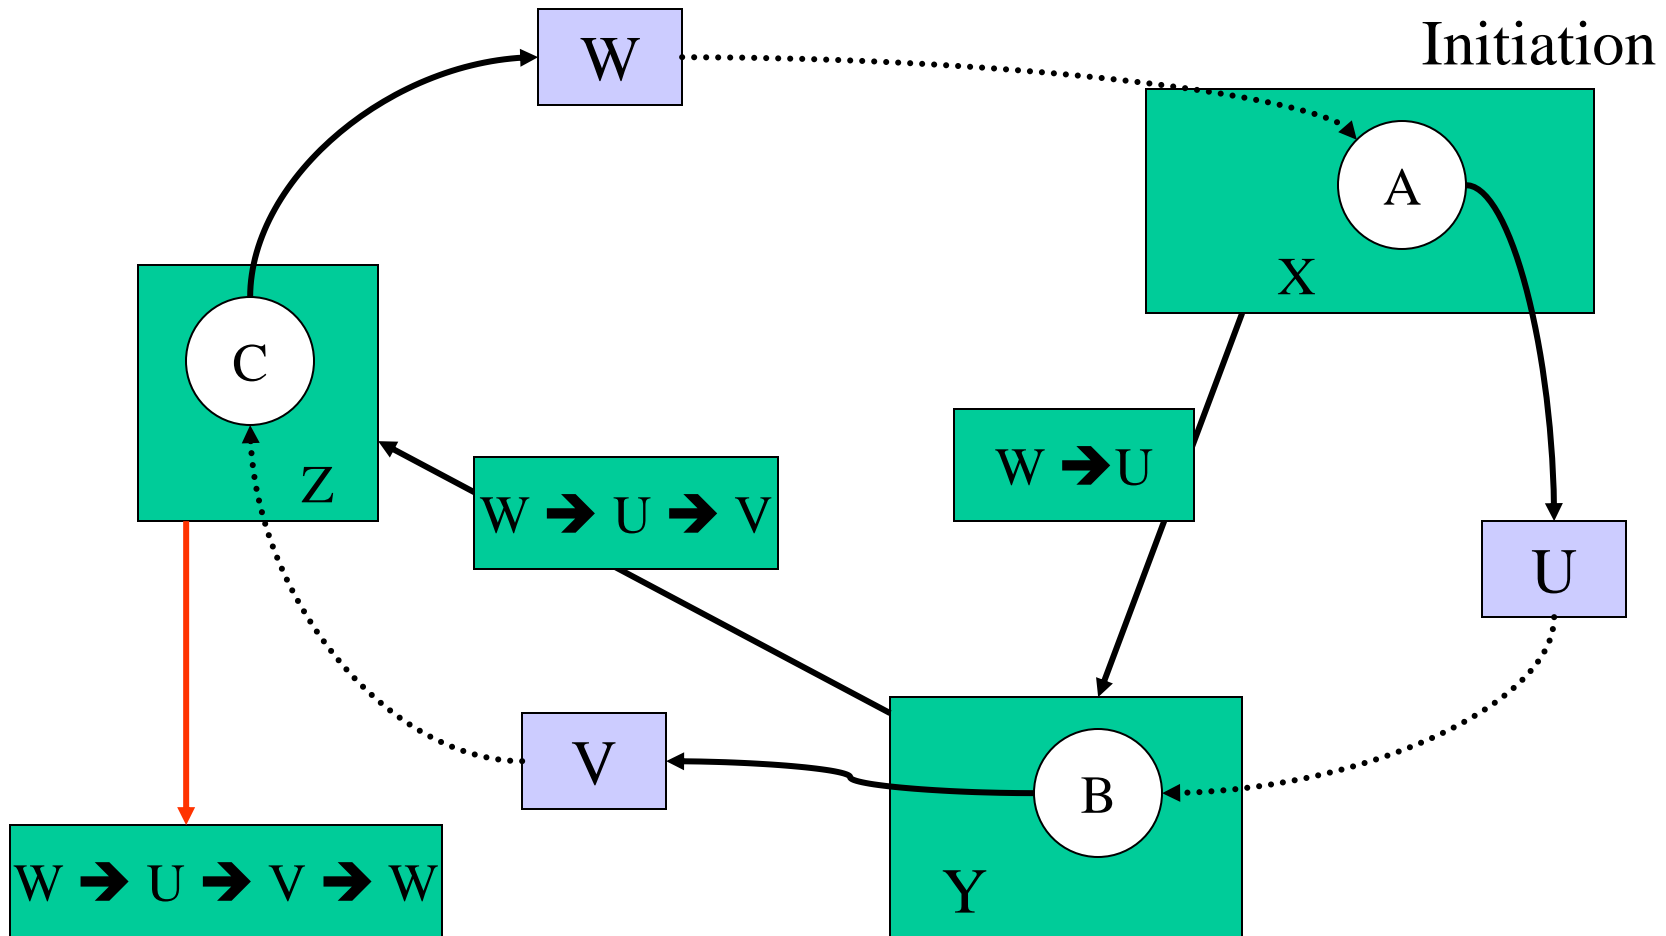
# Distributed transactions

## Deadlocks

- Edge Chasing: detection
  - when receiving probe  $T \rightarrow U$ 
    - Check if U is waiting
    - if U is waiting for V (and V is waiting)  
add V to probe  $T \rightarrow U \rightarrow V$
    - check for loop in probe?
      - yes  $\rightarrow$  deadlock
      - no  $\rightarrow$  forward new probe

# Distributed transactions

## Deadlocks



# Distributed transactions

## Deadlocks

- Edge Chasing: resolution
  - abort one transaction
  - problem?
    - Every waiting transaction can initiate deadlock detection
    - detection may happen at different servers
    - several transactions may be aborted
  - solution: transactions priorities



# Distributed transactions

## Deadlocks

- Edge Chasing: transaction priorities
  - assign priority to each transaction, e.g. using timestamps
  - solution of problem above:
    - abort transaction with lowest priority
    - if different servers detect same cycle, the same transaction will be aborted

# Distributed transactions

## Deadlocks

- Edge Chasing: transaction priorities
  - other improvements
    - number of initiated probe messages ➡
      - detection only initiated when higher priority transaction waits for a lower priority one
    - number of forwarded probe messages ➡
      - probes travel downhill -from transaction with high priority to transactions with lower priorities
      - probe queues required; more complex algorithm

# Overview

- Transactions
- Distributed transactions
  - Flat and nested distributed transactions
  - Atomic commit protocols
  - Concurrency in distributed transactions
  - Distributed deadlocks
  - Transaction recovery (skipped 2023-2024)



fourth edition

## **DISTRIBUTED SYSTEMS** CONCEPTS AND DESIGN

George Coulouris  
Jean Dollimore  
Tim Kindberg



KATHOLIEKE UNIVERSITEIT  
**LEUVEN**

# Distributed Systems: Transactions – Part 3

19/12/2023