



fourth edition

DISTRIBUTED SYSTEMS
CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



KATHOLIEKE UNIVERSITEIT
LEUVEN

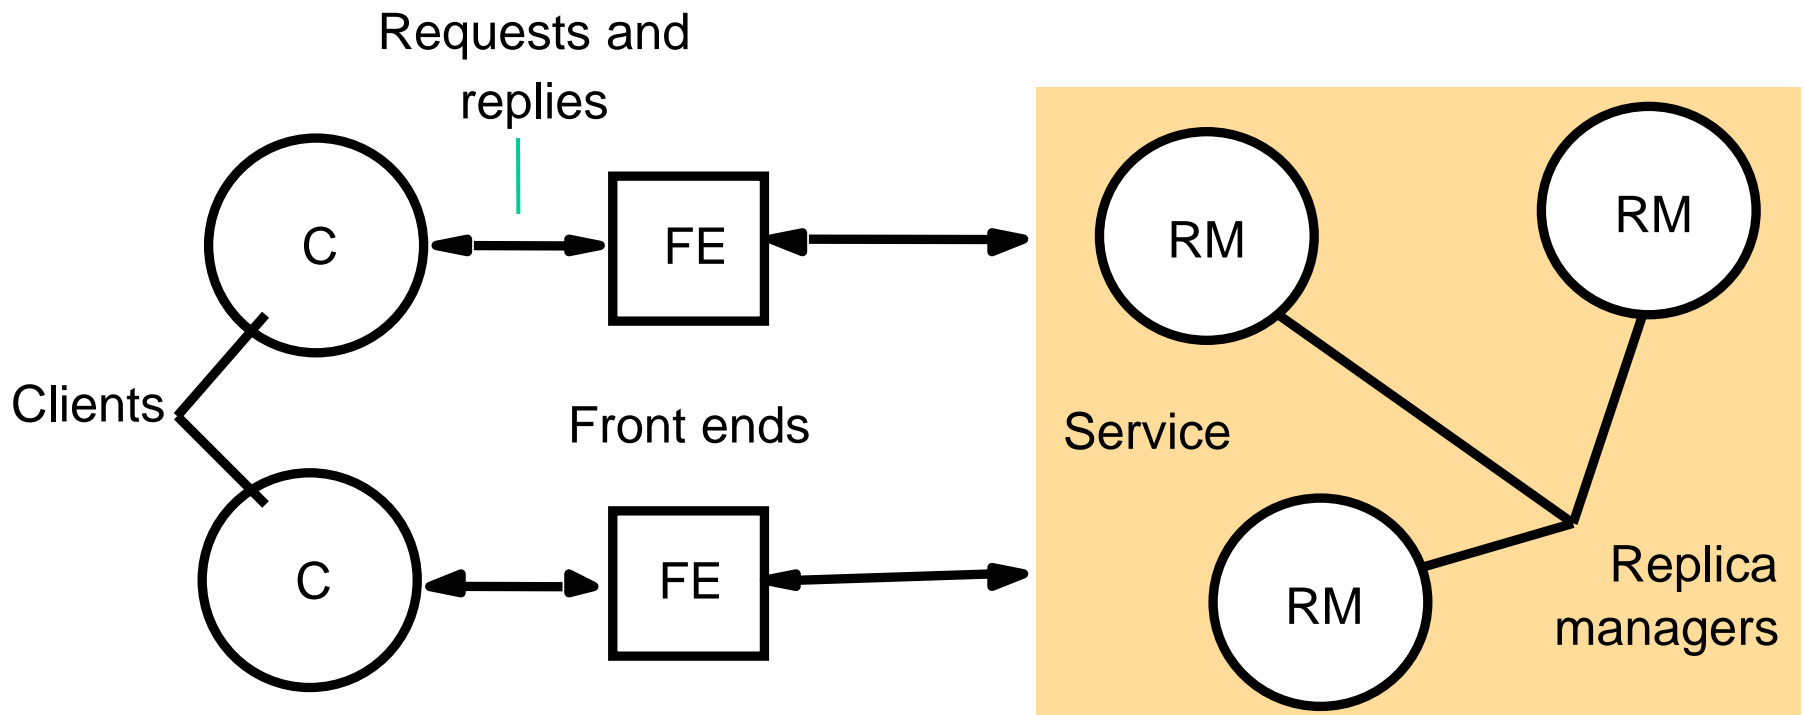
Distributed Systems: Replicated Data

Overview

- Replication
 - System model and group communication
 - Fault-tolerant services
 - Masking failure
 - Highly available services
 - Maximizing service availability

System model and group communication

- Architectural model



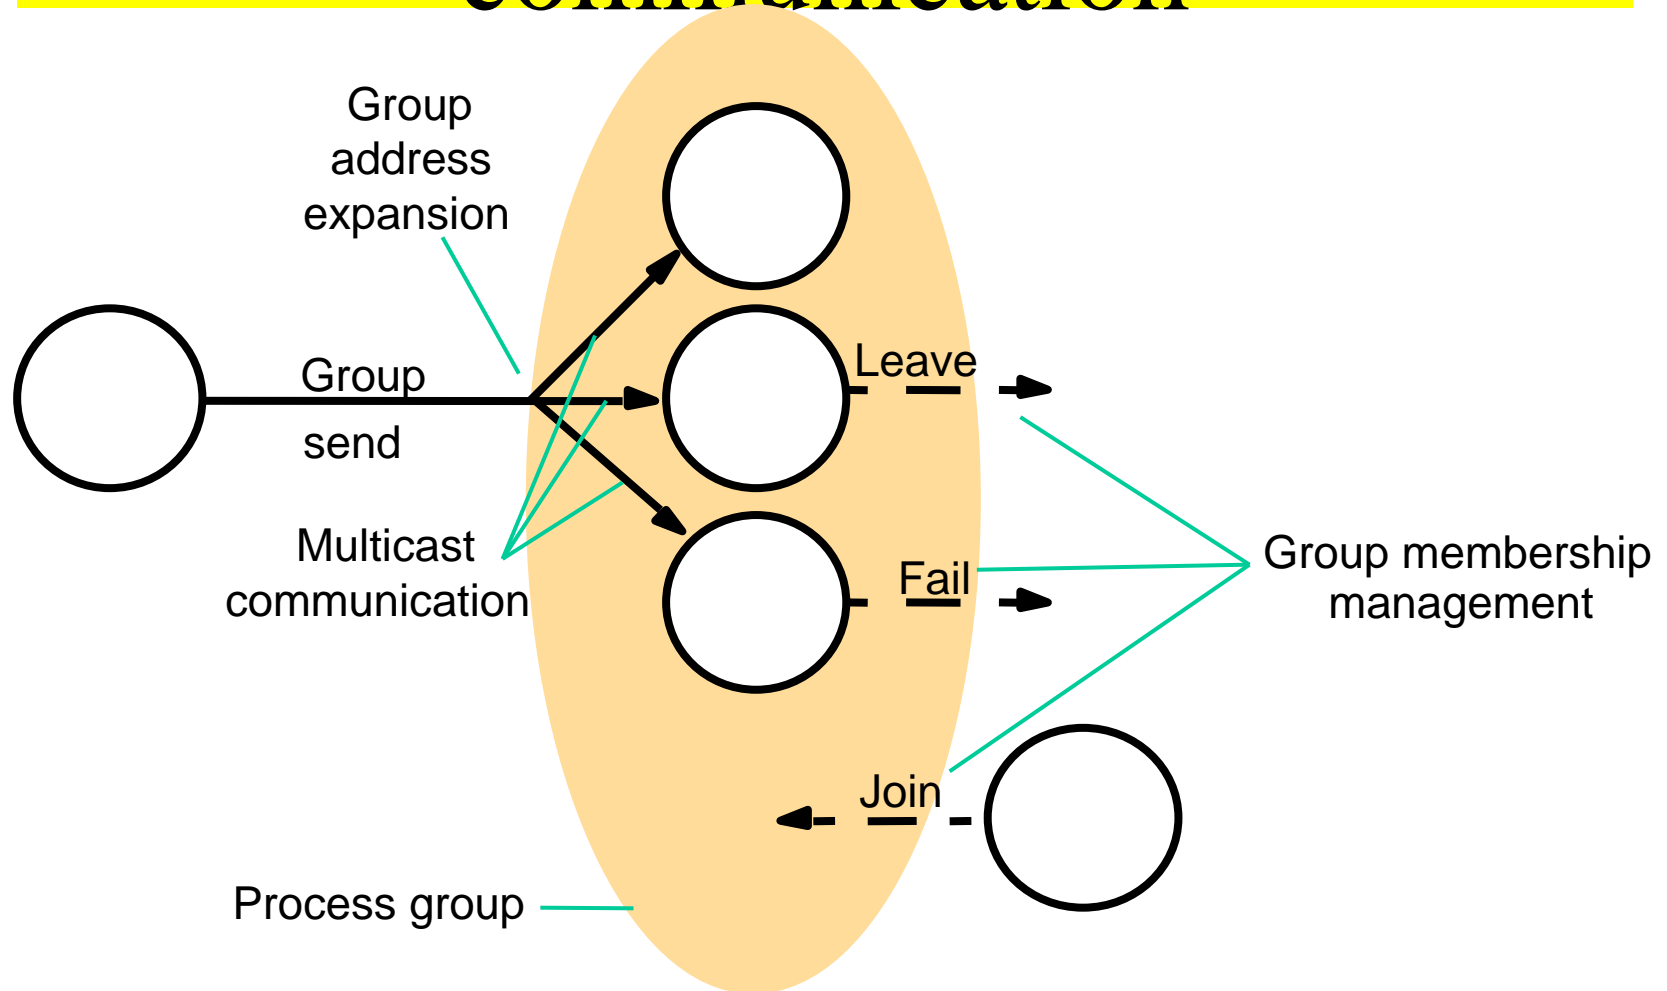
System model and group communication

- 5 phases in the execution of a request:
 - FE issues requests to one or more RMs
 - Coordination: needed to execute requests consistently
 - FIFO
 - Causal
 - Total
 - Execution: by all managers, perhaps tentatively
 - Agreement
 - Response

System model and group communication

- Need for dynamic groups!
- Role of group membership service
 - Interface for group membership changes: create/destroy groups, add process
 - Implementing a failure detector: monitor group members
 - Notifying members of group membership changes
 - Performing group address expansion
- Handling network partitions: group is
 - Reduced: primary-partition
 - Split: partitionable

System model and group communication



System model and group communication


- View delivery
 - To all members when a change in membership occurs
 - $\langle \rangle$ receive view
- Event occurring in a view $v(g)$ at process p
- Basic requirements for view delivery
 - **Order:** if process p delivers $v(g)$ and then $v(g')$ then no process delivers $v(g')$ before $v(g)$
 - **Integrity:** if p delivers $v(g)$ then $p \in v(g)$
 - **Non-triviality:** if q joins group and remains reachable then eventually $q \in v(g)$ at p

System model and group communication


- View-synchronous group communication
 - Reliable multicast + handle changing group views
 - Guarantees
 - **Agreement**: correct processes deliver the same set of messages in any given view
 - **Integrity**: if a process delivers m , it will not deliver it again
 - **Validity**: if the system fails to deliver m to q
then other processes will deliver $v'(g) (=v(g) - \{q\})$
before delivering m

System model and group communication


a (allowed).




b (allowed).



c (disallowed).



d (disallowed).



Overview

- Replication
 - System model and group communication
 - Fault-tolerant services
 - Highly available services
 - *(Not part of the course 2023-2024)*
- Transactions with replicated data

Fault-tolerant services

- Basic Goal: provide a service that is correct despite up to f process failures
- Assumptions:
 - Communication reliable
 - No network partitions
- Meaning of **correct** in case of replication
 - Service keeps responding
 - Clients cannot discover difference with ...
(*transparency*)

Fault-tolerant services

- Naive replication:
 - Client 1 updates account x to 1
 - Client 2 reads account x and sees 0
 - Client 2 updates account y to 2
 - Replication failure
 - Result: Client 2 sees 0 on account x and NOT 1
 - Client 2 sees 2 on account y
 - Result: update of x has been done earlier!!

- Example:

Client 1

setBalance_B(x,1)

setBalance_A(y,2)

Client 2

getBalance_A(y) → 2

getBalance_A(x) → 0

Fault-tolerant services

- Correct behaviour? – *Single copy semantics*
- *Two variants:*
 - Linearizability
 - Strong requirement
 - Sequential consistency
 - Weaker requirement

Fault-tolerant services

- Linearizability
 - Terminology:
 - O_{ij} : client i performs operation j
 - Sequence of operations by one client: $O_{20}, O_{21}, O_{22}, \dots$
 - Virtual interleaving of operations performed by all clients
 - Correctness requirements: \exists interleaved sequence ...
 - Interleaved sequence of operations meets specification of a (single) copy of the objects
 - Order of operations in the interleaving is consistent with the **real times** at which the operations occurred
 - Real time?
 - Yes, we prefer up-to-date information
 - Requires clock synchronization: difficult

Fault-tolerant services

- Sequential consistency
 - Correctness requirements: \exists interleaved sequence ... (red = difference!)
 - Interleaved sequence of operations meets specification of a (single) copy of the objects
 - Order of operations in the interleaving is consistent **with the program order in which each individual client executed them**
 - Example: sequential consistent not linearizable

Client 1

setBalance_B(x,1)

setBalance_A(y,2)


Client 2

getBalance_A(y) \rightarrow 0

getBalance_A(x) \rightarrow 0

Fault-tolerant services

- Passive (primary-backup) replication



Fault-tolerant services


- Passive (primary-backup) replication
 - Sequence of events for handling a client request:
 - **Request:** FE issues request with unique id to primary
 - **Coordination:** request handled atomically in order; if request already handled, re-send response
 - **Execution:** execute request and store response
 - **Agreement:** primary sends updated state to backups and waits for acks
 - **Response:** primary responds to FE; FE hands response back to client
 - Correctness: linearizability
 - *Failures?*

Fault-tolerant services

- Passive (primary-backup) replication
 - Failures?
 - Primary uses view-synchronous group communication
 - Linearizability preserved, if
 - Primary replaced by a unique backup
 - Surviving replica managers agree on which operations had been performed at the replacement point
 - Evaluation:
 - Non-deterministic behaviour of primary supported
 - Large overhead: view-synchronous communication required
 - Variation of the model:
 - Read requests handled by backups: linearizability → sequential consistent

Fault-tolerant services

- Active replication



Fault-tolerant services

- Active replication
 - Sequence of events for handling a client request:
 - **Request:** FE does reliable TO-multicast($g, \langle m, i \rangle$) and waits for reply
 - **Coordination:** every correct RM gets requests in same order
 - **Execution:** every correct RM executes the request; all RMs execute all requests in the same order
 - **Agreement:** not needed
 - **Response:** every RM returns result to FE;
when return result to client?
 - Crash failures: after first response from RM
 - Byzantine failures: after $f+1$ identical responses from RMs
 - Correctness: sequential consistency, not linearizability

Fault-tolerant services

- Active replication
 - Evaluation
 - Reliable + totally ordered multicast \equiv solving consensus
 - ➔ Synchronous system
 - ➔ Asynchronous + failure detectors
 - Overhead!
 - More performance
 - Relax total order in case operations commute:
result of $o_1;o_2 = \text{result } o_2;o_1$
 - Forward read-only request to a single RM

Overview

- Replication
 - System model and group communication
 - Fault-tolerant services
 - Highly available services



Highly available services

- Goal
 - Provide acceptable level of service
 - Use minimal number of RMs
 - Minimize delay for returning result
 - ➔ Weaker consistency \Leftrightarrow single-copy semantics
- Overview (text book)
 - Coda
 - Gossip Architecture <*briefly introcuced, mostly SKIPPED 2023-24*>
 - Bayou <*SKIPPED 2023-24*>



Highly available services

Coda

- Aims: **constant data availability**
 - better performance, e.g. for bulletin boards, databases,...
 - more fault tolerance with increasing scale
 - support mobile and portable computers (disconnected operation)

Approach: AFS + replication



Highly available services

Coda

- Design AFS+
 - file volumes **replicated** on different servers
 - **volume storage group** (VSG) per file volume
 - **Available Volume Storage group** (AVSG) per file volume at a particular instance of time
 - volume **disconnected** when AVSG is empty; due to
 - network failure, partitioning
 - server failures
 - deliberate disconnection of portable workstation



Highly available services

Coda

- Replication and consistency
 - file version
 - integer number associated with file copy
 - incremented when file is changed
 - Coda version vector (CVV)
 - array of (assumed “version”) numbers stored with file copy on a particular server (holding a volume)
 - one value per volume in VSG



Highly available services

Coda

- Replication and consistency: example 1
 - File F stored at 3 servers: S_1, S_2, S_3
 - Initial values for all CVVs: $CVV_i = [1,1,1]$
 - update by C_1 at S_1 and S_2 ; S_3 inaccessible
 - ➔ $CVV_1 = [2,2,1], CVV_2 = [2,2,1], CVV_3 = [1,1,1]$
 - network repaired ➔ conflict detected
 - file copy at S_3 updated
 - ➔ $CVV_1 = [2,2,2], CVV_2 = [2,2,2], CVV_3 = [2,2,2]$



Highly available services

Coda

- Replication and consistency: example 2
 - File F stored at 3 servers: S_1, S_2, S_3
 - Initial values for all CVVs: $CVV_i = [1,1,1]$
 - update by C_1 at S_1 and S_2 ; S_3 inaccessible
 - ➔ $CVV_1 = [2,2,1], CVV_2 = [2,2,1], CVV_3 = [1,1,1]$
 - update by C_2 at S_3 ; S_1 and S_2 inaccessible
 - ➔ $CVV_1 = [2,2,1], CVV_2 = [2,2,1], CVV_3 = [1,1,2]$
 - network repaired ➔ conflict detected
manual intervention or



Highly available services

Coda

- Implementation
 - On open
 - Select one server from AVSG
 - check CCV with all servers in AVSG
 - files in replicated volume remain accessible to a client that can access at least one of the replica
 - load sharing over replicated volumes
 - On close
 - multicast file to AVSG
 - update of CCV
 - manual resolution of conflicts might be necessary



Highly available services

Coda

- Caching: update semantics
 - successful open:

AVSG not empty **and** latest(F, AVSG, 0)

or

AVSG not empty **and** latest(F, AVSG, T) **and**
lostcallback(AVSG, T) **and** incache (F)

or

AVSG empty **and** incache (F)



Highly available services

Coda

- Caching: cache coherence
 - relevant events to detect by Venus within T seconds of their occurrence:
 - enlargement of AVSG
 - shrinking of AVSG
 - lost callback event
 - method: probe message to all servers in VSG of any cached file every T seconds



Highly available services

Coda

- Caching: disconnected operation
 - Cache replacement policy: e.g. least-recently used
 - how support long disconnection of portables:
 - Venus can monitor file referencing
 - users can specify a prioritised list of files to retain on local disk
 - reintegration after disconnection
 - priority for files on server
 - client files in conflict are stored on covolumes; client is informed



Highly available services

Coda

- Performance: Coda <> AFS
 - No replication: no significant difference
 - 3-fold replication & load for 5 users:
load +5%
 - 3-fold replication & load for 50 users
load + 70% for Coda <> +16% for AFS
 - Difference: replication + tuning?
- Discussion
 - Optimistic approach to achieve high availability
 - Use of semantics free conflict detection
(except file directories)



Highly available services

Gossip

- Goal of **Gossip architecture**
 - *Framework* for implementing highly available services
 - Replicate data close to points where groups of clients need it
- Operations:
 - 2 types:
 - Queries: read-only operations
 - Updates: change state (do not read state)
 - FE send operations to any RM
 - selection criterium: available + reasonable response time
 - *Guarantees*



Highly available services

Gossip

- Update ordering:
 - **Causal** least costly
 - **Forced** (= total + causal)
 - **Immediate**
 - Applied in a consistent order relative to any other update at all RMs, independent of order requested for other updates
- Choice
 - Left to application designer!!!
 - Reflects trade-off between consistency and operation cost
 - Implications for users



Highly available services

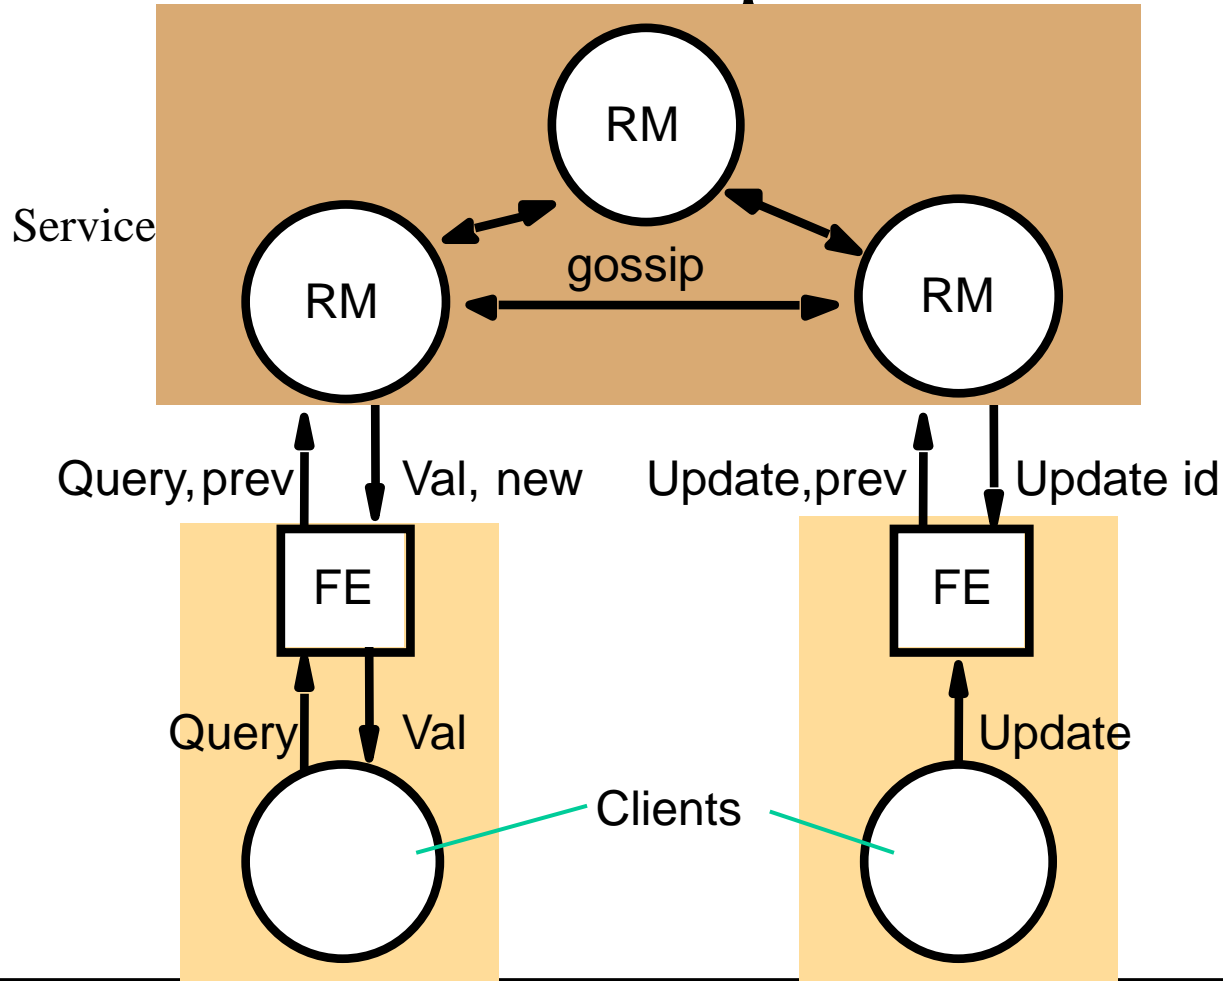
Gossip

- Architecture
 - Clients + FE/client
 - Timestamps added to operations: in next figure
 - Prev: reflects version of latest data values seen by client
 - New: reflects state of responding RM
 - Gossip messages:
 - exchange of operations between RMs



Highly available services

Gossip



Highly available services

Gossip

- Discussion of architecture
 - + Clients can continue to obtain a service even with network partition
 - Relaxed consistency guarantees
 - Inappropriate for updating replicas in near-real time
- Varying properties dependent on choices made in the framework, and by the application developer
- In any very quite different from active & passive replication!





fourth edition

DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



KATHOLIEKE UNIVERSITEIT
LEUVEN

Distributed Systems: Replicated Data - Part 2



Overview

(recap chapter 16-17-18)

- Transactions
- Distributed transactions
- Replication
 - System model and group communication
 - Fault-tolerant services
 - Highly available services

