

Fuzzing (Lecture 3)

Prof. Mathy Vanhoef

DistriNet – KU Leuven – Belgium

Slides partly inspired by: <https://cmu-program-analysis.github.io/2022/lecture-slides/20-fuzzing.pdf>

Security testing of programs

A common technique for security testing is **fuzzing**:

- › Give unexpected or random input to the program
- › Then monitor for crashes, failed assertions, memory leaks,...

Tested program is called the **SUT (System Under Test)**

History of fuzzing

The paper started the field of fuzz testing (1990):

Barton P. Miller, Lars Fredriksen and Bryan So

Study of the Reliability of UNIX Utilities

“

On a dark and stormy night one of the authors was logged on to his workstation on a dial-up line from home and the rain had affected the phone lines; there were frequent **spurious characters on the line**. The author had to race to see if he could type a sensible sequence of characters before the noise scrambled the command. **This line noise** was not surprising; but we were surprised that these spurious characters were causing programs to **crash**.


”

History of fuzzing

The paper started the field of fuzz testing (1990):

Barton P. Miller, Lars Fredriksen and Bryan So

Study of the Reliability of UNIX Utilities



History of fuzzing

The paper started the field of fuzz testing (1990):

Barton P. Miller, Lars Fredriksen and Bryan So

Study of the Reliability of UNIX Utilities

- › Tested ~90 Unix tools by piping random characters into them.
- › 24% of them failed.
- › “... our simple testing technique has discovered a wealth of errors and is **likely to be more commonly used** (at least in the near term) **than more formal** procedures.”

Why fuzzing?

Why is fuzzing useful in practice?

- › Programs **often only undergo functional testing**, i.e., they are tested to handle expected inputs
- › Want to test how programs will react to **unexpected inputs**, since incorrectly handled input can cause vulnerabilities!

Fuzzing is frequently used in practice:

- › AFL++, LibFuzzer, libAFL, Honggfuzz, Boofuzz, and so on...

What bugs can fuzzing detect?

- › Typically, bugs resulting in a **crash**: buffer-overflows, memory leaks, division-by-zero, use-after-free, assert violations, etc.
- › **Root causes**: incorrect argument validation, incorrect type casting, executing untrusted code, etc.
- › Possible **impact**: security, reliability, performance, correctness

Fuzzing recently had many successes

- › Google discovered more than **25,000+ bugs in Chrome...**
...and 36,000+ bugs in more than 550 open-source projects
- › Using SAGE saved Microsoft millions of dollars while creating Windows 7
- › The 2016 DARPA Cyber Grand Challenge winner, Mayhem, heavily relied on white-box fuzzing to find vulnerabilities

Sources:

- <https://google.github.io/clusterfuzz/#trophies>
- “Automated whitebox fuzz testing” by . P. Godefroid, M. Y. Levin, and D. Molnar
- <http://pages.cs.wisc.edu/~bart/fuzz/Foreword1.html>
- “Unleashing Mayhem on binary code” by S. K. Cha, T. Avgerinos, A. Rebert, and D. Brumley

Well-known example: American Fuzzy Lop (AFL)



Partly a “dumb” fuzzer:

- › No model of input. Uses set of **seed inputs**.
- › Given a test input, it **mutates** random bits.

Partly a “smart” fuzzer:

- › Tracks which code in the SUT was executed.
- › Adds input covering new code to the set of interesting inputs

= **coverage-guided fuzzing** = combination of dumb & smart

Example: fuzzing jpeg

- › Start with a single seed input: “hello”
- › Using 7 cores for ~28 hours (i7 2.6 GHz)
- › Interesting inputs are discovered, but not yet a valid jpeg file

```
$ ./djpeg id:002,op:havoc,rep:32,+cov
```

```
Premature end of JPEG file
```

```
Not a JPEG file: starts with 0xff 0xff
```

```
$ ./djpeg id:003,+cov
```

```
Premature end of JPEG file
```

```
JPEG datastream contains no image
```

Example: fuzzing jpeg

```
$ ./djpeg id:000840,sync:fuzzer04
```

```
Corrupt JPEG data: 50 extraneous bytes before marker 0xc4
```

```
Bogus Huffman table definition
```

```
$ ./djpeg id:001032,sync:fuzzer06
```


```
Corrupt JPEG data: 2 extraneous bytes before marker 0xc9
```

```
Quantization table 0x31 was not defined
```

Suddenly valid jpeg's are generated!



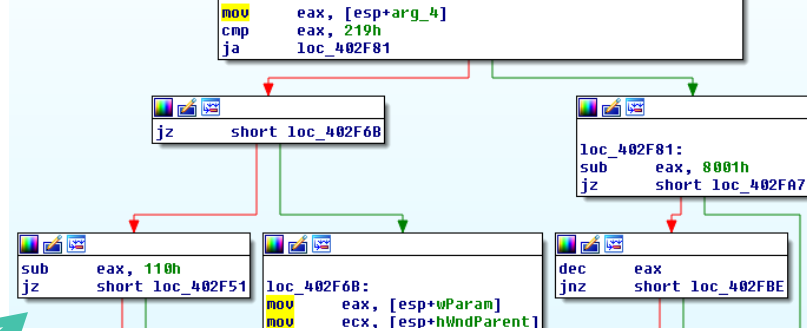
Downside: this can take a long time



How does AFL achieve this?

It monitors which code is executed

- › Doesn't track the actual code path
- › Tracks how many times a branch was taken



Example 1:

Path:

A → B → C → D → E

A → B → D → C → E

Branches:


AB, BC, CD, DE

AB, BD, DC, CE

How does AFL achieve this?

It monitors which code is executed

- › Doesn't track the actual code path
- › Tracks how many times a branch was taken



Example 2:

Path:

A → B → A → C

A → B → A → B → A → C

More hits = different path

Branches:


AB, BA, AC

AB, BA, AC

How does AFL achieve this?

It monitors which code is executed

- › Doesn't track the actual code path
- › Tracks how many times a branch was taken



Queue of “interesting” inputs:

- › Take an input from this queue and mutate it
- › If new path is taken, add mutation to the queue

→ AFL slowly “explores” functionality of program

Taking a step back: types of fuzzers

- › **Black-box** fuzzing: assumes no access to the source code or compiled code. We can only observe input/outputs.
- › **White-box** fuzzing: analyzing the internals of the SUT and the information gathered when executing the SUT.
- › **Gray-box** fuzzing: obtains *some* info internal to the SUT but does not reason about the full semantics of the SUT.

Manes et al.: “Although there usually is some consensus among security experts, the distinction among black-, grey- and white-box fuzzing is not always clear.”

Random vs mutated input

Why not generate purely random input?

- › Purely random data is not very interesting input!

UPnP

```
M-SEARCH * HTTP/1.1
HOST: 239.255.255.250:1900
MAN: "ssdp:discover "
MX: 1
ST: urn:dial-multiscreen-org:service:dial:1
```

Random input is likely
to be (quickly) ignored

```
Asdf;kj4389bnkl;ssa0953t ][asdf q2rjhn
0q23408jfk 902-g9 aDVQ$#Sfgv q543tsdfg
ds dfsgdsfg addfsga
Adfgadfgad 2546ueu6sda asfa[]'']'sdf
```

Random vs mutated input

Why not generate purely random input?

- › Purely random data is not very interesting input!



UPnP

```
M-SEARCH * HTTP/1.1
HOST: 239.255.255.250:1900
MAN: "ssdp:discover "
MX: 1
ST: urn:dial-multiscreen-org:service:dial:1
```

Mutated input more likely
to explore new code paths

```
M-SEARCH * HTTP/1.1
HOAT: 239.255.255.250:190019001900
MAN: "ssdp:@@@@discover "
MXsiEf: 1
ST: urn:dial-multiscreen-org:service:dial:1
```

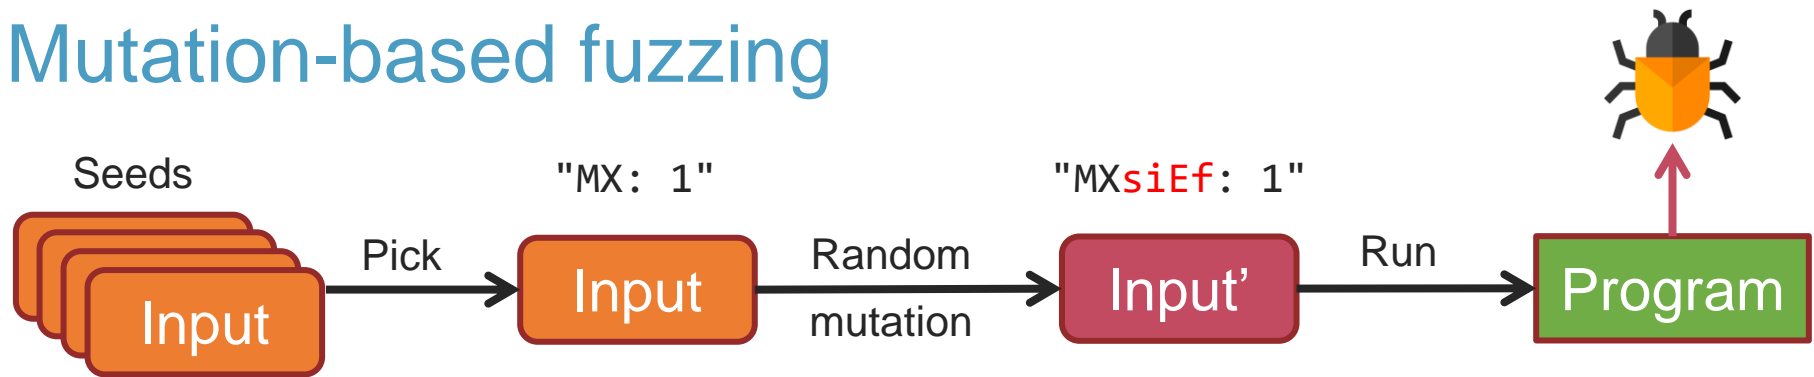
How to mutate binary input?

- › Bit flips, byte flips
- › Change random bytes
- › Insert random byte chunks
- › Delete random byte chunks
- › Set randomly chosen byte chunks to interesting values
 - ›› Examples: INT_MAX, INT_MIN, 0, 1, -1,...

How to mutate text input?

- › Repeat words or add long sequences of characters
- › Insert random symbols or keywords from a dictionary
- › Set randomly chosen words to interesting values
 - ›› Format strings: `"%s%s%s", "%n%n%n", ...`
 - ›› Command injection: `"|touch /tmp/fuzz", ";touch /tmp/fuzz", ...`
 - ›› Directory traversal: `"/.../.../.../.../..."`, ...
 - ›› Binary strings: `"\xde\xad\xbe\xef", "\x00\x00\x00\x00"`, ...
- › What are your ideas?

Mutation-based fuzzing



Valid seed input

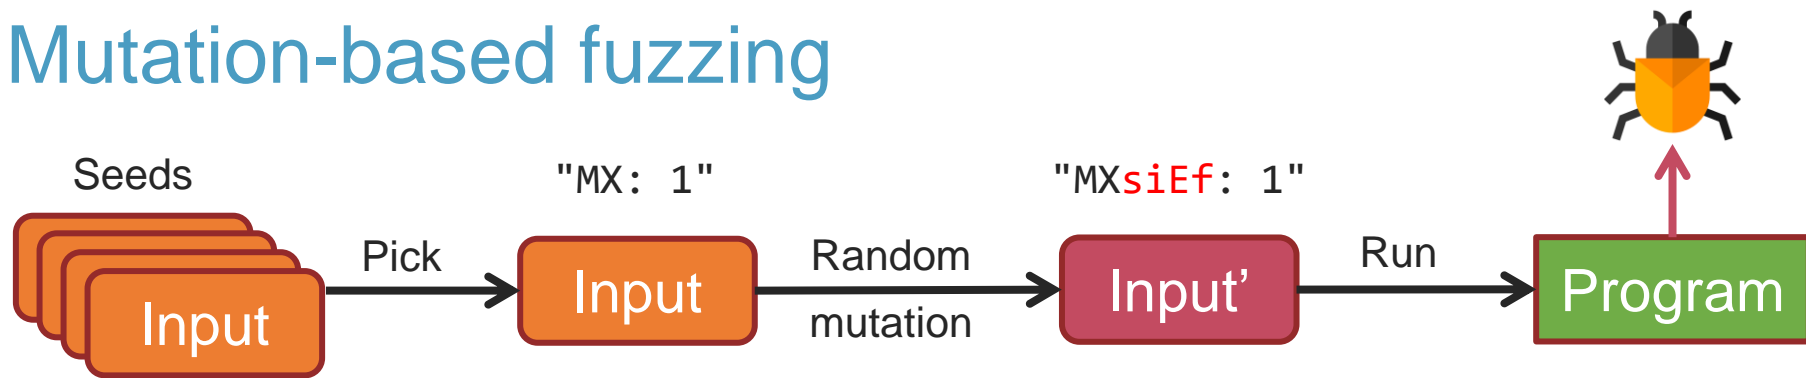
```
M-SEARCH * HTTP/1.1
HOST: 239.255.255.0:1900
MAN: "ssdp:discover "
MX: 1
```

Mutated seed input

```
M-SEARCH * HTTP/1.1
HOAT: 239.255.255.0:19001900
MAN: "ssdp:@@@@discover "
MXsiEff: 1
```

→ Fuzzers using this approach: radamsa, zzuf, etc.

Mutation-based fuzzing



- › Advantage: a **black-box** fuzzing technique which can be used against many products. Fairly easy to set up.
- › Disadvantage is **no code coverage** feedback:
 - › How do you know that new code is being explored? How to know that your mutations are meaningful?
 - › When to stop fuzzing?

What kind of coverage to collect?

- › Function coverage

- » foo(F, F, F);

- › Statement coverage

- » foo(T, T, T);

- › Branch/Decision coverage

- » foo(T, T, T);

- » foo(T, T, F);

- › Condition coverage

- » foo(F, F, T);

- » foo(T, T, F);

```
int foo (bool a, bool b, bool c)
{
    int ret = 0;
    if ((a || b) && c)
    {
        ret = 1;
    }
    return ret;
}
```


What kind of coverage to collect?

- › Modified condition/decision coverage (MC/DC): every condition must be True/False once *and* affect the outcome

- › foo(F, T, F);

- › foo(F, T, T);

- › foo(F, F, T);

- › foo(T, F, T);

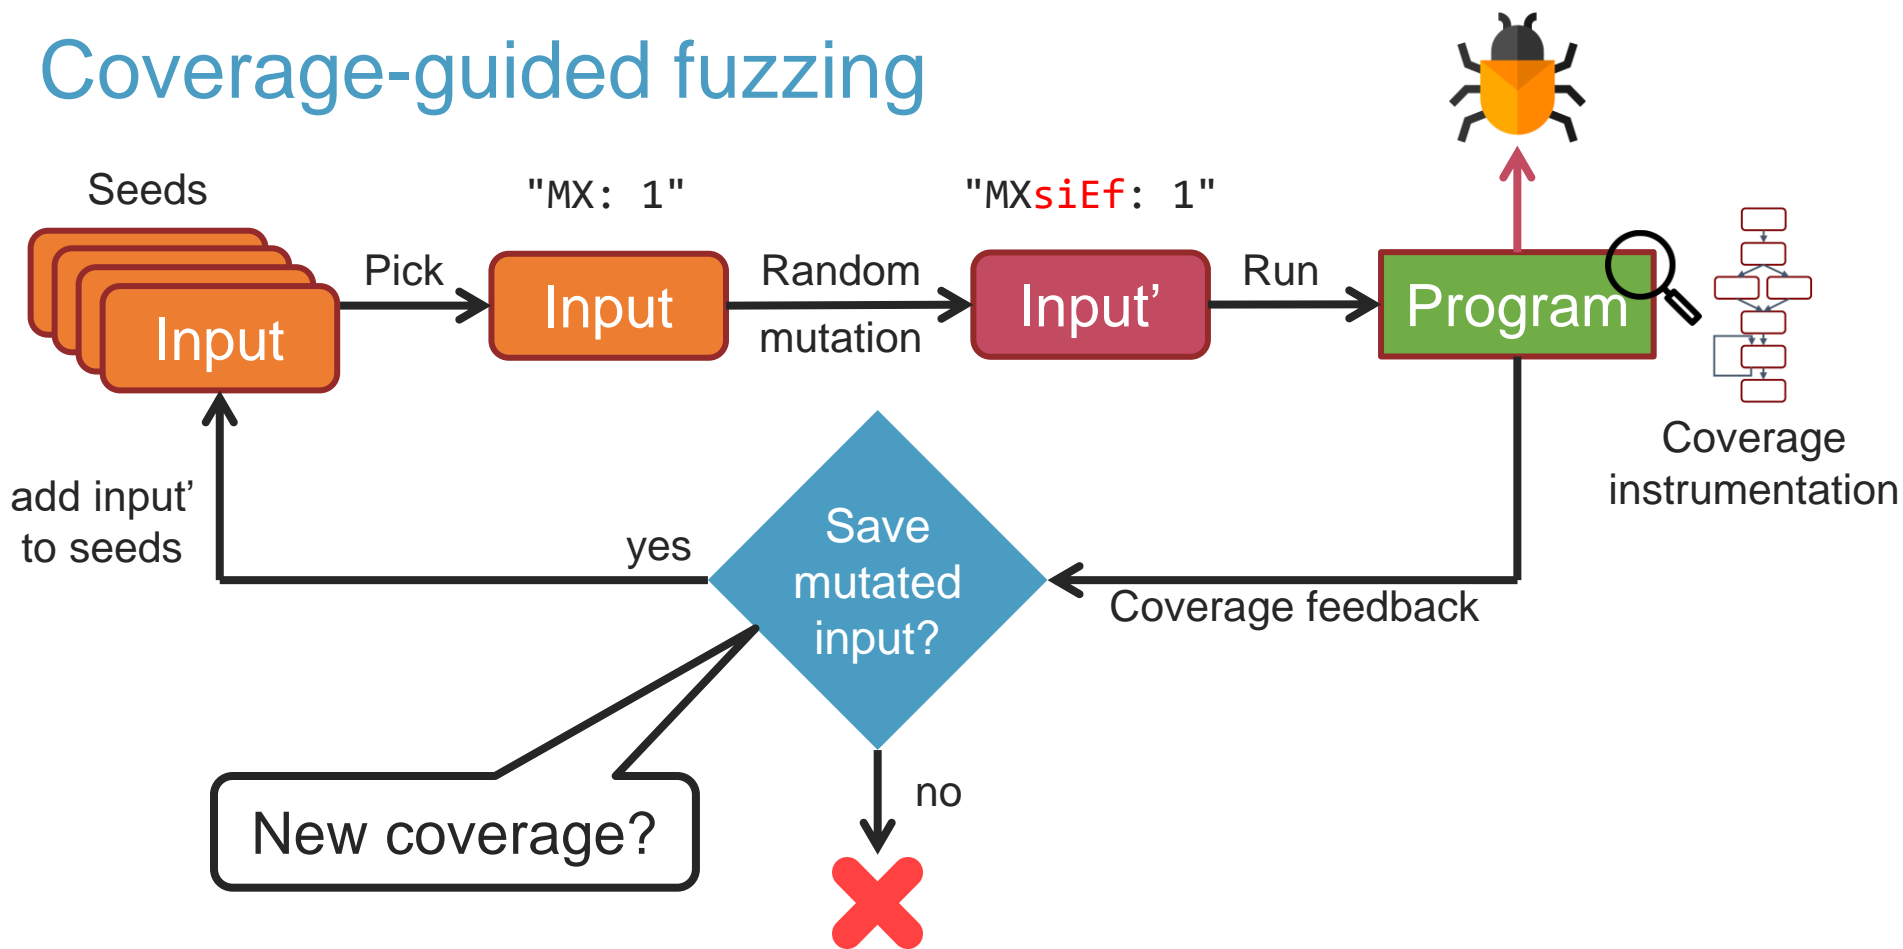
- › Multiple condition coverage

- › Parameter value coverage

- › ...

```
int foo (bool a, bool b, bool c)
{
    int ret = 0;
    if ((a || b) && c)
    {
        ret = 1;
    }
    return ret;
}
```

Coverage-guided fuzzing



Code coverage details in AFL

- › Branch coverage with coarse branch-taken hit counts
- › Code injected at every branch is equivalent to:

```
cur_location = <COMPILE_TIME_RANDOM>; // identify code block
shared_mem[cur_location ^ prev_location]++;
prev_location = cur_location >> 1;
```

- › Every basic code block is assigned a random identifier
 - ›› Simplifies assigning IDs when linking complex projects
 - ›› Keeps the XOR output uniformly distributed

Code coverage details in AFL

- › Branch coverage with coarse branch-taken hit counts
- › Code injected at every branch is equivalent to:

```
cur_location = <COMPILE_TIME_RANDOM>;  
shared_mem[cur_location ^ prev_location]++;  
prev_location = cur_location >> 1;
```

- › shared_mem is added by the instrumentation (typically 64 kB)
- › Bytes in this map represent (branch_src, branch_dst) tuples
 - › Size of the map is chosen so that collisions are sporadic...
 - › ...and so the map is small enough to be rapidly analyzed/compared

Code coverage details in AFL

- › Branch coverage with coarse branch-taken hit counts
- › Code injected at every branch is equivalent to:

```
cur_location = <COMPILE_TIME_RANDOM>;  
shared_mem[cur_location ^ prev_location]++;  
prev_location = cur_location >> 1;
```

- › Shift operation in the last line tracks the directionality of tuples
 - › Without this, $A \wedge B$ would be indistinguishable from $B \wedge A$
- › Shift operation also retains the identity of tight loops
 - › Otherwise, $A \wedge A$ would be equal to $B \wedge B$

Code coverage details in AFL

- › Branch coverage with coarse branch-taken hit counts
- › Code injected at every branch is equivalent to:

```
cur_location = <COMPILE_TIME_RANDOM>;  
shared_mem[cur_location ^ prev_location]++;  
prev_location = cur_location >> 1;
```

- › The tuple hit counts are divided into several buckets:
 - ›› 1, 2, 3, 4-7, 8-15, 16-31, 32-127, 128+
 - ›› Changes within a bucket are ignored. Transition from one bucket to another is flagged as an interesting new code coverage.

Does it work in practice?

The bug-o-rama trophy case

Yeah, it finds bugs. I am focusing chiefly on development and have not been running the fuzzer at a scale, but here are some of the notable vulnerabilities and other uniquely interesting bugs that are attributable to AFL (in large part thanks to the work done by other users):

IJG jpeg ¹	libjpeg-turbo ^{1 2}	libpng ¹
libtiff ^{1 2 3 4 5}	mozjpeg ¹	PHP ^{1 2 3 4 5 6 7 8}
Mozilla Firefox ^{1 2 3 4}	Internet Explorer ^{1 2 3 4}	Apple Safari ¹
Adobe Flash / PCRE ^{1 2 3 4 5 6 7}	sqlite ^{1 2 3 4...}	OpenSSL ^{1 2 3 4 5 6 7}
LibreOffice ^{1 2 3 4}	poppler ^{1 2...}	freetype ^{1 2}

What about fuzzing in general?

bugs chromium New issue All issues label:ClusterFuzz Sign in

1 - 100 of 40265 Next > List Grid Chart

ID	Pri	Type	Component	Status	Summary + Labels
1379144	1	Bug	Blink>JavaScript>Runtime	Untriaged	Abt in v8::internal::Isolate::PushStackTraceAndDie Re
1379142	2	Bug	Internals>Storage	Assigned	dawn_wire_server_and_frontend_fuzzer: Abt in Handl ClusterFuzz
1379138	1	Bug	Blink>JavaScript	Assigned	extension_file_highlighter_fuzzer: Abt in extensions::C Reproducible ClusterFuzz
1379001	1	Bug	Internals>GPU>Dawn	Assigned	tint_vertex_pulling_fuzzer: Null-dereference READ in
1379000	1	Bug	Platform>Extensions	Assigned	
1378999	2	Run	Internals>GPI I>Tint	Assigned	

ClusterFuzz @ Google found 40000+ bugs in chromium!

Source: <https://bugs.chromium.org/p/chromium/issues/list?q=label%3AClusterFuzz&can=1>

Detecting bugs

By default, a fuzzer will only detect crashes

- › State-of-the-art fuzzers can also detect certain logical implementation flaws (mainly academic work)

Not all programs will immediately crash...

- › Buffer overflow may overwrite unused data → no crash
- › Reading memory outside of buffer → will likely succeed
- › Undefined behavior may only lead to a crash on specific platforms

Sanitizers

- › **Address** Sanitizer (ASAN)
- › **Leak**Sanitizer (comes with ASAN)
- › Threat Sanitizer (TSAN): detects **data races and deadlocks** for C++ and Go
- › **Undefined-behavior** Sanitizer (UBSan)
- › MemorySanitizer: detects use of **uninitialized memory**
- › **Hardware-assisted** AddressSanitizer (HWASAN): variant of AddressSanitizer that consumes much less memory


AddressSanitizer

Compile with ``clang -fsanitize=address``

Asan is a memory error detector for C/C++. It finds:

- › User after free (dangling pointer dereference)
- › Heap/stack/global buffer overflow
- › Use after return/scope
- › Initialization order bugs
- › Memory leaks

AddressSanitizer: Overhead



Source <https://github.com/google/sanitizers/wiki/AddressSanitizerPerformanceNumbers>

Crash Triaging

american fuzzy lop 2.36b (████████)		
process timing		overall results
run time : 0 days, 0 hrs, 5 min, 20 sec		cycles done : 0
last new path : 0 days, 0 hrs, 0 min, 9 sec		total paths : 241
last uniq crash : 0 days, 0 hrs, 0 min, 49 sec		uniq crashes : 14
last uniq hang : 0 days, 0 hrs, 0 min, 19 sec		uniq hangs : 22
cycle progress	map coverage	
now processing : 121 (50.21%)	map density : 0.23% / 0.87%	
paths timed out : 0 (0.00%)	count coverage : 2.34 bits/tuple	
stage progress	findings in depth	
now trying : interest 32/8	avored paths : 51 (21.16%)	
stage execs : 3550/8883 (39.96%)	new edges on : 75 (31.12%)	
total execs : 777k	total crashes : 140 (14 unique)	
exec speed : 3560/sec	total hangs : 400 (22 unique)	
fuzzing strategy yields	path geometry	
bit flips : 91/30.7k, 15/30.7k, 6/30.6k	levels : 3	
byte flips : 1/3838, 1/3542, 2/3510	pending : 217	
arithmetics : 42/198k, 3/71.9k, 0/32.0k	pend fav : 38	
known ints : 3/19.1k, 7/84.4k, 22/132k	own finds : 239	
dictionary : 0/0, 0/0, 5/23.3k	imported : n/a	
havoc : 55/106k, 0/0	stability : 100.00%	
trim : 22.95%/1711, 7.22%		
[cpu:301%]		

Crash Triaging

Given crashing inputs x_1 and x_2 , do they trigger the same bug?

- › **Very** difficult to answer in practice
- › Heuristics to check if $\text{bug}(x_1) == \text{bug}(x_2)$:
 - › $\text{exitcode}(x_1) == \text{exitcode}(x_2)$ // or exception or error msg
 - › $\text{coverage}(x_1) == \text{coverage}(x_2)$
 - › $\text{stacktrace}(x_1) == \text{stacktrace}(x_2)$
 - › $\text{newcoverage}(x_1, \text{old}) == \text{newcoverage}(x_2, \text{old})$
 - › $\text{fix}(x_1) == \text{fix}(x_2)$

Crash Triaging in AFL

A crash is considered unique if any of two conditions are met:

- › Crash trace includes a **tuple not seen in previous** crashes
- › Crash trace is **missing a tuple that was always present** in earlier crashes.

Note: may result in unique crash count inflation early on, but it exhibits a very strong self-limiting effect after early stage.

Fuzzers in the build pipeline?

Example is [ClusterFuzzLite](#). Has a “code-change” mode:

- › Will fuzz the code base for every pull request or commit
- › Defaults to fuzzing for 10 minutes
- › Quits after finding a single crash (will not run other fuzzers)
- › Only reports new crashes (likely introduced due to commit)

Fuzzers in parallel to the build pipeline?

ClusterFuzz also has a “batch” fuzzing mode to fuzz for longer durations (called “continuous fuzzing”):

- › Report all crashes, not just new ones.
- › Useful to build an input corpus that can be used during “code-change” fuzzing.
- › Continuous fuzzing:
 - ›› Reuse input corpus of fuzzing old version when fuzzing new version
 - ›› Run the fuzzers for a long time! [Inputs may keep being discovered.](#)

Many challenges when fuzzing in practice

- › Fuzzing heuristics:
 - ›› Which input to change? How many times? Which mutations?
 - ›› Type of code coverage to use?
- › Detecting bugs: what is a bug? Infinite loops? Race conditions? How to know when a bug was found?
- › Debugging crashes:
 - ›› Assuring crashes are reproducible
 - ›› Identifying unique crashes/vulnerabilities
 - ›› Input minimization

Many challenges when fuzzing in practice

- › Fuzzing roadblocks:
 - › Handling magic bytes & checksums. May require writing a test harness.
 - › Dependencies in binary inputs (length of chunks, index into table)
 - › Inputs with complex syntax and semantics (e.g. XML, JSON, C++)
 - › Stateful applications like network protocols (see next lecture)
 - › There's ongoing research to handle these obstacles
- Fuzzing is partly a science and partly an art!

Required reading & references

Required reading:

- › Technical "whitepaper" for afl-fuzz. Can be found in the afl repository in the file [technical_details.txt](#)

Optional reading:

- › “[The Art, Science, and Engineering of Fuzzing: A Survey](#)” by Manes et al., IEEE Transactions on Software Engineering.
- › “[Registered Report: Dissecting American Fuzzy Lop](#)” by Fioraldi et al., International Fuzzing Workshop (2022).