

AFL++ & KLEE demos (Lecture 5 part 1)

Prof. Mathy Vanhoef

DistriNet – KU Leuven – Belgium

AFL++ Demo: “JPEGs out of thin air”

AFL guide (!!): https://aflplus.plus/docs/fuzzing_in_depth/

- › Target: <https://github.com/libjpeg-turbo/libjpeg-turbo>
- › Read build instructions of project → uses cmake
 - ›› If needed, search how to change the compiler
 - ›› Example solution for cmake: <https://stackoverflow.com/a/17275650>
- › Essential commands:

```
mkdir build && cd build
CC=afl-cc CXX=afl-c++ cmake ..
# create in_dir and out_dir and fill with seeds
afl-fuzz -i in_dir/ -o out_dir/ ./djpeg @@
```

AFL++ Demo: “JPEGs out of thin air”

Windows PowerShell

```
american fuzzy lop ++4.09a {default} (./djpeg) [fast]

process timing | overall results
  run time   : 0 days, 0 hrs, 8 min, 36 sec | cycles done :
  last new find : 0 days, 0 hrs, 0 min, 9 sec | corpus count : 200
  last saved crash : none seen yet | saved crashes : 0
  last saved hang : none seen yet | saved hangs : 0
cycle progress | map coverage
  now processing : 0.2 (0.0%) | map density : 4.63% / 8.49%
  runs timed out : 0 (0.00%) | count coverage : 2.22 bits/tuple
stage progress | findings in depth
  now trying : bitflip 1/1 | favored items : 2 (1.00%)
  stage execs : 28.3k/290k (9.73%) | new edges on : 75 (37.50%)
  total execs : 31.4k | total crashes : 0 (0 saved)
  exec speed : 59.28/sec (slow!) | total tmouts : 9 (0 saved)
fuzzing strategy yields | item geometry
  bit flips : 0/32, 0/31, 0/29 | levels : 2
  byte flips : 0/4, 0/3, 0/1 | pending : 197
  arithmetics : 0/224, 0/0, 0/0 | pend fav : 0
  known ints : 0/23, 0/84, 0/46 | own finds : 198
  dictionary : 0/0, 0/0, 0/0, 0/0 | imported : 0
  havoc/splice : 1/148, 0/0 | stability : 100.00%
  py/custom/rq : unused, unused, unused, unused
  trim/eff : 0.01%/1121, 0.00%
strategy: explore | state: in progress ^C
```

[cpu000: 12%]

+++ Testing aborted by user +++

AFL++ Demo: “JPEGs out of thin air”

Why is fuzzing slow?

- › Not running in persistent mode: process is constantly restarted by AFL.
- › Running in docker/VM, this slows things down.
- › Only using single CPU, not fuzzing on each CPU in parallel.

How to discover (more) bugs?

- › Also enable sanitizers during compilation!
- › This slows down fuzzing... but can now find more bugs.

How to fuzz libraries?

Not every program/library accepts input from stdin

- › Need to write custom harness to provide fuzzed inputs to the code/function being tested
- › Typical harness structure:

```
int main() {  
    /* 1. Initialize library */  
    /* 2. Read input from fuzzer */  
    /* 3. Call function being fuzzed with given input */  
}
```

Example harness: fuzzing capstone disassembler

```
int main(int argc, char** argv) {
    uint8_t buf[128]; // The buffer we will pass to the library
    csh h; cs_insn *insn; size_t count;

    ssize_t numread = read(stdin, buf, 128);
    if (cs_open(CS_ARCH_X86, CS_MODE_64, &h) == CS_ERR_OK) {
        count = cs_disasm(h, buf, numread, 0x1000, 0, &insn);
        cs_free(insn, count); // clean up after ourselves
    } else return -1;
    cs_close(&h); // close the capstone library
}
```

Writing fuzz harness can be hard

- › Need to understand library, know which function to call,...
- › Harness may need updates to support new library version

Alternative: developer writes the harness! Use libFuzzer:

- › Developer adds `LLVMFuzzerTestOneInput` function
- › This function is part of the project. Can initialize needed functionality and call the code to be tested.
- › Integrated into the latest clang compiler.

libFuzzer: demo

```
bool FuzzMe(const uint8_t *Data, size_t DataSize) {  
    return DataSize >= 3 && Data[0] == 'F' &&  
        Data[1] == 'U' && Data[2] == 'Z' && Data[3] == 'Z';  
}  
  
extern "C" int LLVMFuzzerTestOneInput(const uint8_t *data, size_t len) {  
    FuzzMe(data, len);  
    return 0;  
}  
  
// clang++ -g -fsanitize=address,fuzzer fuzzme.cpp  
// ./a.out
```


KLEE demos

1. check-sign

- » `-posix-runtime` and `-sym-arg X` to provide a command-line argument consisting of `X` symbolic bytes.

2. strtol-sign



- » `-libc=uclibc` to symbolically analyze libc functions.

3. symmalloc.c

- » KLEE limitation: cannot handle symbolic malloc sizes.
- » Length will be concretized.

General remarks

- › Docker images take up space:

| <input type="checkbox"/> | Name | Tag | Status | Created | Size |
|--------------------------|---|--------|------------------------|--------------|----------|
| <input type="checkbox"/> | aflplusplus/aflplusplus 13fdf7e8b621  | latest | Unused | 8 days ago | 3.06 GB |
| <input type="checkbox"/> | klee/klee 673d72c9b2d7  | latest | In use | 5 months ago | 10.44 GB |

- › First do AFL, then do KLEE. Periodically remove temp files.
- › Windows: first start Docker, then run commands. Otherwise you get errors when starting the instances.

References

Optional reading:

- › [Fuzzing capstone using AFL persistent mode](#)
- › [LLVMFuzzerTestOneInput integrated into capstone](#)