

# Fuzzing Network Protocol Implementations (Lecture 6)

Prof. Mathy Vanhoef

DistriNet – KU Leuven – Belgium

# Types of protocols

## 1. Stateless protocols

- » Previous packets don't influence response to current packet
- » Examples: ARP, ICMP, DNS, HTTP (without cookies),...

## 2. Stateful protocols

- » Code being executed depends on current & previous input
- » Examples: TLS, WPA2, IMAP, SMTP, FTP,...

# Fuzzing stateless protocols

How can we fuzz them?

- › Make them behave like a ~~network service~~ program that processes ~~an incoming packet~~ a file / given input
- › We can then apply the typical fuzzers like AFL

# Fuzzing stateless protocols: Preeny's desock

Preeny's desock module:

- › Converts a network service into a command-line program
- › Intercepts (= “hooks”) socket library functions:
  - ›› Intercepts `socket()`, `bind()`, `listen()`, and `accept()`
  - ›› Returns sockets that are synchronized to `stdin` and `stdout`
  - ›› Use `LD_PRELOAD` to load the desock module when starting a program
- › Popular in practice. Included in [AFL++](#) (called [desock\\_dup](#)).

→ Can be combined with many fuzzers!

# Introduction: stateful protocols

We will now focus on **stateful** network protocols



- › Code being executed depends on current & previous input:

```
void handle_packet(uint8_t *p, size_t len) {  
    switch (current_state) {  
        case INIT: handle_init(p, len); break;  
        case AUTH: handle_auth(p, len); break;  
        // ... other states here ...  
        case DATA: handle_data(p, len); break;  
    }  
}
```

# What about network protocols?

What makes fuzzing **stateful** network protocols special?

- › The code being executed depends on previous input
- › This is in contrast with, e.g., image parsing tools



# How is this handled in practice?

One approach is to fuzz individual functions

- › “Unit test(s)” provide mutated input to an internal function(s)
- › Internal function corresponds a certain state in the protocol
- › The typical fuzzers can then be used again

This is currently a common technique in practice:

- › For instance, used by OpenSSL (TLS) and hostap (Wi-Fi)
- › Downside is that many “unit tests” must be written

# Practical example 1: OpenSSL

The OpenSSL library:

- › Implements the TLS/SSL protocol
- › Commonly used in HTTPS

The TLS protocol consists of two phases:

1. Handshake protocol: negotiate keys, many crypto operations
2. Record protocol: exchange encrypted data. Records with an invalid message authentication code (“MAC”) are dropped.



# Practical example 1: OpenSSL

How is the OpenSSL library being fuzzed?

› Build-in support to fuzz using libFuzzer and AFL

```
int FuzzerTestOneInput(const uint8_t *buf, size_t len)
{
    SSL *client = NULL;
    BIO *in;
    SSL_CTX *ctx;

    if (len == 0)
        return 0;

    /* This only fuzzes the initial flow from the client so far. */
    ctx = SSL_CTX_new(SSLv23_method());
    if (ctx == NULL)
        goto end;

    client = SSL_new(ctx);
    if (client == NULL)
        goto end;
    OPENSSL_assert(SSL_set_min_proto_version(client, 0) == 1);
    OPENSSL_assert(SSL_set_cipher_list(client, "ALL:NULL:@SECLEVEL=0") == 1);
    SSL_set_client_cert_name(client, "localhost");
    in = BIO_new(BIO_s_mem());
    if (in == NULL)
        goto end;
    out = BIO_new(BIO_s_mem());
    if (out == NULL) {
        BIO_free(in);
        goto end;
    }
    SSL_set_bio(client, in, out);
    SSL_set_connect_state(client);
    OPENSSL_assert((size_t)BIO_write(in, buf, len) == len);
    if (SSL_do_handshake(client) == 1) {
        /* Keep reading application data until error or EOF. */
        uint8_t tmp[1024];
        while (1) {
            if (SSL_read(client, tmp, sizeof(tmp)) <= 0) {
                break;
            }
        }
    }
    SSL_free(client);
    ERR_clear_error();
    SSL_CTX_free(ctx);

    return 0;
}
```

```
int FuzzerTestOneInput(const uint8_t *buf, size_t len)
```

```
{
```

```
SSL *client = NULL;
```

```
...
```

```
SSL_set_connect_state(client);
```

```
OPENSSL_assert((size_t)BIO_write(in, buf, len) == len);
```

```
if (SSL_do_handshake(client) == 1) {
```

```
/* Keep reading application data until error or EOF.
```

```
uint8_t tmp[1024];
```

# Practical example 1: OpenSSL

How is the OpenSSL library being fuzzed?

- › Build-in support to fuzz using libFuzzer and AFL
- › An initial input corpus is also provided
- › Handling crypto operations deterministically:
  - ›› Uses option to disable generation of random numbers
  - ›› Checks of correct MACs is not disabled
- › Continuously being fuzzed by Google's OSS-Fuzz

## Side note: Google ClusterFuzz and OSS-Fuzz

ClusterFuzz: a framework to fuzz on large clusters

- › The OSS-Fuzz instance runs on 100,000+ machines
- › Has automatic bug reporting, support for various fuzzers, etc.

OSS-Fuzz: fuzzing of open-source software

- › Build on top of ClusterFuzz. Contains scripts to fuzz each project (building the project & giving it mutated input)
- › Found over 40,500 bugs in 650 open-source projects

## Practical example 2: Wi-Fi hostap

Similar to OpenSSL:

- › Support to fuzz using libFuzzer and AFL
- › Initial input corpus is provided
- › Continuously being fuzzed by Google's OSS-Fuzz
- › No special code changes to facilitate fuzzing

# Practical example 2: Wi-Fi hostap

How does hostap tackle state?

- › Fuzzes functions that handle various connection stages

path: `root/tests/fuzzing`

Name  
README  
`ap-mgmt`  
`asn1`  
`build-test.sh`  
`dpp-uri`  
`eap-aka-peer`  
`eap-mschapv2-peer`  
`eap-sim-peer`  
`eapol-key-auth`  
`eapol-key-supp`  
`eapol-supp`

Relevant example “fuzz unit tests”:

1. `ap-mgmt`: frames sent at start of handshake
2. `eap-mschapv2-peer`: authentication protocol
3. `eapol-key-auth`: session key negotiation stage

# Limitations

Several challenges remain:

- › Must write “unit test” for every (important) function
  - ›› May require deep knowledge of the implementation
  - ›› Interactions between different functions may be missed
  - ›› Not easy to assure that all states are covered
- › Handling checksums or message authentication codes
- › Need good initial input corpus
  - ›› Hostap only has a single input frame for some unit tests...

# Properly fuzzing stateless protocols

There are effectively two input grammars to consider:

1. The grammar defining the allowed **format of packets**
2. The grammar defining the allowed **order of packets**

How to explore both aspects while fuzzing/testing?

- › Assume one grammar is known & explore the other
  - ›› For instance: using state inference tools
- › Many other options exists as well...

# Fuzzers for stateful systems: an overview<sup>1</sup>

1. Grammar-based (generational)
2. Grammar learner
3. Evolutionary
4. Evolutionary grammar-based
5. Evolutionary grammar-learner
6. Machine learning-based
7. Man-in-the-middle based

Typical components:

- › Test harness
- › SUT
- › Anomaly detector

<sup>1</sup> Source: “[An overview of Stateful Fuzzing](#)” by Seyed Andarzian, Cristian Daniele, and Erik Poll.



# Fuzzers for stateful systems: an overview<sup>1</sup>

1. **Grammar-based (generational)**
2. **Grammar learner**
3. Evolutionary
4. Evolutionary grammar-based
5. Evolutionary grammar-learner
6. Machine learning-based
7. Man-in-the-middle based

Typical components:


- › Test harness
- › SUT
- › Anomaly detector

<sup>1</sup> Source: “[An overview of Stateful Fuzzing](#)” by Seyed Andarzian, Cristian Daniele, and Erik Poll.

# Grammar-based fuzzing

Define packet layout and state machine

- › Fuzzer then sends valid packets to reach a target state
- › When in the target state, send malformed/mutated packets



Example:

- › Send packets to reach each state, will eventually test state  $S_5$
- › Then send mutations of  $m_5$
- › Will eventually detect the crash?

# Other grammar-based fuzzers

**BooFuzz** protocol fuzzer (fork/successor of Sulley)

1. The structure of each message is first defined

```
user = Request("user", children=(
```

```
String("key", "USER"),  
  Delim("space", " "),  
  String("val", "anonymous"),  
  Static("end", "\r\n")  
))
```

**Name of the message**

**Name of the field**

**Default field value**

**Field type: impacts mutation during fuzzing**



# Other grammar-based fuzzers

**BooFuzz** protocol fuzzer (fork/successor of Sulley)

1. The structure of each message is first defined

```
passwd = Request("passwd", children=(  
    String("key", "PASS"),  
    Delim("space", " "),  
    String("val", "james"),  
    Static("end", "\r\n"),
```

)) **Note: these block-based grammars are inspired by the (underdocumented?) SPIKE fuzzer**




# Other grammar-based fuzzers

**BooFuzz** protocol fuzzer (fork/successor of Sulley)

2. State machine is defined by connecting messages

```
session.connect(user)
session.connect(user, passwd)
session.connect(passwd, stor)
session.connect(passwd, retr)
```



→ user is sent before fuzzing passwd, user and passwd is sent before fuzzing stor or retr. Doesn't fuzz order of messages.

# Other grammar-based fuzzers



## **Peach** network protocol fuzzer

- › Like Sulley/BooFuzz, but uses XML for the grammar
- › Initially an open-source project.

## Commercial edition received updates and features

- › There (was) a community edition, but it lacked such updates
- › GitLab open-sourced core engine of commercial Peach (2021)
  - › Known as the [GitLab Protocol Fuzzer Community Edition](#)
  - › The commercial version is no longer available...?

## Other grammar-based fuzzers



### **Peach** network protocol fuzzer

- › Like Sulley/BooFuzz, but uses XML for the grammar

Main mutation strategies of Peach:

1. Random: selects  $n$  fields from the data model. These fields are modified using a random mutator function.
2. Sequential: all fields are mutated in order using all possible mutator functions.

Limitation: the order of messages isn't fuzzed.

# Grammar learning: state machine inference

Other downsides of BooFuzz and Peach:

- › Must manually specify state machine *and* packet format
- › Can't detect new interesting inputs based on code coverage

First point can be improved by **inferring the state machine**

- › Will still need to specify packet formats, but the state machine of the implementation is automatically inferred.
- › Can manually or (semi-)automatically inspect the inferred state machine and then use it in BooFuzz or Peach.



# Black-box state inference

Common method is to use algorithms for automata learning

- › **Actively** interact with the SUT to learn its behavior
- › Send packets in a random order and inspect the responses

Infer the state machine based on the responses


- › We can then inspect & use the state machine
- › Successfully applied to discover bugs in TLS, SSH, WPA2,...

# Intuitive intro to state machine inference

Start with traces of length one:

- › ClientHello / ServerHello
- › Update state machine
- › Other packets / FatalAlert+Close
- › Update state machine

Other messages /  
FatalAlert+Close




Traces of length two:

- › ClientHello / Server Hello, ClientHello / FatalAlert+Close
- › Update state machine to handle this case

# Intuitive intro to state machine inference

Start with traces of length one:

- › ClientHello / ServerHello
- › Update state machine
- › Other packets / FatalAlert+Close
- › Update state machine




Traces of length two:

- › ClientHello / Server Hello, ClientHello / FatalAlert+Close
- › Update state machine to handle this case

# Intuitive intro to state machine inference

Continue with traces of length two:


- › Other messages / FatalAlert+Close,  
Any message / empty



# Intuitive intro to state machine inference

Continue with traces of length two:


- › Other messages / FatalAlert+Close,  
Any message / empty
- › ClientHello / ServerHello,  
ClientKeyExchange / empty
- › ClientHello / ServerHello,  
Other messages / FatalAlert+Close



# Intuitive intro to state machine inference

Continue with traces of length two:

- › Other messages / FatalAlert+Close,  
Any message / empty
- › ClientHello / ServerHello,  
ClientKeyExchange / empty
- › ClientHello / ServerHello,  
Other messages / FatalAlert+Close



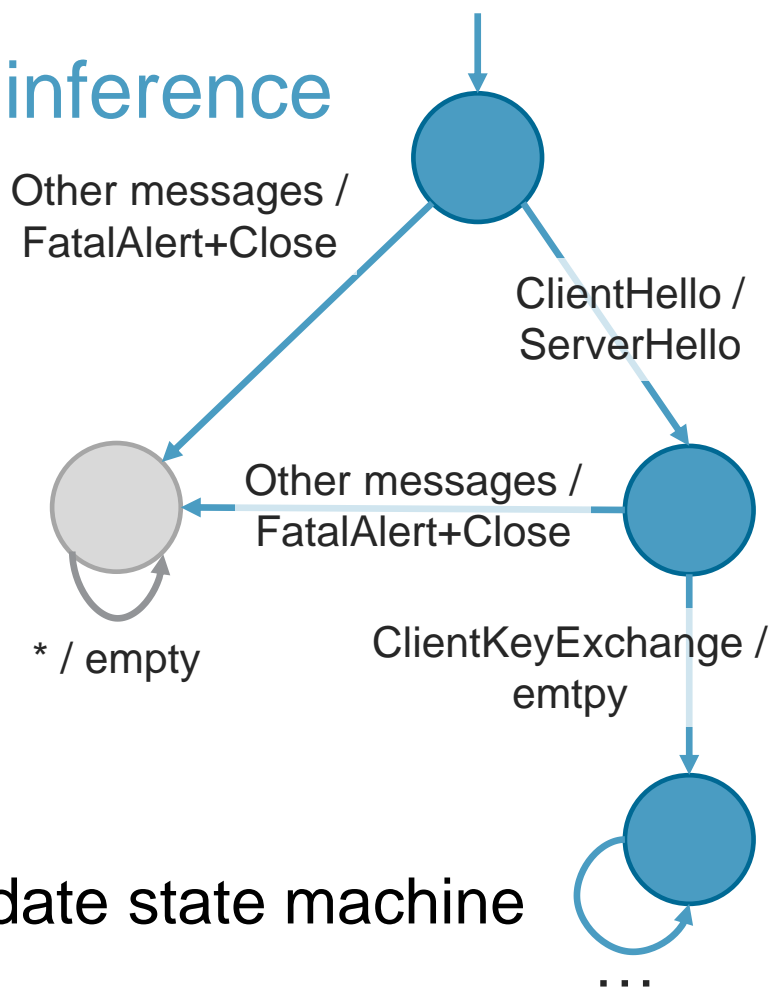
Continue with traces of length 3 & update state machine

# Intuitive intro to state machine inference

Continue with traces of length two:

- › Other messages / FatalAlert+Close,  
Any message / empty
- › ClientHello / ServerHello,  
ClientKeyExchange / empty
- › ClientHello / ServerHello,  
Other messages / FatalAlert+Close

Continue with traces of length 3 & update state machine



# The real learning setup



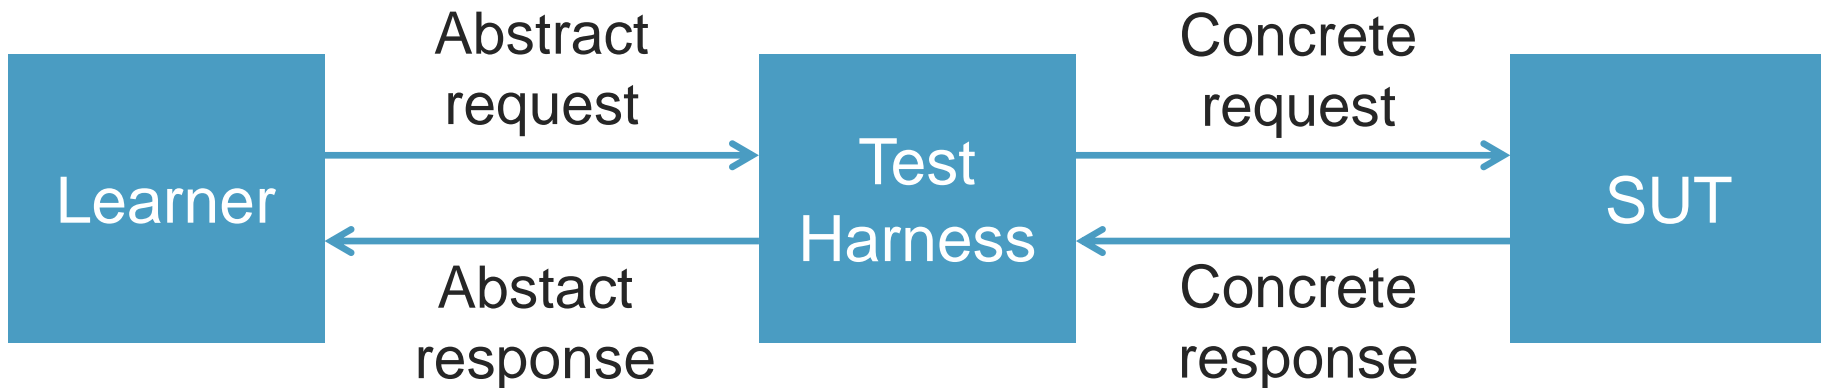
- › What do we learn? A deterministic Mealy machine.
- › How to learn? Use libraries such as LearnLib.
  - ›› They use  $L^*$  or TTT to form a hypothesis for the state machine.
- › Must be able to perform three actions:
  1. Reset the SUT
  2. Send message to SUT & get the output
  3. Check whether the hypothesis (= current state machine) is correct
- › Performing action 2 & 3 is non-trivial!



# Challenge: sending messages to the SUT (1)

State machine uses abstract messages, e.g., “ClientHello”

- › Must be converted to concrete messages = actual bytes!
- › A test harness is used for this conversion:



## Challenge: sending messages to the SUT (2)

The **state harness** must be able to send packets in any order

- › Must consider previous messages that were sent/received
  - › Example: random nonces that were part of the handshake
  - › Example: currently negotiated session key
- › In certain cases, it's unclear which values to use in requests
  - › Example: how to send an encrypted TLS record before a key was negotiated? Use a random key? Use an all-zero key?
- › And we must be able to **receive packets in any order**
  - › Example edge case: we receive an encrypted packet before a key was negotiated. Do we try to decrypt it? With which key?

# Challenge: is the state machine correct?

Learning algorithms need a way to check if their current hypothesis for the state machine is correct

- › But we don't know the state machine...

Two typical solutions:

1. **Random traces**: send some fixed number of random traces and see if the responses match the state machine
2. **Chow's W-method**: guarantees correctness of the state machine given an upper bound on the number of states

# Why is state inference useful?

Use the inferred state machine in **BooFuzz or similar**

- › You will now fuzz the *actual* states of the implementation
- › This may be more/other/different states than in the standard!

**Manually inspect** the state machine for flaws

- › Identified flaws in TLS, DTLS, WPA2, and 4G/LTE

State machine may form a **fingerprint** of the implementation

- › Use unique behavior to detect implementation being used

## Last but not least: frameworks for manual tests

Sometimes you want to manually test specific behavior

- › To reimplement a known attack quickly...
- › You think a library has a specific vulnerability...
- › You want to confirm a flaw in the standard...

This may require implementing large parts of a protocol

- › E.g., might require implementing tedious crypto algorithms
- › Or requires implementing the full handshake if you want to test behavior after authenticating

# TLS testing framework

Use [TLS-Attacker](#) framework for TLS tests:

- › Java-based framework for analyzing TLS libraries
- › Easily define custom TLS protocol flows to test libraries

```
Config config = Config.createConfig();  
WorkflowTrace trace = new WorkflowTrace();  
trace.addTlsAction(new SendAction(new ClientHelloMessage()));  
trace.addTlsAction(new ReceiveAction(new ServerHelloMessage()));  
State state = new State(config, trace);  
DefaultWorkflowExecutor ex = new DefaultWorkflowExecutor(state);  
ex.executeWorkflow();
```

# Testing tools in practice: Wi-Fi security

Wi-Fi Alliance **tests for KRACK** during device certification (2017):

“ industry. There is no evidence of the vulnerability being used against Wi-Fi users maliciously, and Wi-Fi Alliance has taken immediate steps to ensure users can continue to count on Wi-Fi to deliver strong security protections.

- Wi-Fi Alliance now requires testing for this vulnerability within our global certification lab network
- Wi-Fi Alliance has provided a **vulnerability detection tool** for use by any Wi-Fi Alliance member

”

# Testing tools in practice: Wi-Fi security

Wi-Fi Alliance also **seems to test for FragAttacks** flaws (2021):

“ Wi-Fi Alliance has taken immediate steps to ensure users can remain confident in the strong security protections provided by Wi-Fi.

- Wi-Fi CERTIFIED now includes additional testing within our global certification lab network to encourage greater adoption of recommended practices
- Wi-Fi Alliance is broadly communicating implementation guidance to device vendors [..] ”



## Optional reading

- › “[Protocol State Fuzzing of TLS Implementations](#)” by Joeri de Ruiter and Erik Poll, USENIX Security 2015